

CIS Red Hat OpenShift Container Platform Benchmark

v1.9.0 - 12-11-2025

Terms of Use

Please see the below link for our current terms of use:

<https://www.cisecurity.org/cis-securesuite/cis-securesuite-membership-terms-of-use/>

For information on referencing and/or citing CIS Benchmarks in 3rd party documentation (including using portions of Benchmark Recommendations) please contact CIS Legal (legalnotices@cisecurity.org) and request guidance on copyright usage.

NOTE: It is **NEVER** acceptable to host a CIS Benchmark in **ANY** format (PDF, etc.) on a 3rd party (non-CIS owned) site.

Table of Contents

<i>Terms of Use</i>	1
<i>Table of Contents</i>	2
<i>Overview</i>	4
Important Usage Information	4
Target Technology Details	7
Intended Audience	7
Consensus Guidance	8
Typographical Conventions	9
<i>Recommendation Definitions</i>	10
Title	10
Assessment Status	10
Automated	10
Manual	10
Profile	10
Description	10
Rationale Statement	10
Impact Statement	11
Audit Procedure	11
Remediation Procedure	11
Default Value	11
References	11
CIS Critical Security Controls® (CIS Controls®)	11
Additional Information	11
Profile Definitions	12
Acknowledgements	Error! Bookmark not defined.
<i>Recommendations</i>	14
<i>Appendix: Summary Table</i>	295
<i>Appendix: CIS Controls v7 IG 1 Mapped Recommendations</i>	305
<i>Appendix: CIS Controls v7 IG 2 Mapped Recommendations</i>	309
<i>Appendix: CIS Controls v7 IG 3 Mapped Recommendations</i>	315
<i>Appendix: CIS Controls v7 Unmapped Recommendations</i>	322
<i>Appendix: CIS Controls v8 IG 1 Mapped Recommendations</i>	323

Appendix: CIS Controls v8 IG 2 Mapped Recommendations 327
Appendix: CIS Controls v8 IG 3 Mapped Recommendations 333
Appendix: CIS Controls v8 Unmapped Recommendations..... 340
Appendix: Change History..... Error! Bookmark not defined.

Overview

All CIS Benchmarks™ (Benchmarks) focus on technical configuration settings used to maintain and/or increase the security of the addressed technology, and they should be used in **conjunction** with other essential cyber hygiene tasks like:

- Monitoring the base operating system and applications for vulnerabilities and quickly updating with the latest security patches.
- End-point protection (Antivirus software, Endpoint Detection and Response (EDR), etc.).
- Logging and monitoring user and system activity.

In the end, the Benchmarks are designed to be a key **component** of a comprehensive cybersecurity program.

Important Usage Information

All Benchmarks are available free for non-commercial use from the [CIS Website](#). They can be used to manually assess and remediate systems and applications. In lieu of manual assessment and remediation, there are several tools available to assist with assessment:

- [CIS Configuration Assessment Tool \(CIS-CAT® Pro Assessor\)](#)
- [CIS Benchmarks™ Certified 3rd Party Tooling](#)

These tools make the hardening process much more scalable for large numbers of systems and applications.

NOTE: Some tooling focuses only on the Benchmark Recommendations that can be fully automated (skipping ones marked **Manual**). It is important that **ALL** Recommendations (**Automated** and **Manual**) be addressed since all are important for properly securing systems and are typically in scope for audits.

Key Stakeholders

Cybersecurity is a collaborative effort, and cross functional cooperation is imperative within an organization to discuss, test, and deploy Benchmarks in an effective and efficient way. The Benchmarks are developed to be best practice configuration guidelines applicable to a wide range of use cases. In some organizations, exceptions to specific Recommendations will be needed, and this team should work to prioritize the problematic Recommendations based on several factors like risk, time, cost, and labor. These exceptions should be properly categorized and documented for auditing purposes.

Apply the Correct Version of a Benchmark

Benchmarks are developed and tested for a specific set of products and versions and applying an incorrect Benchmark to a system can cause the resulting pass/fail score to be incorrect. This is due to the assessment of settings that do not apply to the target systems. To assure the correct Benchmark is being assessed:

- **Deploy the Benchmark applicable to the way settings are managed in the environment:** An example of this is the Microsoft Windows family of Benchmarks, which have separate Benchmarks for Group Policy, Intune, and Stand-alone systems based upon how system management is deployed. Applying the wrong Benchmark in this case will give invalid results.
- **Use the most recent version of a Benchmark:** This is true for all Benchmarks, but especially true for cloud technologies. Cloud technologies change frequently and using an older version of a Benchmark may have invalid methods for auditing and remediation.

Exceptions

The guidance items in the Benchmarks are called recommendations and not requirements, and exceptions to some of them are expected and acceptable. The Benchmarks strive to be a secure baseline, or starting point, for a specific technology, with known issues identified during Benchmark development are documented in the Impact section of each Recommendation. In addition, organizational, system specific requirements, or local site policy may require changes as well, or an exception to a Recommendation or group of Recommendations (e.g. A Benchmark could Recommend that a Web server not be installed on the system, but if a system's primary purpose is to function as a Webserver, there should be a documented exception to this Recommendation for that specific server).

In the end, exceptions to some Benchmark Recommendations are common and acceptable, and should be handled as follows:

- The reasons for the exception should be reviewed cross-functionally and be well documented for audit purposes.
- A plan should be developed for mitigating, or eliminating, the exception in the future, if applicable.
- If the organization decides to accept the risk of this exception (not work toward mitigation or elimination), this should be documented for audit purposes.

It is the responsibility of the organization to determine their overall security policy, and which settings are applicable to their unique needs based on the overall risk profile for the organization.

Remediation

CIS has developed [Build Kits](#) for many technologies to assist in the automation of hardening systems. Build Kits are designed to correspond to Benchmark's "Remediation" section, which provides the manual remediation steps necessary to make that Recommendation compliant to the Benchmark.

When remediating systems (changing configuration settings on deployed systems as per the Benchmark's Recommendations), please approach this with caution and test thoroughly.

The following is a reasonable remediation approach to follow:

- CIS Build Kits, or internally developed remediation methods should never be applied to production systems without proper testing.
- Proper testing consists of the following:
 - Understand the configuration (including installed applications) of the targeted systems. Various parts of the organization may need different configurations (e.g., software developers vs standard office workers).
 - Read the Impact section of the given Recommendation to help determine if there might be an issue with the targeted systems.
 - Test the configuration changes with representative lab system(s). If issues arise during testing, they can be resolved prior to deploying to any production systems.
 - When testing is complete, initially deploy to a small sub-set of production systems and monitor closely for issues. If there are issues, they can be resolved prior to deploying more broadly.
 - When the initial deployment above is completed successfully, iteratively deploy to additional systems and monitor closely for issues. Repeat this process until the full deployment is complete.

Summary

Using the Benchmarks Certified tools, working as a team with key stakeholders, being selective with exceptions, and being careful with remediation deployment, it is possible to harden large numbers of deployed systems in a cost effective, efficient, and safe manner.

NOTE: As previously stated, the PDF versions of the CIS Benchmarks™ are available for free, non-commercial use on the [CIS Website](#). All other formats of the CIS Benchmarks™ (MS Word, Excel, and [Build Kits](#)) are available for CIS [SecureSuite](#)® members.

CIS-CAT® Pro is also available to CIS [SecureSuite](#)® members.

Target Technology Details

This document provides prescriptive guidance for establishing a secure configuration posture for OpenShift 4.

The set of configuration files mentioned throughout this benchmark are specific to Red Hat's CNCF certified Kubernetes distribution, Red Hat OpenShift Container Platform.

Each section includes information about the default configuration of an OpenShift cluster and a set of recommendations for hardening the configuration, inspired by the CIS Kubernetes benchmark. For each hardening recommendation, information on how to implement the control and/or how to verify or audit the control is provided. In some cases, remediation information is also provided.

The majority of the settings in the hardening guide are in place by default. The audit information for these settings is provided so that you can verify that the cluster admin has not made changes that would be less secure than the OpenShift defaults. A small number of items require configuration. Finally, there are some recommendations that require decisions by the customer, such as audit log size, retention and related settings.

The recommendations that require decisions based on your needs are:

- Configure encryption of data at rest in etcd datastore
- Manage Image Provenance
- Set the `--event-qps` argument as appropriate
- Configure the API server audit log retention
- Configure cluster logging to forward audit logs off the cluster
- Configure the audit log file size
- Adjust the garbage collection settings as needed
- Create custom Security Context Constraints as needed
- Configure Network Policies as appropriate
- In OCP 4.6 and above, configure audit policies as appropriate

Intended Audience

This document is intended for system and application administrators, security specialists, auditors, help desk, and platform deployment personnel who plan to develop, deploy, assess, or secure solutions that incorporate OpenShift 4.

Consensus Guidance

This CIS Benchmark™ was created using a consensus review process comprised of a global community of subject matter experts. The process combines real world experience with data-based information to create technology specific guidance to assist users to secure their environments. Consensus participants provide perspective from a diverse set of backgrounds including consulting, software development, audit and compliance, security research, operations, government, and legal.

Each CIS Benchmark undergoes two phases of consensus review. The first phase occurs during initial Benchmark development. During this phase, subject matter experts convene to discuss, create, and test working drafts of the Benchmark. This discussion occurs until consensus has been reached on Benchmark recommendations. The second phase begins after the Benchmark has been published. During this phase, all feedback provided by the Internet community is reviewed by the consensus team for incorporation in the Benchmark. If you are interested in participating in the consensus process, please visit <https://workbench.cisecurity.org/>.

Typographical Conventions

The following typographical conventions are used throughout this guide:

Convention	Meaning
<code>Stylized Monospace font</code>	Used for blocks of code, command, and script examples. Text should be interpreted exactly as presented.
<code>Monospace font</code>	Used for inline code, commands, UI/Menu selections or examples. Text should be interpreted exactly as presented.
<code><Monospace font in brackets></code>	Text set in angle brackets denote a variable requiring substitution for a real value.
<i>Italic font</i>	Used to reference other relevant settings, CIS Benchmarks and/or Benchmark Communities. Also, used to denote the title of a book, article, or other publication.
Bold font	Additional information or caveats things like Notes , Warnings , or Cautions (usually just the word itself and the rest of the text normal).

Recommendation Definitions

The following defines the various components included in a CIS recommendation as applicable. If any of the components are not applicable it will be noted, or the component will not be included in the recommendation.

Title

Concise description for the recommendation's intended configuration.

Assessment Status

An assessment status is included for every recommendation. The assessment status indicates whether the given recommendation can be automated or requires manual steps to implement. Both statuses are equally important and are determined and supported as defined below:

Automated

Represents recommendations for which assessment of a technical control can be fully automated and validated to a pass/fail state. Recommendations will include the necessary information to implement automation.

Manual

Represents recommendations for which assessment of a technical control cannot be fully automated and requires all or some manual steps to validate that the configured state is set as expected. The expected state can vary depending on the environment.

Profile

A collection of recommendations for securing a technology or a supporting platform. Most benchmarks include at least a Level 1 and Level 2 Profile. Level 2 extends Level 1 recommendations and is not a standalone profile. The Profile Definitions section in the benchmark provides the definitions as they pertain to the recommendations included for the technology.

Description

Detailed information pertaining to the setting with which the recommendation is concerned. In some cases, the description will include the recommended value.

Rationale Statement

Detailed reasoning for the recommendation to provide the user a clear and concise understanding on the importance of the recommendation.

Impact Statement

Any security, functionality, or operational consequences that can result from following the recommendation.

Audit Procedure

Systematic instructions for determining if the target system complies with the recommendation.

Remediation Procedure

Systematic instructions for applying recommendations to the target system to bring it into compliance according to the recommendation.

Default Value

Default value for the given setting in this recommendation, if known. If not known, either not configured or not defined will be applied.

References

Additional documentation relative to the recommendation.

CIS Critical Security Controls® (CIS Controls®)

The mapping between a recommendation and the CIS Controls is organized by CIS Controls version, Safeguard, and Implementation Group (IG). The Benchmark in its entirety addresses the CIS Controls safeguards of (v7) "5.1 - Establish Secure Configurations" and (v8) "4.1 - Establish and Maintain a Secure Configuration Process" so individual recommendations will not be mapped to these safeguards.

Additional Information

Supplementary information that does not correspond to any other field but may be useful to the user.

Profile Definitions

The following configuration profiles are defined by this Benchmark:

- **Level 1**

Items in this profile intend to:

- be practical and prudent;
- provide a clear security benefit; and
- not inhibit the utility of the technology beyond acceptable means.

- **Level 2**

This profile extends the "Level 1" profile. Items in this profile exhibit one or more of the following characteristics:

- are intended for environments or use cases where security is paramount
- acts as defense in depth measure
- may negatively inhibit the utility or performance of the technology.

Acknowledgements

This Benchmark exemplifies the great things a community of users, vendors, and subject matter experts can accomplish through consensus collaboration. The CIS community thanks the entire consensus team with special recognition to the following individuals who contributed greatly to the creation of this guide:

This benchmark exemplifies the great things a community of users, vendors, and subject matter experts can accomplish through consensus collaboration. The CIS community thanks the entire consensus team with special recognition to the following individuals who contributed greatly to the creation of this guide:

Author/s

Randall Mowen
Kirsten Newcomer

Editor/s

Randall Mowen
Lance Bragstad

Contributor/s

The Red Hat Openshift Team
Mark Larinde
Jakub Hrozek
Rutvik Kshirsagar

Special thanks to: Lance Bragstad

Recommendations

1 Control Plane Components

This section consists of security recommendations for the direct configuration of Kubernetes control plane processes. These recommendations assume that the OpenShift cluster has 3 master nodes, as that is the default configuration on installation. These recommendations may not be directly applicable for cluster operators in environments where these components are managed by a 3rd party such as OpenShift Dedicated, Azure Red Hat OpenShift or Red Hat OpenShift Service on AWS.

All Audit and Remediation commands assume that you are logged into the OpenShift cluster with the cluster admin role, cluster bound.

1.1 Master Node Configuration Files

1.1.1 Ensure that the API server pod specification file permissions are set to 600 or more restrictive (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the API server pod specification file has permissions of **600** or more restrictive.

Rationale:

The API server pod specification file controls various parameters that set the behavior of the API server. You should restrict its file permissions to maintain the integrity of the file. The file should be writable only by the administrators on the system.

Impact:

None.

Audit:

OpenShift 4 deploys two API servers: the OpenShift API server and the Kube API server. The OpenShift API server delegates requests for Kubernetes objects to the Kube API server.

The OpenShift API server is managed as a deployment. The pod specification yaml for openshift-apiserver is stored in etcd.

The Kube API Server is managed as a static pod. The pod specification file for the kube-apiserver is created on the control plane nodes at /etc/kubernetes/manifests/kube-apiserver-pod.yaml. The kube-apiserver is mounted via hostpath to the kube-apiserver pods via /etc/kubernetes/static-pod-resources/kube-apiserver-pod.yaml with permissions **600**.

To verify pod specification file permissions for the kube-apiserver, run the following command.

```
for i in $( oc get pods -n openshift-kube-apiserver -l app=openshift-kube-apiserver -o name )
do
    oc exec -n openshift-kube-apiserver $i -- \
    stat -c %a /etc/kubernetes/static-pod-resources/kube-apiserver-pod.yaml
done
```

Verify that the permissions are **600** or more restrictive.

Remediation:

There is no remediation for updating the permissions of `kube-apiserver-pod.yaml`. The file is owned by an OpenShift operator and any changes to the file will result in a degraded cluster state.

Please do not attempt to remediate the permissions of this file.

Default Value:







By default, in OpenShift 4.14, the `kube-apiserver-pod.yaml` has permissions of `600`.

In older versions of OpenShift, the `kube-apiserver-pod.yaml` has permissions of `644`, and is not remediable. Please upgrade to OpenShift 4.14 when possible.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#defining-masters_control-plane
2. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
3. <https://github.com/openshift/library-go/commit/19a42d2bae8ba68761cfad72bf764e10d275ad6e>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.1.2 Ensure that the API server pod specification file ownership is set to root:root (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the API server pod specification file ownership is set to **root:root**.

Rationale:

The API server pod specification file controls various parameters that set the behavior of the API server. You should set its file ownership to maintain the integrity of the file. The file should be owned by **root:root**.

Impact:

None

Audit:

OpenShift 4 deploys two API servers: the OpenShift API server and the Kube API server.

The OpenShift API server is managed as a deployment. The pod specification yaml for openshift-apiserver is stored in etcd.

The Kube API Server is managed as a static pod. The pod specification file for the kube-apiserver is created on the control plane nodes at /etc/kubernetes/manifests/kube-apiserver-pod.yaml. The kube-apiserver is mounted via hostpath to the kube-apiserver pods via /etc/kubernetes/static-pod-resources/kube-apiserver-pod.yaml with ownership **root:root**.

To verify pod specification file ownership for the kube-apiserver, run the following command.

```
#echo "check kube-apiserver pod specification file ownership"

for i in $( oc get pods -n openshift-kube-apiserver -l app=openshift-kube-apiserver -o name )
do
    oc exec -n openshift-kube-apiserver $i -- \
    stat -c %U:%G /etc/kubernetes/static-pod-resources/kube-apiserver-pod.yaml
done
```

Verify that the ownership is set to **root:root**.

Remediation:

No remediation required; file permissions are managed by the operator.







Default Value:

By default, in OpenShift 4, the kube-apiserver-pod.yaml file ownership is set to **root:root**.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#defining-masters_control-plane
2. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
3. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 Ensure the Use of Dedicated Administrative Accounts Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

1.1.3 Ensure that the controller manager pod specification file permissions are set to 600 or more restrictive (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the controller manager pod specification file has permissions of **600** or more restrictive.

Rationale:

The controller manager pod specification file controls various parameters that set the behavior of the Controller Manager on the master node. You should restrict its file permissions to maintain the integrity of the file. The file should be writable by only the administrators on the system.

Impact:

None.

Audit:

OpenShift 4 deploys two controller managers: the OpenShift Controller manager and the Kube Controller manager.

The OpenShift Controller manager is managed as a deployment. The pod specification yaml for openshift-controller-manager is stored in etcd.

The Kube Controller manager is managed as a static pod. The pod specification file for the openshift-kube-controller-manager is created on control plane nodes at `/etc/kubernetes/manifests/kube-controller-manager-pod.yaml`. It is mounted via hostpath to the kube-controller-manager pods via `/etc/kubernetes/static-pod-resources/kube-controller-manager-pod.yaml` with permissions 0644.

To verify pod specification file permissions for the kube-controller-manager, run the following command.

```
for i in $( oc get pods -n openshift-kube-controller-manager -o name -l
app=kube-controller-manager)
do
    oc exec -n openshift-kube-controller-manager $i -- \
    stat -c %a /etc/kubernetes/static-pod-resources/kube-controller-manager-
pod.yaml
done
```

Verify that the permissions are **600** or more restrictive.

Remediation:

There is no remediation for updating the permissions of `kube-controller-manager-pod.yaml`. The file is owned by an OpenShift operator and any changes to the file will result in a degraded cluster state.

Please do not attempt to remediate the permissions of this file.

Default Value:







By default, in OpenShift 4.14, the `kube-controller-manager-pod.yaml` has permissions of `600`.

In older versions of OpenShift, the `kube-controller-manager-pod.yaml` has permissions of `644`, and it not remediable. Please upgrade to OpenShift 4.14 when possible.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#defining-masters_control-plane
2. <https://github.com/openshift/library-go/commit/19a42d2bae8ba68761cfad72bf764e10d275ad6e>
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-controller-manager/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 <u>Configure Data Access Control Lists</u> Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 <u>Protect Information through Access Control Lists</u> Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.1.4 Ensure that the controller manager pod specification file ownership is set to root:root (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the controller manager pod specification file ownership is set to **root:root**.

Rationale:

The controller manager pod specification file controls various parameters that set the behavior of various components of the master node. You should set its file ownership to maintain the integrity of the file. The file should be owned by **root:root**.

Impact:

None

Audit:

OpenShift 4 deploys two controller managers: the OpenShift Controller manager and the Kube Controller manager.

The OpenShift Controller manager is managed as a deployment. The pod specification yaml for openshift-controller-manager is stored in etcd.

The Kube Controller manager is managed as a static pod. The pod specification file for the openshift-kube-controller-manager is created on control plane nodes at `/etc/kubernetes/manifests/kube-controller-manager-pod.yaml`. It is mounted via hostpath to the kube-controller-manager pods via `/etc/kubernetes/static-pod-resources/kube-controller-manager-pod.yaml` with ownership **root:root**.

Run the following command.

```
#echo "openshift-kube-controller-manager pod specification file ownership"

for i in $( oc get pods -n openshift-kube-controller-manager -o name -l
app=kube-controller-manager)
do
    oc exec -n openshift-kube-controller-manager $i -- \
    stat -c %U:%G /etc/kubernetes/static-pod-resources/kube-controller-manager-
pod.yaml
done
```

Verify that the ownership is set to **root:root**.

Remediation:

No remediation required; file permissions are managed by the operator.







Default Value:

By default, in OpenShift 4, the `kube-controller-manager-pod.yaml` file ownership is set to `root:root`.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#defining-masters_control-plane
2. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#cluster-openshift-controller-manager-operator_red-hat-operators
3. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#kube-controller-manager-operator_red-hat-operators
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-controller-manager/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 Ensure the Use of Dedicated Administrative Accounts Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

1.1.5 Ensure that the scheduler pod specification file permissions are set to 600 or more restrictive (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the scheduler pod specification file has permissions of **600** or more restrictive.

Rationale:

The scheduler pod specification file controls various parameters that set the behavior of the Scheduler service in the master node. You should restrict its file permissions to maintain the integrity of the file. The file should be writable by only the administrators on the system.

Impact:

None.

Audit:

In OpenShift 4 the kube-scheduler is deployed as a static pod and its pod specification file is created on control plane nodes at **/etc/kubernetes/manifests/kube-scheduler-pod.yaml**. It is mounted via hostpath to the kube-controller-manager pods via **/etc/kubernetes/static-pod-resources/kube-scheduler-pod.yaml** with permissions **644**.

To verify, run the following command.

```
for i in $(oc get pods -n openshift-kube-scheduler -l app=openshift-kube-scheduler -o name)
do
  oc exec -n openshift-kube-scheduler $i -- \
    stat -c %a /etc/kubernetes/static-pod-resources/kube-scheduler-pod.yaml
done
```

Verify that the permissions are **600** or more restrictive.

Remediation:

There is no remediation for updating the permissions of **kube-scheduler-pod.yaml**. The file is owned by an OpenShift operator and any changes to the file will result in a degraded cluster state.

Please do not attempt to remediate the permissions of this file.

Default Value:







By default, in OpenShift 4.14, the `kube-scheduler-pod.yaml` has permissions of `600`.

In older versions of OpenShift, the `kube-scheduler-pod.yaml` has permissions of `644`, and is not remediable. Please upgrade to OpenShift 4.14 when possible.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#defining-masters_control-plane
2. <https://github.com/openshift/library-go/commit/19a42d2bae8ba68761cfad72bf764e10d275ad6e>
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-scheduler/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.1.6 Ensure that the scheduler pod specification file ownership is set to root:root (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the scheduler pod specification file ownership is set to **root:root**.

Rationale:

The scheduler pod specification file controls various parameters that set the behavior of the **kube-scheduler** service in the master node. You should set its file ownership to maintain the integrity of the file. The file should be owned by **root:root**.

Impact:

None

Audit:

In OpenShift 4, the kube-scheduler is deployed as a static pod and its pod specification file is created on control plane nodes at `/etc/kubernetes/manifests/kube-scheduler-pod.yaml`. It is mounted via hostpath to the kube-controller-manager pods via `/etc/kubernetes/static-pod-resources/kube-scheduler-pod.yaml` with ownership **root:root**.

Run the following command.

```
#Verify openshift-kube-scheduler ownership
for i in $(oc get pods -n openshift-kube-scheduler -l app=openshift-kube-scheduler -o name)
do
  oc exec -n openshift-kube-scheduler $i -- \
  stat -c %U:%G /etc/kubernetes/static-pod-resources/kube-scheduler-pod.yaml
done
```

Verify that the ownership is set to **root:root**.

Remediation:

No remediation required; file permissions are managed by the operator.







Default Value:

By default, in OpenShift 4, **kube-scheduler-pod.yaml** file ownership is set to **root:root**.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#defining-masters_control-plane
2. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#cluster-kube-scheduler-operator_red-hat-operators
3. <https://docs.openshift.com/container-platform/latest/nodes/scheduling/nodes-scheduler-about.html>
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-scheduler/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 Ensure the Use of Dedicated Administrative Accounts Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

1.1.7 Ensure that the etcd pod specification file permissions are set to 600 or more restrictive (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the `/etc/kubernetes/manifests/etcd.yaml` file has permissions of `600` or more restrictive.

Rationale:

The etcd pod specification file `/etc/kubernetes/manifests/etcd.yaml` controls various parameters that set the behavior of the `etcd` service in the master node. `etcd` is a highly-available key-value store which Kubernetes uses for persistent storage of all of its REST API object. You should restrict its file permissions to maintain the integrity of the file. The file should be writable by only the administrators on the system.

Impact:

None.

Audit:

In OpenShift 4, starting with OpenShift 4.4, the etcd pod specification file is generated by the cluster etcd operator.

The default etcd pod specification file is available here: [openshift/cluster-etcd-operator](https://github.com/openshift/cluster-etcd-operator)

Run the following command.

```
for i in $(oc get pods -n openshift-etcd -l app=etcd -o name | grep etcd )
do
  echo "check pod $i"
  oc rsh -n openshift-etcd $i \
    stat -c %a /etc/kubernetes/manifests/etcd-pod.yaml
done
```

Verify that the permissions are `600` or more restrictive.

Remediation:

There is no remediation for updating the permissions of `etcd-pod.yaml`. The file is owned by an OpenShift operator and any changes to the file will result in a degraded cluster state.

Please do not attempt to remediate the permissions of this file.

Default Value:

By default, in OpenShift 4, `/etc/kubernetes/manifests/etcd-pod.yaml` file has permissions of `644`.







By default, in OpenShift 4.14, the `etcd-pod.yaml` has permissions of `600`.

In older versions of OpenShift, the `etcd-pod.yaml` has permissions of `644`, and is not remediable. Please upgrade to OpenShift 4.14 when possible.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#defining-masters_control-plane
2. <https://github.com/openshift/cluster-etcd-operator/blob/master/bindata/etcd/pod.yaml>
3. <https://etcd.io/>
4. <https://kubernetes.io/docs/tasks/administer-cluster/configure-upgrade-etcd/>
5. <https://github.com/openshift/library-go/commit/19a42d2bae8ba68761cfad72bf764e10d275ad6e>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.1.8 Ensure that the etcd pod specification file ownership is set to root:root (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the `/etc/kubernetes/manifests/etcd.yaml` file ownership is set to `root:root`.

Rationale:

The etcd pod specification file `/etc/kubernetes/manifests/etcd.yaml` controls various parameters that set the behavior of the `etcd` service in the master node. `etcd` is a highly-available key-value store which Kubernetes uses for persistent storage of all of its REST API object. You should set its file ownership to maintain the integrity of the file. The file should be owned by `root:root`.

Impact:

None.

Audit:

In OpenShift 4, starting with OpenShift 4.4, the etcd pod specification file is generated by the cluster etcd operator. The pod specification file is created on control plane nodes at `/etc/kubernetes/manifests/etcd-member.yaml` with ownership `root:root`.

The default etcd pod specification file is available here: [openshift/cluster-etcd-operator](#)

Run the following command :

```
#Verify openshift-etcd ownership
for i in $(oc get pods -n openshift-etcd -l app=etcd -o name | grep etcd )
do
    echo "check pod $i"
    oc rsh -n openshift-etcd $i \
        stat -c %U:%G /etc/kubernetes/manifests/etcd-pod.yaml
done
```

Verify that the ownership is set to `root:root`.

Remediation:

No remediation required; file permissions are managed by the operator.







Default Value:

By default, in OpenShift 4, `/etc/kubernetes/manifests/etcd-member.yaml` file ownership is set to `root:root`.

References:

1. <https://coreos.com/etcd>
2. <https://kubernetes.io/docs/admin/etcd/>
3. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#defining-masters_control-plane

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 Ensure the Use of Dedicated Administrative Accounts Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

1.1.9 Ensure that the Container Network Interface file permissions are set to 600 or more restrictive (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the Container Network Interface files have permissions of **600** or more restrictive.

Rationale:

Container Network Interface provides various networking options for overlay networking. You should consult their documentation and restrict their respective file permissions to maintain the integrity of those files. Those files should be writable by only the administrators on the system.

Impact:

None

Audit:

The Cluster Network Operator (CNO) deploys and manages the cluster network components on an OpenShift Container Platform cluster. This includes the Container Network Interface (CNI) default network provider plug-in selected for the cluster during installation. OpenShift Container Platform uses the Multus CNI plug-in to allow chaining of CNI plug-ins. The default Pod network must be configured during cluster installation. By default, the CNO deploys the OpenShift SDN as the default Pod network.

Ensure that the Container Network Interface file permissions, multus, openshift-sdn and Open vSwitch (OVS) file permissions are set to 644 or more restrictive. The SDN components are deployed as DaemonSets across the master/worker nodes with controllers limited to control plane nodes. OpenShift deploys OVS as a network overlay by default. Various configurations (ConfigMaps and other files managed by the operator via hostpath but stored on the container hosts) are stored in the following locations:

CNI/Multus (pod multus):

/host/etc/cni/net.d = CNI_CONF_DIR /host/var/run/multus/cni/net.d = multus config dir

SDN (pod ovs; daemonset; app=ovs):

/var/lib/cni/networks/openshift-sdn /var/run/openshift-sdn

OVS (container openvswitch):

/var/run/openvswitch /etc/openvswitch /run/openvswitch

Run the following commands to verify the permissions when using CNI multus:

```
for i in $(oc get pods -n openshift-multus -l app=multus -oname); do oc exec -n openshift-multus $i -- /bin/bash -c "stat -c \"%a %n\" /host/etc/cni/net.d/*.conf"; done

for i in $(oc get pods -n openshift-multus -l app=multus -oname); do oc exec -n openshift-multus $i -- /bin/bash -c "stat -c \"%a %n\" /host/var/run/multus/cni/net.d/*.conf"; done
```

Run the following commands to verify the permissions when using OpenShift SDN:

```
for i in $(oc get pods -n openshift-sdn -l app=sdn -oname); do oc exec -n openshift-sdn $i -- find /var/lib/cni/networks/openshift-sdn -type f -exec stat -c %a {} \;; done

for i in $(oc get pods -n openshift-sdn -l app=sdn -oname); do oc exec -n openshift-sdn $i -- find /var/run/openshift-sdn -type f -exec stat -c %a {} \;; done
```

Run the following commands to verify the permissions when using OVS:

```
for i in $(oc get pods -n openshift-sdn -l app=ovs -oname); do oc exec -n openshift-sdn $i -- find /var/run/openvswitch -type f -exec stat -c %a {} \;; done

for i in $(oc get pods -n openshift-sdn -l app=ovs -oname); do oc exec -n openshift-sdn $i -- find /etc/openvswitch -type f -exec stat -c %a {} \;; done

for i in $(oc get pods -n openshift-sdn -l app=ovs -oname); do oc exec -n openshift-sdn $i -- find /run/openvswitch -type f -exec stat -c %a {} \;; done
```

Verify the returned file permissions are **600** or more restrictive.

Remediation:

No remediation required; file permissions are managed by the operator.

Default Value:







In OpenShift 4, the default values are:

```
/host/etc/cni/net.d/00-multus.conf = 600
/host/var/run/multus/cni/net.d/80-openshift-network.conf = 644
/var/lib/cni/networks/openshift-sdn/* = 644 /var/run/openshift-sdn/cniserver/config.json = 444 /var/run/openvswitch/ovs-vswitchd.pid = 644 /etc/openvswitch/conf.db = 644 /etc/openvswitch/system-id.conf = 644 /etc/openvswitch/.conf.db~lock~ = 600 /run/openvswitch/ovs-vswitchd.pid = 644 /run/openvswitch/ovsdb-server.pid = 644
```

References:

1. <https://docs.openshift.com/container-platform/4.3/networking/cluster-network-operator.html>
2. <https://kubernetes.io/docs/concepts/cluster-administration/networking/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 <u>Configure Data Access Control Lists</u> Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 <u>Protect Information through Access Control Lists</u> Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.1.10 Ensure that the Container Network Interface file ownership is set to root:root (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the Container Network Interface files have ownership set to **root:root**.

Rationale:

Container Network Interface provides various networking options for overlay networking. You should consult their documentation and restrict their respective file permissions to maintain the integrity of those files. Those files should be owned by **root:root**.

Impact:

None.

Audit:

The Cluster Network Operator (CNO) deploys and manages the cluster network components on an OpenShift Container Platform cluster. This includes the Container Network Interface (CNI) default network provider plug-in selected for the cluster during installation. OpenShift Container Platform uses the Multus CNI plug-in to allow chaining of CNI plug-ins. The default Pod network must be configured during cluster installation. By default, the CNO deploys the OpenShift SDN as the default Pod network.

Ensure that the **multu** and **openshift-sdn** file ownership is set to root:root and the Open vSwitch (OVS) file ownership is set to **openvswitch:openvswitch**.

The SDN components are deployed as DaemonSets across the master/worker nodes with controllers limited to control plane nodes. OpenShift deploys OVS as a network overlay by default. Various configurations (ConfigMaps and other files managed by the operator via **hostpath** but stored on the container hosts) are stored in the following locations:

CNI:

/host/etc/cni/net.d /host/var/run/multus/cni/net.d

SDN:

/var/lib/cni/networks/openshift-sdn /var/run/openshift-sdn

SDN OVS:

/var/run/openvswitch /etc/openvswitch /run/openvswitch

Run the following commands to verify ownership when using CNI multus:

```
for i in $(oc get pods -n openshift-multus -l app=multus -oname); do oc exec -n openshift-multus $i -- /bin/bash -c "stat -c \"%U:%G %n\" /host/etc/cni/net.d/*.conf"; done

for i in $(oc get pods -n openshift-multus -l app=multus -oname); do oc exec -n openshift-multus $i -- /bin/bash -c "stat -c \"%U:%G %n\" /host/var/run/multus/cni/net.d/*.conf"; done
```

Run the following commands to verify the permissions when using OpenShift SDN:

```
for i in $(oc get pods -n openshift-sdn -l app=sdn -oname); do oc exec -n openshift-sdn $i -- find /var/lib/cni/networks/openshift-sdn -type f -exec stat -c \"%U:%G\" {} \;; done

for i in $(oc get pods -n openshift-sdn -l app=sdn -oname); do oc exec -n openshift-sdn $i -- find /var/run/openshift-sdn -type f -exec stat -c %U:%G {} \;; done
```

Run the following commands to verify the permissions when using OVS:

```
for i in $(oc get pods -n openshift-sdn -l app=ovs -oname); do oc exec -n openshift-sdn $i -- find /var/run/openvswitch -type f -exec stat -c %U:%G {} \;; done

for i in $(oc get pods -n openshift-sdn -l app=ovs -oname); do oc exec -n openshift-sdn $i -- find /etc/openvswitch -type f -exec stat -c %U:%G {} \;; done

for i in $(oc get pods -n openshift-sdn -l app=ovs -oname); do oc exec -n openshift-sdn $i -- find /run/openvswitch -type f -exec stat -c %U:%G {} \;; done
```

Verify file ownership is set to **root:root**.

In deployments using OCS, verify that the file ownership is set to **openvswitch:openvswitch**.

/var/run/openvswitch = openvswitch:openvswitch /etc/openvswitch = openvswitch:openvswitch /run/openvswitch = openvswitch:openvswitch

Remediation:

No remediation required; file permissions are managed by the operator.

Default Value:







In OpenShift 4, the default file ownership is root:root for CNI Multus and SDN and openvswitch:openvswitch for the OVS plugin.

/host/etc/cni/net.d/00-multus.conf = root:root
/host/var/run/multus/cni/net.d/80-openshift-network.conf = root:root
/var/lib/cni/networks/openshift-sdn = root:root /var/run/openshift-sdn = root:root
/var/run/openvswitch = openvswitch:openvswitch
/etc/openvswitch = openvswitch:openvswitch /run/openvswitch = openvswitch:openvswitch

References:

1. <https://docs.openshift.com/container-platform/latest/networking/cluster-network-operator.html>
2. <https://kubernetes.io/docs/concepts/cluster-administration/networking/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 <u>Restrict Administrator Privileges to Dedicated Administrator Accounts</u> Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 <u>Ensure the Use of Dedicated Administrative Accounts</u> Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

1.1.11 Ensure that the etcd data directory permissions are set to 700 or more restrictive (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the etcd data directory has permissions of **700** or more restrictive.

Rationale:

etcd is a highly-available key-value store used by Kubernetes deployments for persistent storage of all of its REST API objects. This data directory should be protected from any unauthorized reads or writes. It should not be readable or writable by any group members or the world.

Impact:

None

Audit:

In OpenShift 4, **etcd** members are deployed on the master nodes as static pods. The pod specification file is created on control plane nodes at `/etc/kubernetes/manifests/etcd-member.yaml`. The **etcd** database is stored on the container host in `/var/lib/etcd` and mounted to the **etcd-member** container via the host path mount `data-dir` with the same filesystem path (`/var/lib/etcd`). The permissions for this directory on the container host is **700**.

Starting with OCP 4.4, **etcd** is managed by the **cluster-etcd-operator**. The **etcd** operator will help to automate restoration of master nodes. There is also a new **etcdctl** container in the **etcd** static pod for quick debugging. cluster-admin rights are required to exec into **etcd** containers.

Run the following commands.

```
for i in $(oc get pods -n openshift-etcd -l app=etcd -oname); do oc exec -n openshift-etcd -c etcd $i -- stat -c %a%n /var/lib/etcd/member; done
```

Verify that the permissions are **700**.

Remediation:

No remediation required. File permissions are managed by the **etcd** operator.







Default Value:

By default, **etcd** data directory has permissions of **700**.

References:

1. https://docs.openshift.com/container-platform/4.3/architecture/control-plane.html#defining-masters_control-plane
2. <https://etcd.io/#data-dir>
3. <https://kubernetes.io/docs/tasks/administer-cluster/configure-upgrade-etcd/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.1.12 Ensure that the etcd data directory ownership is set to root:root (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the **etcd** data directory ownership is set to **root:root**.

Rationale:

etcd is a highly-available key-value store used by Kubernetes deployments for persistent storage of all of its REST API objects. This data directory should be protected from any unauthorized reads or writes. It should be owned by **root:root**.

In the OpenShift Container Platform, the ownership and permissions of the etcd data directory (usually located at `/var/lib/etcd`) are handled by the cluster-etcd-operator. This automation ensures the correct and secure configuration is maintained.

NOTE: The only users that exist on an RHCOS OpenShift node are root and core. This is intentional, as regular management of the underlying RHCOS cluster nodes is designed to be performed via the OpenShift API itself. The core user is a member of the wheel group, which gives it permission to use **sudo** for running privileged commands. Adding additional users at the node level is highly discouraged.

Impact:

None

Audit:

In OpenShift 4, **etcd** members are deployed on the master nodes as static pods. The **etcd** database is stored on the master nodes in `/var/lib/etcd` and mounted to the **etcd-member** container via the host path mount data-dir with the same filesystem path (`/var/lib/etcd`). The ownership for this directory on the **etcd-member** container and on the container host is **root:root**.

Starting with OCP 4.4, **etcd** is managed by the **cluster-etcd-operator**. The **etcd** operator will help to automate restoration of master nodes. There is also a new **etcdctl** container in the **etcd** static pod for quick debugging. cluster-admin rights are required to exec into **etcd** containers.

Run the following command.

```
for i in $(oc get pods -n openshift-etcd -l app=etcd -oname); do oc exec -n openshift-etcd -c etcd $i -- stat -c %U:%G /var/lib/etcd/member; done
```

Verify that the ownership is set to **root:root**.

Remediation:

No remediation required; file ownership is managed by the operator.







Default Value:

By default, in OpenShift 4, **etcd** data directory ownership is set to **root:root**.

References:

1. https://docs.openshift.com/container-platform/4.3/architecture/control-plane.html#defining-masters_control-plane
2. <https://etcd.io/#data-dir>
3. <https://kubernetes.io/docs/tasks/administer-cluster/configure-upgrade-etcd/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 Ensure the Use of Dedicated Administrative Accounts Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

1.1.13 Ensure that the kubeconfig file permissions are set to 600 or more restrictive (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the `kubeconfig` file has permissions of `600` or more restrictive.

Rationale:

The administrator `kubeconfig` file defines various settings for the administration of the cluster. You should restrict its file permissions to maintain the integrity of the file. The file should be writable by only the administrators on the system.

Impact:

None.

Audit:

In OpenShift 4 the admin config file is stored in `/etc/kubernetes/kubeconfig` with permissions `600`.

Run the following command:

```
for i in $(oc get nodes -o name)
do
    oc debug $i -- <<EOF
        chroot /host
        stat -c%a /etc/kubernetes/kubeconfig
    EOF
done
```

Verify that the permissions are `600` or more restrictive.

Remediation:

There is no remediation for updating the permissions of `kubeconfig`. The file is owned by an OpenShift operator and any changes to the file will result in a degraded cluster state.

Please do not attempt to remediate the permissions of this file.

Default Value:







By default, in OpenShift 4.14, the `kubeconfig` has permissions of `600`.

In older versions of OpenShift, the `kubeconfig` has permissions of `644`, and is not remediable. Please upgrade to OpenShift 4.14 when possible.

References:

1. https://docs.openshift.com/container-platform/4.5/cli_reference/openshift_cli/administrator-cli-commands.html
2. <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.1.14 Ensure that the kubeconfig file ownership is set to root:root (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the **kubeconfig** file ownership is set to **root:root**.

Rationale:

The **kubeconfig** file contains the admin credentials for the cluster. You should set its file ownership to maintain the integrity of the file. The file should be owned by **root:root**.

Impact:

None.

Audit:

In OpenShift 4 the admin config file is stored in **/etc/kubernetes/kubeconfig** with ownership **root:root**.

Run the following command.

```
for i in $(oc get nodes -o name)
do
  echo $i
  oc debug $i -- <<EOF
    chroot /host
    stat -c %U:%G /etc/kubernetes/kubeconfig
EOF
done
```

Verify that the ownership is set to **root:root**.

Remediation:

No remediation required; file permissions are managed by the operator.







Default Value:

By default, in OpenShift 4, **kubeconfig** file ownership is set to **root:root**.

References:

1. https://docs.openshift.com/container-platform/latest/cli_reference/openshift_cli/administrator-cli-commands.html
2. <https://kubernetes.io/docs/reference/setup-tools/kubeadm/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 <u>Restrict Administrator Privileges to Dedicated Administrator Accounts</u> Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 <u>Ensure the Use of Dedicated Administrative Accounts</u> Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

1.1.15 Ensure that the Scheduler kubeconfig file permissions are set to 600 or more restrictive (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the Scheduler **kubeconfig** file has permissions of **600** or more restrictive.

Rationale:

You should restrict the **kubeconfig** file permissions to maintain the integrity of the file. The file should be writable by only the administrators on the system.

Impact:

None.

Audit:

The **kubeconfig** file for **kube-scheduler** is stored in the ConfigMap **scheduler-kubeconfig** in the namespace **openshift-kube-scheduler**. The **kubeconfig** file is referenced in the pod via **hostpath** and is stored in **/etc/kubernetes/static-pod-resources/configmaps/scheduler-kubeconfig/kubeconfig** with permissions **600**.

Run the following command :

```
for i in $(oc get pods -n openshift-kube-scheduler -l app=openshift-kube-scheduler -o name)
do
  oc exec -n openshift-kube-scheduler $i -- \
  stat -c %a /etc/kubernetes/static-pod-resources/configmaps/scheduler-kubeconfig/kubeconfig
done
```

Verify that the permissions are **600** or more restrictive.

Remediation:

There is no remediation for updating the permissions of the **kubeconfig** file. The file is owned by an OpenShift operator and any changes to the file will result in a degraded cluster state.

Please do not attempt to remediate the permissions of this file.

Default Value:







By default, in OpenShift 4.14, the **kubeconfig** has permissions of **600**.

In older versions of OpenShift, the **kubeconfig** has permissions of **644**, and is not remediable. Please upgrade to OpenShift 4.14 when possible.

References:

1. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#cluster-kube-scheduler-operator_red-hat-operators
2. <https://docs.openshift.com/container-platform/latest/nodes/scheduling/nodes-scheduler-about.html>
3. <https://kubernetes.io/docs/concepts/scheduling-eviction/kube-scheduler/>
4. <https://issues.redhat.com//browse/OCBUGS-14323>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.1.16 Ensure that the Scheduler kubeconfig file ownership is set to root:root (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the **kubeconfig** file ownership is set to **root:root**.

Rationale:

You should set the **kubeconfig** file ownership to maintain the integrity of the file. The file should be owned by **root:root**.

Impact:

None.

Audit:

The **kubeconfig** file for **kube-scheduler** is stored in the ConfigMap **scheduler-kubeconfig** in the namespace **openshift-kube-scheduler**. The file **kubeconfig** is referenced in the pod via **hostpath** and is stored in **/etc/kubernetes/static-pod-resources/configmaps/scheduler-kubeconfig/kubeconfig** with ownership **root:root**.

Run the following command.

```
for i in $(oc get pods -n openshift-kube-scheduler -l app=openshift-kube-scheduler -o name)
do
  oc exec -n openshift-kube-scheduler $i -- \
    stat -c %U:%G /etc/kubernetes/static-pod-resources/configmaps/scheduler-kubeconfig/kubeconfig
done
```

Verify that the ownership is set to **root:root**.

Remediation:

No remediation required; file permissions are managed by the operator.

Default Value:







By default, **scheduler-kubeconfig/kubeconfig** file ownership is set to **root:root**.

References:

1. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#cluster-kube-scheduler-operator_red-hat-operators

2. <https://docs.openshift.com/container-platform/latest/nodes/scheduling/nodes-scheduler-about.html>
3. <https://kubernetes.io/docs/concepts/scheduling-eviction/kube-scheduler/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	<u>5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts</u> Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	<u>4.3 Ensure the Use of Dedicated Administrative Accounts</u> Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

1.1.17 Ensure that the Controller Manager kubeconfig file permissions are set to 600 or more restrictive (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the **kubeconfig** file mounted into the Controller Manager has permissions of **600** or more restrictive.

Rationale:

You should restrict the **kubeconfig** file permissions to maintain the integrity of the file. The file should be writable by only the administrators on the system.

Impact:

None.

Audit:

The **kubeconfig** file for **kube-controller-manager** is stored in the ConfigMap **controller-manager-kubeconfig** in the namespace **openshift-kube-controller-manager**. The **kubeconfig** file is referenced in the pod via **hostpath** and is stored in **/etc/kubernetes/static-pod-resources/configmaps/controller-manager-kubeconfig/kubeconfig** with permissions **600**.

Run the following command.

```
for i in $(oc get pods -n openshift-kube-controller-manager -l app=kube-controller-manager -oname)
do
  oc exec -n openshift-kube-controller-manager $i -- \
  stat -c %a /etc/kubernetes/static-pod-resources/configmaps/controller-manager-kubeconfig/kubeconfig
done
```

Verify that the permissions are **600** or more restrictive.

Remediation:

There is no remediation for updating the permissions of the **kubeconfig** file. The file is owned by an OpenShift operator and any changes to the file will result in a degraded cluster state.

Please do not attempt to remediate the permissions of this file.

Default Value:







By default, in OpenShift 4.14, the **kubeconfig** has permissions of **600**.

In older versions of OpenShift, the **kubeconfig** has permissions of **644**, and is not remediable. Please upgrade to OpenShift 4.14 when possible.

References:

1. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#kube-controller-manager-operator_red-hat-operators
2. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#cluster-openshift-controller-manager-operator_red-hat-operators
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-controller-manager/>
4. <https://issues.redhat.com//browse/OCPBUGS-14323>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.1.18 Ensure that the Controller Manager kubeconfig file ownership is set to root:root (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the **kubeconfig** file ownership is set to **root:root**.

Rationale:

You should set the **kubeconfig** file ownership to maintain the integrity of the file. The file should be owned by **root:root**.

Impact:

None.

Audit:

Run the following command:

```
for i in $(oc get pods -n openshift-kube-controller-manager -l app=kube-
controller-manager -oname)
do
  oc exec -n openshift-kube-controller-manager $i -- \
    stat -c %U:%G /etc/kubernetes/static-pod-resources/configmaps/controller-
manager-kubeconfig/kubeconfig
done
```

Verify that the ownership is set to **root:root**.

Remediation:

No remediation required; file permissions are managed by the operator.







Default Value:

By default, the Controller Manager **kubeconfig** file ownership is set to **root:root**.

References:

1. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#kube-controller-manager-operator_red-hat-operators
2. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#cluster-openshift-controller-manager-operator_red-hat-operators
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-controller-manager/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 <u>Restrict Administrator Privileges to Dedicated Administrator Accounts</u> Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 <u>Ensure the Use of Dedicated Administrative Accounts</u> Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

1.1.19 Ensure that the OpenShift PKI directory and file ownership is set to root:root (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the OpenShift PKI directory and file ownership is set to **root:root**.

Rationale:

OpenShift makes use of a number of certificates as part of its operation. You should verify the ownership of the directory containing the PKI information and all files in that directory to maintain their integrity. The directory and files should be owned by **root:root**.

Impact:

None

Audit:

Keys for control plane components deployed as static pods, **kube-apiserver**, **kube-controller-manager**, and **openshift-kube-scheduler** are stored in the directory **/etc/kubernetes/static-pod-certs/secrets**. The directory and file ownership are set to **root:root**.

Run the following command.

```
# Should return root:root for all files and directories

for i in $(oc -n openshift-kube-apiserver get pod -l app=openshift-kube-apiserver -o jsonpath='{.items[*].metadata.name}');
do
    echo $i static-pod-certs
    oc exec -n openshift-kube-apiserver $i -c kube-apiserver -- \
    find /etc/kubernetes/static-pod-certs -type d -wholename '*/secrets*'
-exec stat -c %U:%G {} \;
    oc exec -n openshift-kube-apiserver $i -c kube-apiserver -- \
    find /etc/kubernetes/static-pod-certs -type f -wholename '*/secrets*'
-exec stat -c %U:%G {} \;
    echo $i static-pod-resources
    oc exec -n openshift-kube-apiserver $i -c kube-apiserver -- \
    find /etc/kubernetes/static-pod-resources -type d -wholename
'*/secrets*' -exec stat -c %U:%G {} \;
    oc exec -n openshift-kube-apiserver $i -c kube-apiserver -- \
    find /etc/kubernetes/static-pod-resources -type f -wholename
'*/secrets*' -exec stat -c %U:%G {} \;
done
```

Verify that the ownership of all files and directories in this hierarchy is set to **root:root**.

Remediation:

No remediation required; file permissions are managed by the operator.







Default Value:

By default, the `static-pod-resources/secrets` and `static-pod-certs` directories and all of the files and directories contained within it, are set to be owned by the root user.

References:

1. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 Ensure the Use of Dedicated Administrative Accounts Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.1.20 Ensure that the OpenShift PKI certificate file permissions are set to 600 or more restrictive (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that OpenShift PKI certificate files have permissions of **644** or more restrictive.

Rationale:

OpenShift makes use of a number of certificate files as part of the operation of its components. The permissions on these files should be set to 644 or more restrictive to protect their integrity.

Impact:

None

Audit:

Certificates for control plane components like **kube-apiserver**, **kube-controller-manager**, and **kube-scheduler** are stored in the directory **/etc/kubernetes/static-pod-certs/secrets**. Certificate files all have permissions **600**.

Run the following command.

```
# Should 600 or more restrictive

for i in $(oc -n openshift-kube-apiserver get pod -l app=openshift-kube-apiserver -o jsonpath='{.items[*].metadata.name}');
do
    echo $i static-pod-certs
    oc exec -n openshift-kube-apiserver $i -c kube-apiserver -- \
    find /etc/kubernetes/static-pod-certs -type f -wholename
    '*/secrets/*.crt' -exec stat -c %a {} \;
done
```

Verify that the permissions are **600**.

Remediation:

No remediation required; file permissions are managed by the operator.







Default Value:

By default, the certificates used by OpenShift are set to have permissions of **600**.

References:

1. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 <u>Configure Data Access Control Lists</u> Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 <u>Protect Information through Access Control Lists</u> Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.1.21 Ensure that the OpenShift PKI key file permissions are set to 600 (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the OpenShift PKI key files have permissions of **600**.

Rationale:

OpenShift makes use of a number of key files as part of the operation of its components. The permissions on these files should be set to **600** to protect their integrity and confidentiality.

Impact:

None

Audit:

Keys for control plane components like **kube-apiserver**, **kube-controller-manager**, **ube-scheduler** and **etcd** are stored with their respective static pod configurations in the directory **/etc/kubernetes/static-pod-certs/secrets**. Key files all have permissions **600**.

Run the following command.

```
for i in $(oc -n openshift-kube-apiserver get pod -l app=openshift-kube-apiserver -o jsonpath='{.items[*].metadata.name}');
do
    echo $i static-pod-certs
    oc exec -n openshift-kube-apiserver $i -c kube-apiserver -- \
    find /etc/kubernetes/static-pod-certs -type f -wholename
    '*/secrets/*.key' -exec stat -c %a {} \;
done
```

Verify that the permissions are **600**.

Remediation:

No remediation required; file permissions are managed by the operator.







Default Value:

By default, the keys used by OpenShift are set to have permissions of **600**

References:

1. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 <u>Configure Data Access Control Lists</u> Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 <u>Protect Information through Access Control Lists</u> Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

1.2 API Server

This section contains recommendations relating to API server configuration flags.

OpenShift includes two API servers, the OpenShift API server and the Kubernetes API server. All API calls are directed to the Open Shift API server and then Kubernetes objects are delegated to **kube-apiserver**.

The cluster configuration resource (CR) APIServer holds configuration settings (like serving certificates, client CA and CORS domains) shared by all API servers in the system, among them especially **kube-apiserver** and **openshift-apiserver**. The canonical name of an instance is 'cluster'. Changes to the API server configurations should be done in the APIServer custom resource definition (CRD):
apiservers.config.openshift.io

The OpenShift API Server is managed by the **openshift-apiserver-operator**

The Kubernetes API Server is managed by the **openshift-kube-apiserver-operator**

Both are managed by the Cluster Version Operator

1.2.1 Ensure that anonymous requests are authorized (Manual)

Profile Applicability:

- Level 1

Description:

When anonymous requests to the API server are allowed, they must be authorized.

Rationale:

When enabled, requests that are not rejected by other configured authentication methods are treated as anonymous requests. These requests are then served by the API server. You should rely on authentication to authorize anonymous requests.

If you are using RBAC authorization, it is generally considered reasonable to allow anonymous access to the API Server for health checks and discovery purposes, and hence this recommendation is not scored. However, you should consider whether anonymous discovery is an acceptable risk for your purposes.

Impact:

Anonymous requests are assigned to the **system:unauthenticated** group which allows the system to determine which actions are allowed.

Audit:

OpenShift allows anonymous requests (then authorizes them). OpenShift allows anonymous requests to the API server to support information discovery and **webhook** integrations. OpenShift provides its own fully integrated authentication and authorization mechanism. If no access token or certificate is presented, the authentication layer assigns the **system:anonymous** virtual user and the **system:unauthenticated** virtual group to the request. This allows the authorization layer to determine which requests, if any, an anonymous user is allowed to make.

```
oc get clusterrolebindings -o json | jq '.items[] | select(.subjects[]?.kind == "Group" and .subjects[]?.name == "system:unauthenticated") | .metadata.name' | uniq
```

Returns what unauthenticated users can do, which is the following:

```
"self-access-reviewers"  
"system:oauth-token-deleters"  
"system:openshift:public-info-viewer"  
"system:public-info-viewer"  
"system:scope-impersonation"  
"system:webhooks"
```

Remediation:

None. The default configuration should not be modified.

Default Value:







By default, anonymous access is enabled and assigned to the **system:unauthenticated** group, which allows the system to determine which actions are allowed.

If the default behavior is changed, platform components will not work properly, in particular Elasticsearch and Prometheus. The **oauth-proxy** deployed as part of these components makes anonymous use of **/.well-known/oauth-authorization-server** endpoint, granted by **system:discovery** role.

References:

1. <https://docs.openshift.com/container-platform/4.5/authentication/understanding-authentication.html>
2. <https://docs.openshift.com/container-platform/4.5/authentication/using-rbac.html>
3. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#cluster-authentication-operator_red-hat-operators
4. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
5. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators
6. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
7. <https://kubernetes.io/docs/reference/access-authn-authz/authentication/#anonymous-requests>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1190, T1210	TA0001, TA0008	M1025

1.2.2 Use https for kubelet connections (Manual)

Profile Applicability:

- Level 1

Description:

Use https for kubelet connections.

Rationale:

Connections from **apiserver** to **kubelets** could potentially carry sensitive data such as secrets and keys. It is thus important to use in-transit encryption for any communication between the **apiserver** and **kubelets**.

Impact:

You require TLS to be configured on **apiserver** as well as **kubelets**.

Audit:

OpenShift does not use the **--kubelet-https** argument. OpenShift utilizes X.509 certificates for authentication of the control-plane components. OpenShift configures the API server to use an internal certificate authority (CA) to validate the user certificate sent during TLS negotiation. If the validation of the certificate is successful, the request is authenticated and user information is derived from the certificate subject fields.

To verify the kubelet client certificates are present, run the following command:

```
oc get configmap config -n openshift-kube-apiserver -ojson | jq -r
'.data["config.yaml"]' | jq '.apiServerArguments["kubelet-client-
certificate"]'

oc get configmap config -n openshift-kube-apiserver -ojson | jq -r
'.data["config.yaml"]' | jq '.apiServerArguments["kubelet-client-key"]'

oc -n openshift-apiserver describe secret serving-cert

# Run the following command and the output should return true or no output at
all

oc get configmap config -n openshift-kube-apiserver -ojson | jq -r
'.data["config.yaml"]' | jq '.apiServerArguments["kubelet-https"]'
```

Verify that the kubelet client-certificate and kubelet client-key files are present.

```
client-certificate:

'/etc/kubernetes/static-pod-resources/kube-apiserver-certs/secrets/kubelet-
client/tls.crt'

client-key:

'/etc/kubernetes/static-pod-resources/kube-apiserver-certs/secrets/kubelet-
client/tls.key'
```

Verify that the serving-cert for the **openshift-apiserver** is type **kubernetes.io/tls** and that returned Data includes **tls.crt** and **tls.key**.

Remediation:

No remediation is required. OpenShift platform components use X.509 certificates for authentication. OpenShift manages the CAs and certificates for platform components. This is not configurable.





Default Value:

By default, kubelet connections are encrypted.

References:

1. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
2. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators
3. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.13/bindata/assets/config/defaultconfig.yaml#L124-L127>
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet-authentication-authorization/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.10 <u>Encrypt Sensitive Data in Transit</u> Encrypt sensitive data in transit. Example implementations can include: Transport Layer Security (TLS) and Open Secure Shell (OpenSSH).			
v7	14.4 <u>Encrypt All Sensitive Information in Transit</u> Encrypt all sensitive information in transit.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1048, T1189	TA0001, TA0010	M1041

1.2.3 Ensure that the kubelet uses certificates to authenticate (Manual)

Profile Applicability:

- Level 1

Description:

Enable certificate based kubelet authentication.

Rationale:

The apiserver, by default, does not authenticate itself to the kubelet's HTTPS endpoints. The requests from the apiserver are treated anonymously. You should set up certificate-based kubelet authentication to ensure that the apiserver authenticates itself to kubelets when submitting requests.

Impact:

Require TLS to be configured on the apiserver as well as kubelets.

Audit:

OpenShift does not use the `--kubelet-client-certificate` or the `kubelet-client-key` arguments. OpenShift utilizes X.509 certificates for authentication of the control-plane components. OpenShift configures the API server to use an internal certificate authority (CA) to validate the user certificate sent during TLS negotiation. If the CA validation of the certificate is successful, the request is authenticated and user information is derived from the certificate subject fields.

To verify the certificates are present, run the following command:

```
#for OpenShift 4.6 and above

oc get configmap config -n openshift-kube-apiserver -ojson | jq -r
'.data["config.yaml"]' | jq '.apiServerArguments["kubelet-client-
certificate"]'

oc get configmap config -n openshift-kube-apiserver -ojson | jq -r
'.data["config.yaml"]' | jq '.apiServerArguments["kubelet-client-key"]'

oc -n openshift-apiserver describe secret serving-cert
```

Verify that the kubelet client-certificate and kubelet client-key files are present.

client-certificate:

`/etc/kubernetes/static-pod-certs/secrets/kubelet-client/tls.crt`

client-key:

`/etc/kubernetes/static-pod-certs/secrets/kubelet-client/tls.key`

Verify that the serving-cert for the **openshift-apiserver** is type **kubernetes.io/tls** and that returned Data includes **tls.crt** and **tls.key**.

Remediation:

No remediation is required. OpenShift platform components use X.509 certificates for authentication. OpenShift manages the CAs and certificates for platform components. This is not configurable.




Default Value:

By default, kubelet authentication is managed with X.509 certificates.

References:

1. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.13/bindata/assets/config/defaultconfig.yaml#L124-L127>
2. <https://kubernetes.io/docs/admin/kube-apiserver/>
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet-authentication-authorization/>
4. <https://kubernetes.io/docs/concepts/architecture/control-plane-node-communication/>
5. https://docs.openshift.com/container-platform/4.13/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
6. https://docs.openshift.com/container-platform/4.13/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	16.11 Leverage Vetted Modules or Services for Application Security Components Leverage vetted modules or services for application security components, such as identity management, encryption, and auditing and logging. Using platform features in critical security functions will reduce developers' workload and minimize the likelihood of design or implementation errors. Modern operating systems provide effective mechanisms for identification, authentication, and authorization and make those mechanisms available to applications. Use only standardized, currently accepted, and extensively reviewed encryption algorithms. Operating systems also provide mechanisms to create and maintain secure audit logs.			
v7	1.8 Utilize Client Certificates to Authenticate Hardware Assets Use client certificates to authenticate hardware assets connecting to the organization's trusted network.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1048, T1189	TA0001, TA0010	M1041

1.2.4 Verify that the kubelet certificate authority is set as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

Verify kubelet's certificate before establishing connection.

Rationale:

The connections from the apiserver to the kubelet are used for fetching logs for pods, attaching (through kubectl) to running pods, and using the kubelet's port-forwarding functionality. These connections terminate at the kubelet's HTTPS endpoint. By default, the apiserver does not verify the kubelet's serving certificate, which makes the connection subject to man-in-the-middle attacks, and unsafe to run over untrusted and/or public networks.

Impact:

You require TLS to be configured on apiserver as well as kubelets.

Audit:

OpenShift does not use the `--kubelet-certificate-authority` flag. OpenShift utilizes X.509 certificates for authentication of the control-plane components. OpenShift configures the API server to use an internal certificate authority (CA) to validate the user certificate sent during TLS negotiation. If the CA validation of the certificate is successful, the request is authenticated and user information is derived from the certificate subject fields.

To verify, run the following command:


```
# For 4.6 and above

oc get configmap config -n openshift-kube-apiserver -ojson | jq -r
'.data["config.yaml"]' | jq '.apiServerArguments["kubelet-certificate-
authority"]'

# For 4.11 and above
$ oc get configmap config -n openshift-kube-apiserver -o json | jq -r
'.data["config.yaml"]' | jq '.apiServerArguments["kubelet-client-
certificate"]'
[
  "/etc/kubernetes/static-pod-certs/secrets/kubelet-client/tls.crt"
]
$ oc get configmap config -n openshift-kube-apiserver -o json | jq -r
'.data["config.yaml"]' | jq '.apiServerArguments["kubelet-client-key"]'
[
  "/etc/kubernetes/static-pod-certs/secrets/kubelet-client/tls.key"
]
```

Remediation:

No remediation is required. OpenShift platform components use X.509 certificates for authentication. OpenShift manages the CAs and certificates for platform components. This is not configurable.

Default Value:

By default, kubelet authentication is managed with X.509 certificates.

References:

1. <https://docs.openshift.com/container-platform/4.13/operators/operator-reference.html>
2. <https://docs.openshift.com/container-platform/4.13/operators/operator-reference.html>
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet-authentication-authorization/>
5. <https://kubernetes.io/docs/concepts/architecture/control-plane-node-communication/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	<u>16.11 Leverage Vetted Modules or Services for Application Security Components</u> Leverage vetted modules or services for application security components, such as identity management, encryption, and auditing and logging. Using platform features in critical security functions will reduce developers' workload and minimize the likelihood of design or implementation errors. Modern operating systems provide effective mechanisms for identification, authentication, and authorization and make those mechanisms available to applications. Use only standardized, currently accepted, and extensively reviewed encryption algorithms. Operating systems also provide mechanisms to create and maintain secure audit logs.		●	●
v7	<u>1.8 Utilize Client Certificates to Authenticate Hardware Assets</u> Use client certificates to authenticate hardware assets connecting to the organization's trusted network.			●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1048, T1189	TA0001, TA0010	M1041

1.2.5 Ensure that the `--authorization-mode` argument is not set to `AlwaysAllow` (Manual)

Profile Applicability:

- Level 1

Description:

Do not always authorize all requests.

Rationale:

The API Server, can be configured to allow all requests. This mode should not be used on any production cluster.

Impact:

Only authorized requests will be served.

Audit:

It is not possible to configure an OpenShift cluster to allow all requests. OpenShift is configured at bootstrap time to use RBAC to authorize requests. Role-based access control (RBAC) objects determine what actions a user is allowed to perform on what objects in an OpenShift cluster. Cluster administrators manage RBAC for the cluster. Project owners can manage RBAC for their individual OpenShift projects. The OpenShift API server configmap does not use the `authorization-mode` flag.

To verify, run the following commands:

```
$ oc get configmap config -n openshift-kube-apiserver -ojson | jq -r  
' .data["config.yaml"]' | jq '.apiServerArguments."authorization-mode"'  
[  
  "Scope",  
  "SystemMasters",  
  "RBAC",  
  "Node"  
]` ``
```

Remediation:

None. RBAC is always on and the OpenShift API server does not use the values assigned to the flag `authorization-mode`.






Default Value:

OpenShift uses RBAC by default.

References:

1. <https://docs.openshift.com/container-platform/4.5/authentication/using-rbac.html>
2. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
3. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators
4. <https://kubernetes.io/docs/admin/kube-apiserver/>
5. <https://kubernetes.io/docs/reference/access-authn-authz/authorization/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	9.2 Ensure Only Approved Ports, Protocols and Services Are Running Ensure that only network ports, protocols, and services listening on a system with validated business needs, are running on each system.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1133	TA0001	M1026

1.2.6 Verify that RBAC is enabled (Manual)

Profile Applicability:

- Level 1

Description:

Turn on Role Based Access Control.

Rationale:

Role Based Access Control (RBAC) allows fine-grained control over the operations that different entities can perform on different objects in the cluster. It is recommended to use the RBAC authorization mode.

Impact:

When RBAC is enabled you will need to ensure that appropriate RBAC settings (including Roles, RoleBindings, ClusterRoles, and ClusterRoleBindings) are configured to allow appropriate access.

Audit:

OpenShift is configured at bootstrap time to use Role-Based Access Control (RBAC) to authorize requests. RBAC objects determine what actions a user is allowed to perform on what objects in an OpenShift cluster. Cluster administrators manage RBAC for the cluster. Project owners can manage RBAC for their individual OpenShift projects.

Use the following command to verify the API server is configured to use RBAC for authorization:

```
$ oc get configmap config -n openshift-kube-apiserver -ojson | jq -r '.data["config.yaml"]' | jq '.apiServerArguments."authorization-mode"'
```

Verify the list returned contains **RBAC** as an authorization option.

Remediation:

None.

Default Value:

OpenShift uses RBAC by default. OpenShift includes default roles and role bindings. Custom roles and role bindings can be added for additional granularity.

Please refer to the OpenShift [documentation](#) for more information on RBAC.

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/index.html>

2. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.5/bindata/bootkube/manifests/cluster-role-binding-kube-apiserver.yaml>
3. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.6/bindata/v4.1.0/config/defaultconfig.yaml#L17-L21>
4. <https://kubernetes.io/docs/reference/access-authn-authz/rbac/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	6.8 Define and Maintain Role-Based Access Control Define and maintain role-based access control, through determining and documenting the access rights necessary for each role within the enterprise to successfully carry out its assigned duties. Perform access control reviews of enterprise assets to validate that all privileges are authorized, on a recurring schedule at a minimum annually, or more frequently.			●
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.	●	●	●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078, T1548	TA0001, TA0004	M1018

1.2.7 Ensure that the *APIPriorityAndFairness* feature gate is enabled (Manual)

Profile Applicability:

- Level 1

Description:

Limit the rate at which the API server accepts requests.

Rationale:

A misbehaving workload could overwhelm and DoS the API Server, making it unavailable. This particularly applies to a multi-tenant cluster, where there might be a small percentage of misbehaving tenants which could have a significant impact on the performance of the cluster overall. Hence, it is recommended to limit the rate of events that the API server will accept.

Impact:

None, as the OpenShift kubelet has been fixed to send fewer requests.

Audit:

OpenShift 4.5 and forward uses the api priority and fairness feature to limit the rate at which the API server accepts requests.

Run the following command:

```
#Verify the APIPriorityAndFairness feature-gate
oc get kubeapiservers.operator.openshift.io cluster -o json | jq
'.spec.observedConfig.apiServerArguments'
```

For 4.5, verify that the feature-gate is turned on for the APIServer priority and fairness: **APIPriorityAndFairness=true**. In OCP 4.5 and earlier, the default set of admission plugins are compiled into the **apiserver** and are not visible in the configuration yaml.

Remediation:

No remediation is required.





Default Value:

By default, the OpenShift kubelet has been fixed to send fewer requests. Version 4.6+ it is enabled by default.

References:

1. <https://docs.openshift.com/container-platform/4.13/operators/operator-reference.html>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	12.6 <u>Use of Secure Network Management and Communication Protocols</u> Use secure network management and communication protocols (e.g., 802.1X, Wi-Fi Protected Access 2 (WPA2) Enterprise or greater).			
v7	8.3 <u>Enable Operating System Anti-Exploitation Features/ Deploy Anti-Exploit Technologies</u> Enable anti-exploitation features such as Data Execution Prevention (DEP) or Address Space Layout Randomization (ASLR) that are available in an operating system or deploy appropriate toolkits that can be configured to apply protection to a broader set of applications and executables.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1609	TA0002	M1028

1.2.8 Ensure that the admission control plugin AlwaysAdmit is not set (Manual)

Profile Applicability:

- Level 1

Description:

Do not allow all requests.

Rationale:

Setting admission control plugin **AlwaysAdmit** allows all requests and does not filter any requests.

The **AlwaysAdmit** admission controller was deprecated in Kubernetes v1.13. Its behavior was equivalent to turning off all admission controllers.

Impact:

Only requests explicitly allowed by the admissions control plugins would be served.

Audit:

This controller is disabled by default in OpenShift and cannot be enabled. It has also been deprecated by the Kubernetes community as it behaves as if there were no controller.

Use the following command to verify the configured admission controllers:

```
oc -n openshift-kube-apiserver get configmap config -o json | jq -r '.data."config.yaml"' | jq '.apiServerArguments."enable-admission-plugins"'
```

The output should not include **AlwaysAdmit**.

Remediation:

None.

Default Value:







This **AlwaysAdmit** controller is disabled by default in OpenShift and cannot be enabled.

References:

1. <https://docs.openshift.com/container-platform/latest/architecture/admission-plugins.html>
2. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.6/bindata/v4.1.0/config/defaultconfig.yaml#L34-L78>

3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
4. <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/#alwaysadmit>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 <u>Configure Data Access Control Lists</u> Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 <u>Protect Information through Access Control Lists</u> Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1133	TA0001	M1026

1.2.9 Ensure that the admission control plugin AlwaysPullImages is not set (Manual)

Profile Applicability:

- Level 1

Description:

Always pull images.

Rationale:

Setting admission control policy to **AlwaysPullImages** forces every new pod to pull the required images every time. In a multi-tenant cluster users can be assured that their private images can only be used by those who have the credentials to pull them. Without this admission control policy, once an image has been pulled to a node, any pod from any user can use it simply by knowing the image's name, without any authorization check against the image ownership. When this plug-in is enabled, images are always pulled prior to starting containers, which means valid credentials are required.

However, turning on this admission plugin can introduce new kinds of cluster failure modes. OpenShift 4 master and infrastructure components are deployed as pods. Enabling this feature can result in cases where loss of contact to an image registry can cause a redeployed infrastructure pod (**oauth-server** for example) to fail on an image pull for an image that is currently present on the node. We use **PullIfNotPresent** so that a loss of image registry access does not prevent the pod from starting. If it becomes **PullAlways**, then an image registry access outage can cause key infrastructure components to fail.

This can be managed per container. When OpenShift Container Platform creates containers, it uses the container's **imagePullPolicy** to determine if the image should be pulled prior to starting the container. There are three possible values for **imagePullPolicy**: **Always**, **IfNotPresent**, **Never**. If a container's **imagePullPolicy** parameter is not specified, OpenShift Container Platform sets it based on the image's tag. If the tag is latest, OpenShift Container Platform defaults **imagePullPolicy** to **Always**. Otherwise, OpenShift Container Platform defaults **imagePullPolicy** to **IfNotPresent**.

Impact:

Credentials would be required to pull the private images every time. Also, in trusted environments, this might increase load on network, registry, and decrease speed.

This setting could impact offline or isolated clusters, which have images pre-loaded and do not have access to a registry to pull in-use images. This setting is not appropriate for clusters which use this configuration.

Audit:

Use the following command to obtain a list of configured admission controllers:

```
oc -n openshift-kube-apiserver get configmap config -o json | jq -r '.data."config.yaml"' | jq '.apiServerArguments."enable-admission-plugins"'
```

Verify the list does not include **AlwaysPullImages**.

Remediation:

None.







Default Value:

When OpenShift Container Platform creates containers, it uses the container's **imagePullPolicy** to determine if the image should be pulled prior to starting the container.

References:

1. https://docs.openshift.com/container-platform/latest/openshift_images/managing_images/image-pull-policy.html
2. <https://docs.openshift.com/container-platform/latest/architecture/admission-plugins.html>
3. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.6/bindata/v4.1.0/config/defaultconfig.yaml#L34-L78>
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
5. <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/#alwayspullimages>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1610	TA0002	M1038, M1050

1.2.10 Ensure that the admission control plugin ServiceAccount is set (Manual)

Profile Applicability:

- Level 1

Description:

Automate service accounts management.

Rationale:

When you create a pod, if you do not specify a service account, it is automatically assigned the default service account in the same namespace. You should create your own service account and let the API server manage its security tokens.

Impact:

None.

Audit:

The **ServiceAccount** admission control plugin is enabled by default. Every service account has an associated user name that can be granted roles, just like a regular user. The user name for each service account is derived from its project and the name of the service account. Service accounts are required in each project to run builds, deployments, and other pods. The default service accounts that are automatically created for each project are isolated by the project namespace.

Use the following command to obtain a list of configured admission controllers:

```
oc -n openshift-kube-apiserver get configmap config -o json | jq -r  
'data.config.yaml' | jq '.apiServerArguments.enable-admission-plugins'
```

Verify the list includes **ServiceAccount**.

Remediation:

None.

Default Value:







By default, OpenShift configures the **ServiceAccount** admission controller.

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/understanding-and-creating-service-accounts.html>
2. <https://docs.openshift.com/container-platform/latest/architecture/admission-plugins.html>

3. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.6/bindata/v4.1.0/config/defaultconfig.yaml#L34-L78>
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
5. <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/#serviceaccount>
6. <https://kubernetes.io/docs/tasks/configure-pod-container/configure-service-account/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078.001, T1552	TA0001, TA0006	M1043

1.2.11 Ensure that the admission control plugin NamespaceLifecycle is set (Manual)

Profile Applicability:

- Level 1

Description:

Reject creating objects in a namespace that is undergoing termination.

Rationale:

Setting admission control policy to **NamespaceLifecycle** ensures that objects cannot be created in non-existent namespaces, and that namespaces undergoing termination are not used for creating the new objects. This is recommended to enforce the integrity of the namespace termination process and also for the availability of the newer objects.

Impact:

None.

Audit:

OpenShift enables the **NamespaceLifecycle** plugin by default.

Use the following command to obtain a list of configured admission controllers:

```
oc -n openshift-kube-apiserver get configmap config -o json | jq -r '.data."config.yaml"' | jq '.apiServerArguments."enable-admission-plugins"'
```

Verify that **NamespaceLifecycle** is present in the returned list.

Remediation:

None.







Default Value:

OpenShift configures **NamespaceLifecycle** admission controller by default.

References:

1. <https://docs.openshift.com/container-platform/latest/architecture/admission-plugins.html>
2. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.6/bindata/v4.1.0/config/defaultconfig.yaml#L34-L78>
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
4. <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/#namespace-lifecycle>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.1 <u>Establish and Maintain a Secure Configuration Process</u> Establish and maintain a secure configuration process for enterprise assets (end-user devices, including portable and mobile, non-computing/IoT devices, and servers) and software (operating systems and applications). Review and update documentation annually, or when significant enterprise changes occur that could impact this Safeguard.			
v7	14.6 <u>Protect Information through Access Control Lists</u> Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1578	TA0005	M1018

1.2.12 Ensure that the admission control plugin SecurityContextConstraint is set (Manual)

Profile Applicability:

- Level 1

Description:

Reject creating pods that do not match Pod Security Policies.

Rationale:

A Pod Security Policy is a cluster-level resource that controls the actions that a pod can perform and what it has the ability to access. The **PodSecurityPolicy** objects define a set of conditions that a pod must run with in order to be accepted into the system. Pod Security Policies are composed of settings and strategies that control the security features a pod has access to and hence this must be used to control pod access permissions.

Note: When the **PodSecurityPolicy** admission plugin is in use, there needs to be at least one **PodSecurityPolicy** in place for ANY pods to be admitted. See section 5.2 for recommendations on **PodSecurityPolicy** settings.

Impact:

Default Security Context Constraint objects are present on the cluster and granted by default based on roles. Custom SCCs can be created and granted as needed.

Audit:

OpenShift provides the same protection via the SecurityContextConstraints admission plugin, which is enabled by default. The **PodSecurityPolicy** admission control plugin is disabled by default as it is still beta and not yet supported with OpenShift. Security Context Constraints (SCCs) and Pod Security Policy cannot be used on the same cluster.

Use the following command to obtain a list of configured admission controllers:

```
oc -n openshift-kube-apiserver get configmap config -o json | jq -r '.data."config.yaml"' | jq '.apiServerArguments."enable-admission-plugins"'
```

Verify that **security.openshift.io/SecurityContextConstraint** is returned in the list.

Remediation:

None.






Default Value:

By default, the **SecurityContextConstraints** admission controller is configured and cannot be disabled.

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/managing-security-context-constraints.html>
2. <https://docs.openshift.com/container-platform/latest/architecture/admission-plugins.html>
3. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.6/bindata/v4.1.0/config/defaultconfig.yaml#L34-L78>
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
5. <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/#podsecuritypolicy>
6. <https://kubernetes.io/docs/concepts/policy/pod-security-policy/#enabling-pod-security-policies>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	16.11 Leverage Vetted Modules or Services for Application Security Components Leverage vetted modules or services for application security components, such as identity management, encryption, and auditing and logging. Using platform features in critical security functions will reduce developers' workload and minimize the likelihood of design or implementation errors. Modern operating systems provide effective mechanisms for identification, authentication, and authorization and make those mechanisms available to applications. Use only standardized, currently accepted, and extensively reviewed encryption algorithms. Operating systems also provide mechanisms to create and maintain secure audit logs.			
v7	9.4 Apply Host-based Firewalls or Port Filtering Apply host-based firewalls or port filtering tools on end systems, with a default-deny rule that drops all traffic except those services and ports that are explicitly allowed.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1610	TA0002	M1038, M1050

1.2.13 Ensure that the admission control plugin NodeRestriction is set (Manual)

Profile Applicability:

- Level 1

Description:

Limit the **Node** and **Pod** objects that a kubelet could modify.

Rationale:

Using the **NodeRestriction** plug-in ensures that the kubelet is restricted to the **Node** and **Pod** objects that it could modify as defined. Such kubelets will only be allowed to modify their own **Node** API object, and only modify **Pod** API objects that are bound to their node.

Impact:

None.

Audit:

In OpenShift, the **NodeRestriction** admission plugin is enabled by default.

Use the following command to obtain a list of configured admission controllers:

```
oc -n openshift-kube-apiserver get configmap config -o json | jq -r  
'data.config.yaml' | jq '.apiServerArguments.enable-admission-plugins'
```

Verify the list includes **NodeRestriction**.

Remediation:

None.

Default Value:

In OpenShift, the **NodeRestriction** admission plugin is enabled by default and cannot be disabled.

References:

1. <https://docs.openshift.com/container-platform/latest/architecture/admission-plugins.html>
2. <https://github.com/openshift/origin/blob/release-4.5/vendor/k8s.io/kubernetes/cmd/kubeadm/app/phases/controlplane/manifests.go#L132>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	16.11 <u>Leverage Vetted Modules or Services for Application Security Components</u> Leverage vetted modules or services for application security components, such as identity management, encryption, and auditing and logging. Using platform features in critical security functions will reduce developers' workload and minimize the likelihood of design or implementation errors. Modern operating systems provide effective mechanisms for identification, authentication, and authorization and make those mechanisms available to applications. Use only standardized, currently accepted, and extensively reviewed encryption algorithms. Operating systems also provide mechanisms to create and maintain secure audit logs.		●	●
v7	9.4 <u>Apply Host-based Firewalls or Port Filtering</u> Apply host-based firewalls or port filtering tools on end systems, with a default-deny rule that drops all traffic except those services and ports that are explicitly allowed.	●	●	●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1609, T1610	TA0002	M1038

1.2.14 Ensure that the `--insecure-bind-address` argument is not set (Manual)

Profile Applicability:

- Level 1

Description:

Do not bind the insecure API service.

Rationale:

If you bind the `apiserver` to an insecure address, basically anyone who could connect to it over the insecure port, would have unauthenticated and unencrypted access to your master node. The `apiserver` doesn't do any authentication checking for insecure binds and traffic to the Insecure API port is not encrypted, allowing attackers to potentially read sensitive data in transit.

Impact:

Connections to the API server will require valid authentication credentials.

Audit:

The `openshift-kube-apiserver` is served over HTTPS with authentication and authorization; the secure API endpoint for the `openshift-kube-apiserver` is bound to `0.0.0.0:6443` by default. Note that the `openshift-apiserver` is not running in the host network namespace. The port is not exposed on the node, but only through the pod network.

Use the following command to obtain a list of configured feature gates on the API server.

```
oc get kubeapiservers.operator.openshift.io cluster -ojson | jq
'.spec.observedConfig.apiServerArguments."feature-gates"'
```

Verify that `InsecureBindAddress=true` is not in the returned list.

Next, query Kubernetes API server endpoints:

```
oc -n openshift-kube-apiserver get endpoints -o
jsonpath='{.items[*].subsets[*].ports[*].port}'
```

Verify the API server port for the Kubernetes API server is using `6443`.

Next, query the OpenShift API server endpoints:

```
oc -n openshift-apiserver get endpoints -o
jsonpath='{.items[*].subsets[*].ports[*].port}'
```

Verify the API server port for the OpenShift API server is using `8443`.

Note that the **openshift-apiserver** is not running in the host network namespace. The port is not exposed on the node, but only through the pod network.

Remediation:

None.





Default Value:

By default, the **openshift-kube-apiserver** is served over HTTPS with authentication and authorization; the secure API endpoint is **bound** to **0.0.0.0:6443**. Note that the **openshift-apiserver** is not running in the host network namespace. The port is not exposed on the node, but only through the pod network.

References:

1. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.5/bindata/v4.1.0/config/defaultconfig.yaml#L104-L105>
2. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
3. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	16.11 <u>Leverage Vetted Modules or Services for Application Security Components</u> Leverage vetted modules or services for application security components, such as identity management, encryption, and auditing and logging. Using platform features in critical security functions will reduce developers' workload and minimize the likelihood of design or implementation errors. Modern operating systems provide effective mechanisms for identification, authentication, and authorization and make those mechanisms available to applications. Use only standardized, currently accepted, and extensively reviewed encryption algorithms. Operating systems also provide mechanisms to create and maintain secure audit logs.			
v7	9.2 <u>Ensure Only Approved Ports, Protocols and Services Are Running</u> Ensure that only network ports, protocols, and services listening on a system with validated business needs, are running on each system.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1106	TA0002	M1035

1.2.15 Ensure that the `--insecure-port` argument is set to 0 (Manual)

Profile Applicability:

- Level 1

Description:

Do not bind to insecure port.

Rationale:

Setting up the `apiserver` to serve on an insecure port would allow unauthenticated and unencrypted access to your master node. This would allow attackers who could access this port, to easily take control of the cluster.

Impact:

All components that use the API must connect via the secured port, authenticate themselves, and be authorized to use the API.

This includes:

- kube-controller-manager
- kube-proxy
- kube-scheduler
- kubelets

Audit:

The `openshift-kube-apiserver` is served over HTTPS with authentication and authorization; the secure API endpoint is bound to `0.0.0.0:6443` by default. By default the `insecure-port` argument is set to 0. Note that the `openshift-apiserver` is not running in the host network namespace. The port is not exposed on the node, but only through the pod network.

Run the following command:

```
oc -n openshift-kube-apiserver get endpoints -o jsonpath='{.items[*].subsets[*].ports[*].port}'
```

Verify the return value is `6443`.

Remediation:

None.

Default Value:

By default, the **openshift-kube-server** is served over HTTPS with authentication and authorization; the secure API endpoint is bound to **0.0.0.0:6443** and the **insecure-port** has been removed in Kubernetes [1.20+](#).

References:

1. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.6/bindata/v4.1.0/config/defaultconfig.yaml#L102-L103>
2. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.5/bindata/v4.1.0/config/defaultconfig.yaml#L103-L105>
3. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.5/bindata/v4.1.0/kube-apiserver/pod.yaml#L155-L157>
4. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
5. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators
6. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	16.11 Leverage Vetted Modules or Services for Application Security Components Leverage vetted modules or services for application security components, such as identity management, encryption, and auditing and logging. Using platform features in critical security functions will reduce developers' workload and minimize the likelihood of design or implementation errors. Modern operating systems provide effective mechanisms for identification, authentication, and authorization and make those mechanisms available to applications. Use only standardized, currently accepted, and extensively reviewed encryption algorithms. Operating systems also provide mechanisms to create and maintain secure audit logs.		●	●
v7	9.4 Apply Host-based Firewalls or Port Filtering Apply host-based firewalls or port filtering tools on end systems, with a default-deny rule that drops all traffic except those services and ports that are explicitly allowed.	●	●	●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1609	TA0002	M1035

1.2.16 Ensure that the `--secure-port` argument is not set to 0 (Manual)

Profile Applicability:

- Level 1

Description:

Do not disable the secure port.

Rationale:

The secure port is used to serve https with authentication and authorization. If you disable it, no https traffic is served and all traffic is served unencrypted.

Impact:

You need to set the API Server up with the right TLS certificates.

Audit:

The `openshift-kube-apiserver` is served over HTTPS with authentication and authorization; the secure API endpoint is bound to `0.0.0.0:6443` by default. In OpenShift, the only supported way to access the API server pod is through the load balancer and then through the internal service. The value is set by the `bindAddress` argument under the `servingInfo` parameter.

Run the following command:

```
oc get kubeapiservers.operator.openshift.io cluster -o json | jq
'.spec.observedConfig.servingInfo.bindAddress'
```

Verify the bind address is `0.0.0.0:6443`.

```
oc get pods -n openshift-kube-apiserver -l app=openshift-kube-apiserver -o
jsonpath='{.items[*].spec.containers[?(@.name=="kube-
apiserver")].ports[*].containerPort}'
```

Verify the ports returned are `6443`.

Remediation:

None.

Default Value:





By default, the `openshift-kube-apiserver` is served over HTTPS with authentication and authorization; the secure API endpoint is bound to `0.0.0.0:6443`. Note that the `openshift-apiserver` is not running in the host network namespace. The port is not exposed on the node, but only through the pod network.

The OpenShift platform manages the TLS certificates for the API servers. External access is only available through the load balancer and then through the internal service.

References:

1. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.6/bindata/v4.1.0/config/defaultconfig.yaml#L102-L103>
2. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.5/bindata/v4.1.0/config/defaultconfig.yaml#L103-L105>
3. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
4. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators
5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.10 Encrypt Sensitive Data in Transit Encrypt sensitive data in transit. Example implementations can include: Transport Layer Security (TLS) and Open Secure Shell (OpenSSH).			
v7	14.4 Encrypt All Sensitive Information in Transit Encrypt all sensitive information in transit.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1609	TA0002	M1035

1.2.17 Ensure that the healthz endpoint is protected by RBAC (Manual)

Profile Applicability:

- Level 1

Description:

Disable profiling, if not needed.

Rationale:

Profiling allows for the identification of specific performance bottlenecks. It generates a significant amount of program data that could potentially be exploited to uncover system and program details. If you are not experiencing any bottlenecks and do not need the profiler for troubleshooting purposes, it is recommended to turn it off to reduce the potential attack surface.

Impact:

Profiling information would not be available.

Audit:

Profiling is enabled by default in OpenShift. The API server operators expose Prometheus metrics via the metrics service. Profiling data is sent to `healthzPort`, the port of the localhost `healthz` endpoint. Changing this value may disrupt components that monitor the kubelet health. The default port value is 10248, and the `healthz BindAddress` is `127.0.0.1`

To ensure the collected data is not exploited, profiling endpoints are exposed at each master port and secured via RBAC (see `cluster-debugger` role). By default, the profiling endpoints are accessible only by users bound to `cluster-admin` or `cluster-debugger` role.

Profiling cannot be disabled.

To verify the configuration, run the following command:

Run the following command to verify the Kubernetes API server endpoints:

```
oc -n openshift-kube-apiserver describe endpoints
```

Use the following steps to ensure Kubernetes API server metrics are protected by RBAC:

```
oc project openshift-kube-apiserver
export POD=$(oc get pods -n openshift-kube-apiserver -l app=openshift-kube-apiserver -o jsonpath='{.items[0].metadata.name}')
export PORT=$(oc get pods -n openshift-kube-apiserver -l app=openshift-kube-apiserver -o jsonpath='{.items[0].spec.containers[0].ports[0].hostPort}')
```

Verify unauthenticated access returns an **HTTP 403**:

```
oc rsh -n openshift-kube-apiserver $POD curl https://localhost:$PORT/metrics -k
```

Create a service account to test RBAC

```
oc create -n openshift-kube-apiserver sa permission-test-sa
export SA_TOKEN=$(oc sa -n openshift-kube-apiserver get-token permission-test-sa)
```

Verify that a service account cannot access metrics endpoints:

```
oc rsh -n openshift-kube-apiserver $POD curl https://localhost:$PORT/metrics -H "Authorization: Bearer $SA_TOKEN" -k
```

Verify a cluster administrator can access metrics:

```
export CLUSTER_ADMIN_TOKEN=$(oc whoami -t)
oc rsh -n openshift-kube-apiserver $POD curl https://localhost:$PORT/metrics -H "Authorization: Bearer $CLUSTER_ADMIN_TOKEN" -k
```

Clean up service account and environment variables:

```
oc delete -n openshift-kube-apiserver sa permission-test-sa
unset CLUSTER_ADMIN_TOKEN SA_TOKEN POD PORT
```

Remediation:

None.





Default Value:

By default, profiling is enabled and protected by RBAC.

References:

1. <https://github.com/openshift/kubernetes-kubelet/blob/master/config/v1beta1/types.go#L259-L277>
2. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.5/bindata/v4.1.0/kube-apiserver/pod.yaml#L71-L84>
3. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
4. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators
5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
6. <https://github.com/kubernetes/community/blob/master/contributors/devel/profiling.md>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	6.8 <u>Define and Maintain Role-Based Access Control</u> Define and maintain role-based access control, through determining and documenting the access rights necessary for each role within the enterprise to successfully carry out its assigned duties. Perform access control reviews of enterprise assets to validate that all privileges are authorized, on a recurring schedule at a minimum annually, or more frequently.			
v7	14.6 <u>Protect Information through Access Control Lists</u> Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078, T1548	TA0001, TA0004	M1018

1.2.18 Ensure that the `--audit-log-path` argument is set (Manual)

Profile Applicability:

- Level 1

Description:

Enable auditing on the Kubernetes API Server and set the desired audit log path.

Rationale:

Auditing the Kubernetes API Server provides a security-relevant chronological set of records documenting the sequence of activities that have affected the system by individual users, administrators or other components of the system. Even though currently, Kubernetes provides only basic audit capabilities, it should be enabled. You can enable it by setting an appropriate audit log path.

Impact:

None

Audit:

OpenShift audit works at the API server level, logging all requests coming to the server. API server audit is on by default.

Run the following command to find the Kubernetes API audit log path:

```
oc get configmap config -n openshift-kube-apiserver -ojson | jq -r  
'data["config.yaml"]' | jq '.apiServerArguments["audit-log-path"]'
```

Verify the path is `/var/log/kube-apiserver/audit.log`.

Use the following to verify the audit log exists:

```
export POD=$(oc get pods -n openshift-kube-apiserver -l app=openshift-kube-  
apiserver -o jsonpath='{.items[0].metadata.name}')
```

```
oc rsh -n openshift-kube-apiserver -c kube-apiserver $POD ls /var/log/kube-  
apiserver/audit.log  
echo $?
```

Verify a return code of `0`.

Run the following command to find the OpenShift API audit log path:

```
oc get configmap config -n openshift-apiserver -ojson | jq -r  
'data["config.yaml"]' | jq '.apiServerArguments["audit-log-path"]'
```

Verify the path is `/var/log/openshift-apiserver/audit.log`.

Use the following to verify the audit log exists:

```
export POD=$(oc get pods -n openshift-apiserver -l apiserver=true -o
jsonpath='{.items[0].metadata.name}')
oc rsh -n openshift-apiserver $POD ls /var/log/openshift-apiserver/audit.log
echo $?
```

Verify a return code of 0.

Remediation:

None required. This is managed by the cluster **apiserver** operator.








Default Value:

By default, auditing is enabled.

References:

1. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators
2. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
4. <https://kubernetes.io/docs/tasks/debug-application-cluster/audit/>
5. <https://github.com/kubernetes/features/issues/22>
6. <https://docs.openshift.com/container-platform/4.13/security/audit-log-view.html>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	8.5 Collect Detailed Audit Logs Configure detailed audit logging for enterprise assets containing sensitive data. Include event source, date, username, timestamp, source addresses, destination addresses, and other useful elements that could assist in a forensic investigation.			
v7	6.2 Activate audit logging Ensure that local logging has been enabled on all systems and networking devices.			
v7	6.3 Enable Detailed Logging Enable system logging to include detailed information such as an event source, date, user, timestamp, source addresses, destination addresses, and other useful elements.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1543	TA0003, TA0004	M1047

1.2.19 Ensure that the audit logs are forwarded off the cluster for retention (Manual)

Profile Applicability:

- Level 1

Description:

Retain the logs for at least 30 days or as appropriate.

Rationale:

Retaining logs for at least 30 days ensures that you can go back in time and investigate or correlate any events. Set your audit log retention period to 30 days or as per your business requirements.

Impact:

None.

Audit:

OpenShift audit works at the API server level, logging all requests coming to the server. Audit is on by default. Best practice is to ship audit logs off the cluster for retention.

OpenShift includes the optional Cluster Logging operator and the `elasticSearch` operator. OpenShift cluster logging can be configured to send logs to destinations outside of your OpenShift Container Platform cluster instead of the default `elasticsearch` log store using the following methods:

- Sending logs using the `fluentd` forward protocol. You can create a `ConfigMap` to use the `fluentdForward` protocol to securely send logs to an external logging aggregator that accepts the `fluentdForward` protocol.
- Sending logs using syslog. You can create a `ConfigMap` to use the syslog protocol to send logs to an external syslog (RFC 3164) server.

Alternatively, you can use the Log Forwarding API, currently in Technology Preview. The Log Forwarding API, which is easier to configure than the `fluentdForward` protocol and syslog, exposes configuration for sending logs to the internal `elasticsearch` log store and to external `fluentd` log aggregation solutions.

You cannot use the `ConfigMap` methods and the Log Forwarding API in the same cluster.

Verify that audit log forwarding is configured as appropriate.

Remediation:

Follow the documentation for log forwarding. [Forwarding logs to third party systems.](#)






Default Value:

By default, auditing is enabled.

References:

1. https://docs.openshift.com/container-platform/latest/logging/log_collection_forwarding/log-forwarding.html

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	8.1 <u>Establish and Maintain an Audit Log Management Process</u> Establish and maintain an audit log management process that defines the enterprise's logging requirements. At a minimum, address the collection, review, and retention of audit logs for enterprise assets. Review and update documentation annually, or when significant enterprise changes occur that could impact this Safeguard.			
v7	6.4 <u>Ensure adequate storage for logs</u> Ensure that all systems that store logs have adequate storage space for the logs generated.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1543	TA0003, TA0004	M1047

1.2.20 Ensure that the `maximumRetainedFiles` argument is set to 10 or as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

Retain 10 or an appropriate number of old log files.

Rationale:

Kubernetes automatically rotates the log files. Retaining old log files ensures that you would have sufficient log data available for carrying out any investigation or correlation. For example, if you have set file size of 100 MB and the number of old log files to keep as 10, you would have approximately 1 GB of log data that you could potentially use for your analysis.

Impact:

None.

Audit:

OpenShift audit works at the API server level, logging all requests coming to the server.

Run the following command to verify the number of retained log files:

```
oc get configmap config -n openshift-kube-apiserver -ojson | jq -r '.data["config.yaml"]' | jq -r '.apiServerArguments["audit-log-maxbackup"][]?'
```

Verify the output is at least **10**.

Remediation:

None.

Default Value:






By default, auditing is enabled and the maximum audit log backup is set to **10**.

References:

1. <https://access.redhat.com/solutions/4262201>
2. <https://docs.openshift.com/container-platform/latest/security/audit-log-policy-config.html>
3. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/master/bindata/v4.1.0/config/defaultconfig.yaml#L165-168>

4. <https://github.com/openshift/cluster-authentication-operator/blob/master/bindata/oauth-apiserver/deploy.yaml>
5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
6. <https://kubernetes.io/docs/tasks/debug-application-cluster/audit/>
7. <https://github.com/kubernetes/features/issues/22>
8. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	8.3 <u>Ensure Adequate Audit Log Storage</u> Ensure that logging destinations maintain adequate storage to comply with the enterprise's audit log management process.			
v7	6.4 <u>Ensure adequate storage for logs</u> Ensure that all systems that store logs have adequate storage space for the logs generated.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1543	TA0003, TA0004	M1047

1.2.21 Configure Kubernetes API Server Maximum Audit Log Size (Manual)

Profile Applicability:

- Level 1

Description:

Audit logs are rotated upon reaching a maximum size, which is set to 100 MB or greater by default.

Rationale:

OpenShift automatically rotates the log files. Retaining old log files ensures that you would have sufficient log data available for carrying out any investigation or correlation. If you have set file size of 100 MB and the number of old log files to keep as 10, you would have approximately 1 GB of log data that you could potentially use for your analysis.

Impact:

Overriding can be done by using `unsupportedConfigOverrides`, it is NOT supported by Red Hat and will block future OpenShift Container Platform Upgrades.

Audit:

OpenShift audit works at the API server level, logging all requests coming to the server. Configure via `audit-log-maxsize`.

Run the following command:

```
oc get configmap config -n openshift-kube-apiserver -ojson | jq -r '.data["config.yaml"]' | jq -r '.apiServerArguments["audit-log-maxsize"][]?'
```

Verify that the `audit-log-maxsize` argument is set to **100** or greater

Remediation:

The `audit-log-maxsize` parameter is set by default by Red Hat and not supported to change.

Default Value:

By default, auditing is enabled.

References:

1. <https://access.redhat.com/solutions/4262201>
2. <https://docs.openshift.com/container-platform/4.5/nodes/nodes/nodes-nodes-audit-log.html>

3. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
4. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators
5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
6. <https://kubernetes.io/docs/tasks/debug-application-cluster/audit/>
7. <https://github.com/kubernetes/features/issues/22>
8. https://github.com/ComplianceAsCode/content/blob/master/applications/openshift/api-server/api_server_audit_log_maxsize/rule.yml
9. <https://access.redhat.com/solutions/5993251>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	8.3 <u>Ensure Adequate Audit Log Storage</u> Ensure that logging destinations maintain adequate storage to comply with the enterprise's audit log management process.	●	●	●
v7	6.4 <u>Ensure adequate storage for logs</u> Ensure that all systems that store logs have adequate storage space for the logs generated.		●	●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1543	TA0003, TA0004	M1047

1.2.22 Ensure that the `--request-timeout` argument is set (Manual)

Profile Applicability:

- Level 1

Description:

The API server minimum request timeout defines the minimum number of seconds a handler must keep a request open before timing it out.

Rationale:

Setting global request timeout allows extending the API server request timeout limit to a duration appropriate to the user's connection speed. By default, it is set to 3600 seconds in OpenShift 4. Allowing users to set this timeout limit to be too small can be insufficient for some connections and too large can exhaust the API server resources making it prone to Denial-of-Service attack. Hence, it is not supported to adjust this value in OpenShift 4.

Impact:

None

Audit:

OpenShift configures the `min-request-timeout` flag via `apiServerArguments[min-request-timeout]`, which overrides `request-timeout` and provides a more balanced timeout approach.

Run the following command:

```
oc get configmap config -n openshift-kube-apiserver -ojson | jq -r
'.data["config.yaml"]' |
jq '.apiServerArguments["min-request-timeout"]'
```

Verify that the `min-request-timeout` argument is set to 3600.

Remediation:

None

Default Value:







By default, `min-request-timeout` is set to 3600 seconds in OpenShift 4

References:

1. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
2. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators

3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
4. <https://github.com/kubernetes/kubernetes/pull/51415>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.1 Establish and Maintain a Secure Configuration Process Establish and maintain a secure configuration process for enterprise assets (end-user devices, including portable and mobile, non-computing/IoT devices, and servers) and software (operating systems and applications). Review and update documentation annually, or when significant enterprise changes occur that could impact this Safeguard.			
v7	5.1 Establish Secure Configurations Maintain documented, standard security configuration standards for all authorized operating systems and software.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1499	TA0040	M1037

1.2.23 Ensure that the `--service-account-lookup` argument is set to `true` (Manual)

Profile Applicability:

- Level 1

Description:

Validate service account before validating token.

Rationale:

If `--service-account-lookup` is not enabled, the `apiserver` only verifies that the authentication token is valid, and does not validate that the service account token mentioned in the request is actually present in `etcd`. This allows using a service account token even after the corresponding service account is deleted. This is an example of time of check to time of use security issue.

Impact:

None.

Audit:

OpenShift denies access for any OAuth Access token that does not exist in its `etcd` data store. OpenShift does not use the `service-account-lookup` flag even when it is set.

Run the following command:

```
oc get configmap config -n openshift-kube-apiserver -ojson | jq -r '.data["config.yaml"]' | jq '.apiServerArguments."service-account-lookup"[]'
```

Verify that if the `--service-account-lookup` argument exists it is set to `true`.

Remediation:

None.

Default Value:





Service account lookup is enabled by default.

References:

1. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
2. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators

3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
4. <https://github.com/kubernetes/kubernetes/issues/24167>
5. https://en.wikipedia.org/wiki/Time-of-check_to_time-of-use

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 <u>Configure Data Access Control Lists</u> Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.7 <u>Enforce Access Control to Data through Automated Tools</u> Use an automated tool, such as host-based Data Loss Prevention, to enforce access controls to data even when data is copied off a system.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078	TA0001	M1026

1.2.24 Ensure that the `--service-account-key-file` argument is set as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

Explicitly set a service account public key file for service accounts on the apiserver.

Rationale:

By default, if no `--service-account-key-file` is specified to the `apiserver`, it uses the private key from the TLS serving certificate to verify service account tokens. To ensure that the keys for service account tokens could be rotated as needed, a separate public/private key pair should be used for signing service account tokens. Hence, the public key should be specified to the `apiserver` with `--service-account-key-file`.

Impact:

The corresponding private key must be provided to the controller manager. You would need to securely maintain the key file and rotate the keys based on your organization's key rotation policy.

Audit:

OpenShift API server does not use the `service-account-key-file` argument. OpenShift does not reuse the `apiserver` TLS key. The `ServiceAccount` token authenticator is configured with `serviceAccountConfig.publicKeyFiles`. OpenShift automatically manages and rotates the keys.

Run the following command:

```
oc get configmap config -n openshift-kube-apiserver -ojson | \
jq -r '.data["config.yaml"]' | \
jq -r '.serviceAccountPublicKeyFiles[]'
```

Verify that the following is returned.

```
/etc/kubernetes/static-pod-resources/configmaps/sa-token-signing-certs
/etc/kubernetes/static-pod-resources/configmaps/bound-sa-token-
signing-certs
```

Remediation:

The OpenShift API server does not use the `service-account-key-file` argument. The `ServiceAccount` token `authenticator` is configured with `serviceAccountConfig.publicKeyFiles`. OpenShift does not reuse the `apiserver` TLS key. This is not configurable.






Default Value:

The OpenShift API server does not use the `service-account-key-file` argument. The `ServiceAccount` token authenticator is configured with `serviceAccountConfig.publicKeyFiles`. OpenShift does not reuse the apiserver TLS key.

References:

1. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#kube-apiserver-operator_red-hat-operators
2. https://docs.openshift.com/container-platform/4.5/operators/operator-reference.html#openshift-apiserver-operator_red-hat-operators
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
4. <https://github.com/kubernetes/kubernetes/issues/24167>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.2 Use Unique Passwords Use unique passwords for all enterprise assets. Best practice implementation includes, at a minimum, an 8-character password for accounts using MFA and a 14-character password for accounts not using MFA.			
v7	4.4 Use Unique Passwords Where multi-factor authentication is not supported (such as local administrator, root, or service accounts), accounts will use passwords that are unique to that system.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078	TA0001	M1026

1.2.25 Ensure that the `--etcd-certfile` and `--etcd-keyfile` arguments are set as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

`etcd` should be configured to make use of TLS encryption for client connections.

Rationale:

`etcd` is a highly-available key value store used by Kubernetes deployments for persistent storage of all of its REST API objects. These objects are sensitive in nature and should be protected by client authentication. This requires the API server to identify itself to the `etcd` server using a client certificate and key.

Impact:

TLS and client certificate authentication are configured by default for `etcd`.

Audit:

OpenShift uses X.509 certificates to provide secure communication to `etcd`. OpenShift configures these automatically. OpenShift does not use the `etcd-certfile` or `etcd-keyfile` flags. Certificates are used for encrypted communication between `etcd` member peers, as well as encrypted client traffic. The following certificates are generated and used by `etcd` and other processes that communicate with `etcd`:

- Peer certificates: Used for communication between `etcd` members.
- Client certificates: Used for encrypted server-client communication. Client certificates are currently used by the API server only, and no other service should connect to `etcd` directly except for the proxy. Client secrets (`etcd-client`, `etcd-metric-client`, `etcd-metric-signer`, and `etcd-signer`) are added to the `openshift-config`, `openshift-monitoring`, and `openshift-kube-apiserver` namespaces.
- Server certificates: Used by the `etcd` server for authenticating client requests.
- Metric certificates: All metric consumers connect to proxy with `metric-client` certificates.

Run the following command to check the location of the `etc-certfile`:

```
oc get configmap config -n openshift-kube-apiserver -ojson | \
jq -r '.data["config.yaml"]' | \
jq -r '.apiServerArguments["etcd-certfile"]'
```

Verify that `/etc/kubernetes/static-pod-resources/secrets/etcd-client/tls.crt` is returned.

Run the following command to check the location of the **etc-keyfile**:

```
oc get configmap config -n openshift-kube-apiserver -ojson | \
jq -r '.data["config.yaml"]' | \
jq -r '.apiServerArguments["etcd-keyfile"]'
```

Verify that **/etc/kubernetes/static-pod-resources/secrets/etcd-client/tls.key** is returned.

Remediation:

OpenShift automatically manages TLS and client certificate authentication for **etcd**. This is not configurable.

Default Value:

By default, OpenShift uses X.509 certificates to provide secure communication to **etcd**. OpenShift configures these automatically. OpenShift does not use the **etcd-certfile** or **etcd-keyfile** flags. OpenShift generates the necessary files and sets the arguments appropriately.

References:

1. https://docs.openshift.com/container-platform/latest/security/certificate_types_descriptions/etcd-certificates.html
2. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
3. <https://etcd.io/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.11 <u>Encrypt Sensitive Data at Rest</u> Encrypt sensitive data at rest on servers, applications, and databases containing sensitive data. Storage-layer encryption, also known as server-side encryption, meets the minimum requirement of this Safeguard. Additional encryption methods may include application-layer encryption, also known as client-side encryption, where access to the data storage device(s) does not permit access to the plain-text data.		●	●
v7	14.8 <u>Encrypt Sensitive Information at Rest</u> Encrypt all sensitive information at rest using a tool that requires a secondary authentication mechanism not integrated into the operating system, in order to access the information.			●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

1.2.26 Ensure that the `--tls-cert-file` and `--tls-private-key-file` arguments are set as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

Setup TLS connection on the API server.

Rationale:

API server communication contains sensitive parameters that should remain encrypted in transit. Configure the API server to serve only HTTPS traffic.

Impact:

TLS and client certificate authentication must be configured for your Kubernetes cluster deployment. By default, OpenShift uses X.509 certificates to provide secure connections between the API server and node/kubelet. OpenShift Container Platform monitors certificates for proper validity, for the cluster certificates it issues and manages. The OpenShift Container Platform manages certificate rotation and the alerting framework has rules to help identify when a certificate issue is about to occur.

Audit:

OpenShift uses X.509 certificates to provide secure connections between API server and node/kubelet by default. OpenShift does not use values assigned to the `tls-cert-file` or `tls-private-key-file` flags. OpenShift generates the certificate files and sets the arguments appropriately.

The API server is accessible by clients external to the cluster at `api.<cluster_name>.<base_domain>`. The administrator must set a custom default certificate to be used by the API server when serving content in order to enable clients to access the API server at a different host name or without the need to distribute the cluster-managed certificate authority (CA) certificates to the clients.

Run the following command to obtain the API server TLS certificate file:

```
oc get configmap config -n openshift-kube-apiserver -ojson | jq -r '.data["config.yaml"]' | jq -r '.apiServerArguments."tls-cert-file"[]'
```

Verify the output is `/etc/kubernetes/static-pod-certs/secrets/service-network-serving-certkey/tls.crt`.

Run the following command to obtain the API server TLS private key file:

```
oc get configmap config -n openshift-kube-apiserver -ojson | jq -r '.data["config.yaml"]' | jq -r '.apiServerArguments."tls-private-key-file"[]'
```

Verify the output is `/etc/kubernetes/static-pod-certs/secrets/service-network-serving-certkey/tls.key`.

Remediation:

None.

Default Value:

By default, OpenShift uses X.509 certificates to provide secure connections between the API server and node/kubelet. OpenShift does not use values assigned to the `tls-cert-file` or `tls-private-key-file` flags.





You may optionally set a custom default certificate to be used by the API server when serving content in order to enable clients to access the API server at a different host name or without the need to distribute the cluster-managed certificate authority (CA) certificates to the clients. Follow the directions in the OpenShift documentation

[User-provided certificates for the API server](#)

References:

1. <https://docs.openshift.com/container-platform/latest/security/certificates/api-server.html>
2. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
3. <https://rootsquash.com/2016/05/10/securing-the-kubernetes-api/>
4. <https://github.com/kelseyhightower/docker-kubernetes-tls-guide>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.10 <u>Encrypt Sensitive Data in Transit</u> Encrypt sensitive data in transit. Example implementations can include: Transport Layer Security (TLS) and Open Secure Shell (OpenSSH).			
v7	14.4 <u>Encrypt All Sensitive Information in Transit</u> Encrypt all sensitive information in transit.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

1.2.27 Ensure that the `--client-ca-file` argument is set as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

Setup TLS connection on the API server.

Rationale:

API server communication contains sensitive parameters that should remain encrypted in transit. Configure the API server to serve only HTTPS traffic. If `--client-ca-file` argument is set, any request presenting a client certificate signed by one of the authorities in the `client-ca-file` is authenticated with an identity corresponding to the CommonName of the client certificate.

Impact:

TLS and client certificate authentication must be configured for your Kubernetes cluster deployment. By default, OpenShift uses X.509 certificates to provide secure connections between the API server and node/kubelet. OpenShift Container Platform monitors certificates for proper validity, for the cluster certificates it issues and manages. The OpenShift Container Platform alerting framework has rules to help identify when a certificate issue is about to occur. These rules consist of the following checks:

- API server client certificate expiration is less than five minutes.

Audit:

OpenShift uses X.509 certificates to provide secure connections between API server and node/kubelet by default. OpenShift configures the `client-ca-file` value and does not use value assigned to the `client-ca-file` flag. OpenShift generates the necessary files and sets the arguments appropriately.

The API server is accessible by clients external to the cluster at `api.<cluster_name>.<base_domain>`. The administrator must set a custom default certificate to be used by the API server when serving content in order to enable clients to access the API server at a different host name or without the need to distribute the cluster-managed certificate authority (CA) certificates to the clients.

Run the following command:

```
oc get configmap config -n openshift-kube-apiserver -ojson | \
jq -r '.data["config.yaml"]' | \
jq -r '.-servingInfo.clientCA'
Version 4.10+

[core@00-50-56-a6-f5-f6 ~]$ oc get configmap config -n openshift-kube-
apiserver -o json | jq -r '.data["config.yaml"]' | jq -r '.apiServerArguments
."client-ca-file"'
[
  "/etc/kubernetes/static-pod-certs/configmaps/client-ca/ca-bundle.crt"
]
```

Verify that the following file exists.

</etc/kubernetes/static-pod-certs/configmaps/client-ca/ca-bundle.crt>

Remediation:

None.

Default Value:

By default, OpenShift configures the [client-ca-file](#) and automatically manages the certificate. It does not use the value assigned to the [client-ca-file](#) flag.



You may optionally set a custom default certificate to be used by the API server when serving content in order to enable clients to access the API server at a different host name or without the need to distribute the cluster-managed certificate authority (CA) certificates to the clients.



Please follow the OpenShift [documentation](#) for providing certificates for OpenShift to use.

References:

1. https://docs.openshift.com/container-platform/latest/security/certificate_types_descriptions/user-provided-certificates-for-api-server.html
2. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
3. <https://rootsquash.com/2016/05/10/securing-the-kubernetes-api/>
4. <https://github.com/kelseyhightower/docker-kubernetes-tls-guide>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.10 <u>Encrypt Sensitive Data in Transit</u> Encrypt sensitive data in transit. Example implementations can include: Transport Layer Security (TLS) and Open Secure Shell (OpenSSH).			

Controls Version	Control	IG 1	IG 2	IG 3
v7	14.4 <u>Encrypt All Sensitive Information in Transit</u> Encrypt all sensitive information in transit.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

1.2.28 Ensure that the `--etcd-cafile` argument is set as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

`etcd` should be configured to make use of TLS encryption for client connections.

Rationale:

`etcd` is a highly-available key value store used by Kubernetes deployments for persistent storage of all of its REST API objects. These objects are sensitive in nature and should be protected by client authentication. This requires the API server to identify itself to the `etcd` server using a SSL Certificate Authority file.

Impact:

TLS and client certificate authentication must be configured for `etcd`.

Audit:

OpenShift uses X.509 certificates to provide secure communication to `etcd`. OpenShift does not use values assigned to the `etcd-cafile` argument. OpenShift generates the `etcd-cafile` and sets the arguments appropriately in the API server. OpenShift includes multiple certificate authorities (CAs) providing independent chains of trust, increasing the security posture of the cluster. The certificates generated by each CA are used to identify a particular OpenShift platform component to another OpenShift platform component. Communication with `etcd` is secured by the `etcd` serving CA.

Run the following command

```
oc get configmap config -n openshift-kube-apiserver -ojson | \
jq -r '.data["config.yaml"]' | \
jq -r '.apiServerArguments["etcd-cafile"]'
```

Verify that the following is returned

`/etc/kubernetes/static-pod-resources/configmaps/etcd-serving-ca/ca-bundle.crt`

Remediation:

None.





Default Value:

By default, OpenShift uses X.509 certificates to provide secure communication to **etcd**. OpenShift does not use values assigned to **etcd-cafile**. OpenShift generates the **etcd-cafile** and sets the arguments appropriately in the API server. Communication with **etcd** is secured by the **etcd** serving CA.

References:

1. https://docs.openshift.com/container-platform/latest/security/certificate_types_descriptions/etcd-certificates.html
2. <https://docs.openshift.com/container-platform/latest/operators/index.html>
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
4. <https://etcd.io/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.10 <u>Encrypt Sensitive Data in Transit</u> Encrypt sensitive data in transit. Example implementations can include: Transport Layer Security (TLS) and Open Secure Shell (OpenSSH).			
v7	14.4 <u>Encrypt All Sensitive Information in Transit</u> Encrypt all sensitive information in transit.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

1.2.29 Ensure that encryption providers are appropriately configured (Manual)

Profile Applicability:

- Level 1

Description:

Where **etcd** encryption is used, appropriate providers should be configured.

Rationale:

Where **etcd** encryption is used, it is important to ensure that the appropriate set of encryption providers is used. Currently, the **aescbc**, **kms** and **secretbox** are likely to be appropriate options.

Impact:

When you enable **etcd** encryption, the following OpenShift API server and Kubernetes API server resources are encrypted:

- Secrets
- ConfigMaps
- Routes
- OAuth access tokens
- OAuth authorize tokens

When you enable **etcd** encryption, encryption keys are created. These keys are rotated on a weekly basis. You must have these keys in order to restore from an **etcd** backup.

Audit:

OpenShift supports encryption of data at rest of **etcd** datastore, but it is up to the customer to configure. The **aescbc** cipher had been the only supported cipher up to OCP 4.13. Starting with OCP 4.13, the **aescgm** cipher can be used as well. No other ciphers are supported. Keys are stored on the filesystem of the master and automatically rotated.

Run the following command to review the Encrypted status condition for the OpenShift API server to verify that its resources were successfully encrypted:

```
# encrypt the etcd datastore
oc get openshiftapiserver -o=jsonpath='{range
.items[0].status.conditions[?(@.type=="Encrypted")]}{.reason}{"\n"}{.message}
{"\n"}'
```

The output shows EncryptionCompleted upon successful encryption.

- **EncryptionCompleted**

- All resources encrypted: routes.route.openshift.io, oauthaccesstokens.oauth.openshift.io, oauthauthorizetokens.oauth.openshift.io

If the output shows **EncryptionInProgress**, this means that encryption is still in progress. Wait a few minutes and try again.

Remediation:

Follow the OpenShift documentation for [encrypting etcd data](#).

Default Value:

By default, no encryption provider is set.

References:

1. <https://docs.openshift.com/container-platform/latest/security/encrypting-etcd.html>
2. <https://docs.openshift.com/container-platform/latest/operators/index.html>
3. <https://kubernetes.io/docs/tasks/administer-cluster/encrypt-data/>
4. <https://acotten.com/post/kube17-security>
5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-apiserver/>
6. <https://github.com/kubernetes/features/issues/92>
7. <https://kubernetes.io/docs/tasks/administer-cluster/encrypt-data/#providers>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.11 <u>Encrypt Sensitive Data at Rest</u> Encrypt sensitive data at rest on servers, applications, and databases containing sensitive data. Storage-layer encryption, also known as server-side encryption, meets the minimum requirement of this Safeguard. Additional encryption methods may include application-layer encryption, also known as client-side encryption, where access to the data storage device(s) does not permit access to the plain-text data.		●	●
v7	14.8 <u>Encrypt Sensitive Information at Rest</u> Encrypt all sensitive information at rest using a tool that requires a secondary authentication mechanism not integrated into the operating system, in order to access the information.			●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

1.2.30 Ensure that the API Server only makes use of Strong Cryptographic Ciphers (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the API server is configured to only use strong cryptographic ciphers.

Rationale:

TLS ciphers have had a number of known vulnerabilities and weaknesses, which can reduce the protection provided by them. By default Kubernetes supports a number of TLS ciphersuites including some that have security concerns, weakening the protection provided.

Impact:

API server clients that cannot support the custom cryptographic ciphers will not be able to make connections to the API server.

Audit:

Ciphers for the API servers, authentication operator, and ingress controller can be configured using the `tlsSecurityProfile` parameter. The ingress controller provides external access to the API server. There are four TLS security profile types:

- Old
- Intermediate
- Modern
- Custom

Only the Old, Intermediate and Custom profiles are supported at this time. Custom provides the ability to specify individual TLS security profile parameters. Follow the steps in the [documentation](#) to configure the cipher suite for Ingress and the API server.

Run the following command to obtain the TLS cipher suites used by the authentication operator:

```
oc get cm -n openshift-authentication v4-0-config-system-cliconfig -o jsonpath='{.data.v4\0-config-system-cliconfig}' | jq .servingInfo
```

Run the following command to obtain the TLS cipher suites used by the Kubernetes API server operator:

```
oc get kubeapiservers.operator.openshift.io cluster -o json | jq
.spec.observedConfig.servingInfo
```

Run the following command to obtain the TLS cipher suites used by the OpenShift API server operator:

```
oc get openshiftapiservers.operator.openshift.io cluster -o json | jq
.spec.observedConfig.servingInfo
```

Run the following command to obtain the TLS cipher suites used by the OpenShift ingress operator:

```
oc get -n openshift-ingress-operator ingresscontroller/default -o json | jq
.status.tlsProfile
```

Make sure that **tlsSecurityProfile** is not set to **Old** and if set to **Custom**, make sure that the **minTLSVersion** is not set to **VersionTLS10** or **VersionTLS11**.

Verify the **minTLSVersion** is at least **VersionTLS12**.

Note: The HAProxy Ingress controller image does not support TLS 1.3 and because the Modern profile requires TLS 1.3, it is not supported. The Ingress Operator converts the Modern profile to Intermediate. The Ingress Operator also converts the TLS 1.0 of an Old or Custom profile to 1.1, and TLS 1.3 of a Custom profile to 1.2.

Remediation:

None.

Default Value:

By default, OpenShift uses the **Intermediate** TLS profile, which requires a minimum of TLS 1.2.

You can configure TLS security profiles by following the OpenShift TLS [documentation](#).

References:

1. <https://docs.openshift.com/container-platform/latest/security/tls-security-profiles.html>
2. https://docs.openshift.com/container-platform/4.13/rest_api/config_apis/apiserver-config-openshift-io-v1.html
3. <https://github.com/ssllabs/research/wiki/SSL-and-TLS-Deployment-Best-Practices#23-use-secure-cipher-suites>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	<u>4.8 Uninstall or Disable Unnecessary Services on Enterprise Assets and Software</u> Uninstall or disable unnecessary services on enterprise assets and software, such as an unused file sharing service, web application module, or service function.		●	●
v7	<u>1.8 Utilize Client Certificates to Authenticate Hardware Assets</u> Use client certificates to authenticate hardware assets connecting to the organization's trusted network.			●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1041

1.2.31 Ensure unsupported configuration overrides are not used (Manual)

Profile Applicability:

- Level 1

Description:

OpenShift supported an option called `unsupportedConfigOverrides` that allowed users to opt into unsupported behavior. This option is no longer supported by OpenShift and should not be used.

Rationale:

Users should stop using deprecated and unmaintained features in favor of supported features.

Impact:

None. The feature is set to `null` by default and isn't used by default.

Audit:

Make sure the `unsupportedConfigOverrides` in your deployment returns `null` using the following command:

```
oc get kubeapiserver/cluster -o jsonpath='{.spec.unsupportedConfigOverrides}'
```

The output should return `null`. Any other return value is a finding and you should migrate away from that particular configuration.

Remediation:

None.







Default Value:

By default, OpenShift sets this value to `null` and doesn't support overriding configuration with unsupported features.

References:

1. <https://access.redhat.com/solutions/5170671>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	<u>2.2 Ensure Authorized Software is Currently Supported</u> Ensure that only currently supported software is designated as authorized in the software inventory for enterprise assets. If software is unsupported, yet necessary for the fulfillment of the enterprise's mission, document an exception detailing mitigating controls and residual risk acceptance. For any unsupported software without an exception documentation, designate as unauthorized. Review the software list to verify software support at least monthly, or more frequently.			
v7	<u>2.2 Ensure Software is Supported by Vendor</u> Ensure that only software applications or operating systems currently supported by the software's vendor are added to the organization's authorized software inventory. Unsupported software should be tagged as unsupported in the inventory system.			

1.3 Controller Manager

This section contains recommendations relating to Controller Manager configuration flags. In OpenShift 4, the Controller Manager is managed with the cluster Controller Manager Operator. There are two operators: the OpenShift Controller operator and the Kube Controller operator. The OpenShift Controller Manager operator manages the OpenShift Controller Manager. The Kubernetes Controller Manager operator manages and updates the [Kubernetes Controller Manager](#) deployed on top of [OpenShift](#). All calls are directed to the OpenShift Controller Manager and then Kubernetes objects are delegated to the Kubernetes Controller Manager.

1.3.1 Ensure that controller manager healthz endpoints are protected by RBAC (Manual)

Profile Applicability:

- Level 1

Description:

Disable profiling, if not needed.

Rationale:

Profiling allows for the identification of specific performance bottlenecks. It generates a significant amount of program data that could potentially be exploited to uncover system and program details. If you are not experiencing any bottlenecks and do not need the profiler for troubleshooting purposes, it is recommended to turn it off to reduce the potential attack surface.

Impact:

Profiling information would not be available.

Audit:

By default, the Controller Manager operator exposes metrics via the metrics service. Profiling data is sent to **healthzPort**, the port of the localhost **healthz** endpoint. Changing this value may disrupt components that monitor the kubelet health. To ensure the collected data is not exploited, profiling endpoints are secured via RBAC (see cluster-debugger role). By default, the profiling endpoints are accessible only by users bound to cluster-admin or cluster-debugger role.

Profiling can not be disabled.

To verify the configuration, run the following command:

Run the following command to check the **livenessProbe** configuration:

```
oc -n openshift-kube-controller-manager get cm kube-controller-manager-pod -o json | jq -r '.data."pod.yaml"' | jq '.spec.containers[].livenessProbe'
```

Verify the output path is set to **healthz**.

Run the following command to check the **readinessProbe** configuration:

```
oc -n openshift-kube-controller-manager get cm kube-controller-manager-pod -o json | jq -r '.data."pod.yaml"' | jq '.spec.containers[].readinessProbe'
```

Verify the output path is set to **healthz**.

Verify endpoints exist for the Kubernetes controller manager:

```
oc -n openshift-kube-controller-manager describe endpoints
```

Validate that RBAC is enabled and protects controller endpoints. First, switch to the **openshift-kube-controller-namespaces**:

```
oc project openshift-kube-controller-manager
```

Next, get the controller manager pod name and port:

```
export POD=$(oc get pods -l app=kube-controller-manager -o  
jsonpath='{.items[0].metadata.name}')
```

```
export PORT=$(oc get pods -l app=kube-controller-manager -o  
jsonpath='{.items[0].spec.containers[0].ports[0].hostPort}')
```

Attempt to make an insecure **GET** request to the metrics endpoint:

```
oc rsh $POD curl https://localhost:$PORT/metrics -k
```

Verify that an **HTTP 403** is returned.

Create a test service account:

```
oc create sa permission-test-sa
```

Generate a service account token and attempt to access the metrics endpoint:

```
export SA_TOKEN=$(oc create token permission-test-sa)  
oc rsh $POD curl https://localhost:$PORT/metrics -H "Authorization: Bearer  
$SA_TOKEN" -k
```

Verify that an **HTTP 403** is returned.

Login as a cluster administrator and attempt to access the metrics endpoint:

```
export CLUSTER_ADMIN_TOKEN=$(oc whoami -t)  
oc rsh $POD curl https://localhost:$PORT/metrics -H "Authorization: Bearer  
$CLUSTER_ADMIN_TOKEN" -k
```

Verify metrics output is returned. Unset environment variables used in the test and delete the test service account:

```
unset CLUSTER_ADMIN_TOKEN POD PORT SA_TOKEN  
oc delete sa permission-test-sa
```

Remediation:

None.

Default Value:

By default, the operator exposes metrics via metrics service. The metrics are collected from the OpenShift Controller Manager and the Kubernetes Controller Manager and protected by RBAC.

References:

1. <https://docs.openshift.com/container-platform/latest/monitoring/monitoring-overview.html>

2. <https://github.com/openshift/cluster-kube-controller-manager-operator/tree/master>
3. <https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/bootkube/bootstrap-manifests/kube-controller-manager-pod.yaml>
4. https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/bootkube/manifests/00_openshift-kube-controller-manager-ns.yaml
5. <https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/v4.1.0/kube-controller-manager/kubeconfig-cm.yaml>
6. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-controller-manager/>
7. <https://github.com/kubernetes/community/blob/master/contributors/devel/sig-scalability/profiling.md>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	6.8 Define and Maintain Role-Based Access Control Define and maintain role-based access control, through determining and documenting the access rights necessary for each role within the enterprise to successfully carry out its assigned duties. Perform access control reviews of enterprise assets to validate that all privileges are authorized, on a recurring schedule at a minimum annually, or more frequently.			●
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.	●	●	●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078, T1548	TA0001	M1018

1.3.2 Ensure that the `--use-service-account-credentials` argument is set to `true` (Manual)

Profile Applicability:

- Level 1

Description:

Use individual service account credentials for each controller.

Rationale:

The controller manager creates a service account per controller in the `kube-system` namespace, generates a credential for it, and builds a dedicated API client with that service account credential for each controller loop to use. Setting the `--use-service-account-credentials` to `true` runs each control loop within the controller manager using a separate service account credential. When used in combination with RBAC, this ensures that the control loops run with the minimum permissions required to perform their intended tasks.

Impact:

Whatever authorizer is configured for the cluster, it must grant sufficient permissions to the service accounts to perform their intended tasks. When using the RBAC authorizer, those roles are created and bound to the appropriate service accounts in the `kube-system` namespace automatically with default roles and `rolebindings` that are auto-reconciled on startup.

If using other authorization methods (ABAC, Webhook, etc), the cluster deployer is responsible for granting appropriate permissions to the service accounts (the required permissions can be seen by inspecting the `controller-roles.yaml` and `controller-role-bindings.yaml` files for the RBAC roles).

Audit:

In OpenShift, `--use-service-account-credentials` is set to `true` by default for the Controller Manager. The bootstrap configuration and overrides are available here:

[kube-controller-manager-pod](#)

[bootstrap-config-overrides](#)

Run the following command on the master node:

```
oc get configmaps config -n openshift-kube-controller-manager -ojson | \
jq -r '.data["config.yaml"]' | \
jq -r '.extendedArguments["use-service-account-credentials"][]'
```

Verify that the `--use-service-account-credentials` argument is set to `true`.

Remediation:

None.

Default Value:






By default, in OpenShift 4 `--use-service-account-credentials` is set to `true`.

The OpenShift Controller Manager operator manages and updates the OpenShift Controller Manager. The Kubernetes Controller Manager operator manages and updates the [Kubernetes Controller Manager](#) deployed on top of [OpenShift](#). This operator is configured via [KubeControllerManager](#) custom resource.

References:

1. <https://docs.openshift.com/container-platform/latest/operators/operator-reference.html>
2. <https://github.com/openshift/cluster-kube-controller-manager-operator/tree/master>
3. <https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/bootkube/bootstrap-manifests/kube-controller-manager-pod.yaml>
4. <https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/bootkube/config/bootstrap-config-overrides.yaml>
5. <https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/v4.1.0/kube-controller-manager/kubeconfig-cm.yaml>
6. <https://github.com/openshift/cluster-openshift-controller-manager-operator/blob/release-4.5/bindata/v3.11.0/openshift-controller-manager/ds.yaml>
7. <https://github.com/openshift/cluster-openshift-controller-manager-operator/blob/release-4.5/bindata/v3.11.0/openshift-controller-manager/sa.yaml>
8. <https://github.com/openshift/cluster-openshift-controller-manager-operator/blob/release-4.5/bindata/v3.11.0/openshift-controller-manager/separate-sa-role.yaml>
9. <https://github.com/openshift/cluster-openshift-controller-manager-operator/blob/release-4.5/bindata/v3.11.0/openshift-controller-manager/separate-sa-rolebinding.yaml>
10. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-controller-manager/>
11. <https://kubernetes.io/docs/reference/access-authn-authz/service-accounts-admin/>
12. <https://github.com/kubernetes/kubernetes/blob/release-1.6/plugin/pkg/auth/authorizer/rbac/bootstrappolicy/testdata/controller-roles.yaml>
13. <https://github.com/kubernetes/kubernetes/blob/release-1.6/plugin/pkg/auth/authorizer/rbac/bootstrappolicy/testdata/controller-role-bindings.yaml>
14. <https://kubernetes.io/docs/reference/access-authn-authz/rbac/#controller-roles>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.2 <u>Use Unique Passwords</u> Use unique passwords for all enterprise assets. Best practice implementation includes, at a minimum, an 8-character password for accounts using MFA and a 14-character password for accounts not using MFA.			
v7	4.4 <u>Use Unique Passwords</u> Where multi-factor authentication is not supported (such as local administrator, root, or service accounts), accounts will use passwords that are unique to that system.			

1.3.3 Ensure that the `--service-account-private-key-file` argument is set as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

Explicitly set a service account private key file for service accounts on the controller manager.

Rationale:

To ensure that keys for service account tokens can be rotated as needed, a separate public/private key pair should be used for signing service account tokens. The private key should be specified to the controller manager with `--service-account-private-key-file` as appropriate.

Impact:

You would need to securely maintain the key file and rotate the keys based on your organization's key rotation policy.

Audit:

OpenShift starts the Kubernetes Controller Manager with `service-account-private-key-file` set to `/etc/kubernetes/static-pod-resources/secrets/service-account-private-key/service-account.key`.

The bootstrap configuration and overrides are available here:

[kube-controller-manager-pod](#)

[bootstrap-config-overrides](#)

Run the following command:

```
oc get configmaps config -n openshift-kube-controller-manager -ojson | \
jq -r '.data["config.yaml"]' | \
jq -r '.extendedArguments["service-account-private-key-file"][]'
```

Verify that the following is returned

`/etc/kubernetes/static-pod-resources/secrets/service-account-private-key/service-account.key`

Remediation:

None.






Default Value:

By default, OpenShift starts the controller manager with `service-account-private-key-file` set to `/etc/kubernetes/static-pod-resources/secrets/service-account-private-key/service-account.key`. OpenShift manages the service account credentials for the scheduler automatically.

References:

1. <https://docs.openshift.com/container-platform/latest/operators/operator-reference.html>
2. https://docs.openshift.com/container-platform/4.13/security/certificate_types_descriptions/control-plane-certificates.html
3. <https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/bootkube/bootstrap-manifests/kube-controller-manager-pod.yaml>
4. <https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/bootkube/config/bootstrap-config-overrides.yaml>
5. <https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/v4.1.0/kube-controller-manager/kubeconfig-cm.yaml>
6. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-controller-manager/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.2 Use Unique Passwords Use unique passwords for all enterprise assets. Best practice implementation includes, at a minimum, an 8-character password for accounts using MFA and a 14-character password for accounts not using MFA.			
v7	4.4 Use Unique Passwords Where multi-factor authentication is not supported (such as local administrator, root, or service accounts), accounts will use passwords that are unique to that system.			

1.3.4 Ensure that the `--root-ca-file` argument is set as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

Allow pods to verify the API server's serving certificate before establishing connections.

Rationale:

Processes running within pods that need to contact the API server must verify the API server's serving certificate. Failing to do so could be a subject to man-in-the-middle attacks.

Providing the root certificate for the API server's serving certificate to the controller manager with the `--root-ca-file` argument allows the controller manager to inject the trusted bundle into pods so that they can verify TLS connections to the API server.

Impact:

OpenShift clusters manage and maintain certificate authorities and certificates for cluster components.

Audit:

Certificates for OpenShift platform components are automatically created and rotated by the OpenShift Container Platform.

Run the following command:

```
oc get configmaps config -n openshift-kube-controller-manager -ojson | \
jq -r '.data["config.yaml"]' | \
jq -r '.extendedArguments["root-ca-file"][]'
```

Verify that the `--root-ca-file` argument exists and is set to `/etc/kubernetes/static-pod-resources/configmaps/serviceaccount-ca/ca-bundle.crt`.

Remediation:

None.

Default Value:




By default, OpenShift sets the Kubernetes Controller Manager `root-ca-file` to `/etc/kubernetes/static-pod-resources/configmaps/serviceaccount-ca/ca-bundle.crt`.

Certificates for OpenShift platform components are automatically created and rotated by the OpenShift Container Platform.

References:

1. <https://docs.openshift.com/container-platform/4.13/operators/operator-reference.html>
2. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#cluster-openshift-controller-manager-operator_cluster-operators-ref
3. https://docs.openshift.com/container-platform/4.13/security/certificate_types_descriptions/control-plane-certificates.html
4. <https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/bootkube/bootstrap-manifests/kube-controller-manager-pod.yaml>
5. <https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/bootkube/config/bootstrap-config-overrides.yaml>
6. <https://github.com/openshift/cluster-kube-controller-manager-operator/blob/release-4.5/bindata/v4.1.0/kube-controller-manager/kubeconfig-cm.yaml>
7. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-controller-manager/>
8. <https://github.com/kubernetes/kubernetes/issues/11000>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	13.10 Perform Application Layer Filtering Perform application layer filtering. Example implementations include a filtering proxy, application layer firewall, or gateway.			
v7	4.4 Use Unique Passwords Where multi-factor authentication is not supported (such as local administrator, root, or service accounts), accounts will use passwords that are unique to that system.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

1.4 Scheduler

This section contains recommendations relating to Scheduler configuration flags.

In OpenShift 4, the Scheduler is managed with the Kubernetes Scheduler Operator. The Kubernetes Scheduler Operator manages and updates the Kubernetes Scheduler deployed on top of OpenShift Container Platform. The operator is installed with the Cluster Version Operator (CVO). The Kubernetes Scheduler Operator contains the following components:

- Operator
- Bootstrap manifest renderer
- Installer based on static pods
- Configuration observer

By default, the Operator exposes Prometheus metrics through the metrics service

1.4.1 Ensure that the healthz endpoints for the scheduler are protected by RBAC (Manual)

Profile Applicability:

- Level 1

Description:

Disable profiling, if not needed.

Rationale:

Profiling allows for the identification of specific performance bottlenecks. It generates a significant amount of program data that could potentially be exploited to uncover system and program details. If you are not experiencing any bottlenecks and do not need the profiler for troubleshooting purposes, it is recommended to turn it off to reduce the potential attack surface.

Impact:

Profiling information would not be available.

Audit:

In OpenShift 4, The Kubernetes Scheduler operator manages and updates the Kubernetes Scheduler deployed on top of OpenShift. By default, the operator exposes metrics via metrics service. The metrics are collected from the Kubernetes Scheduler operator. Profiling data is sent to **healthzPort**, the port of the localhost **healthz** endpoint. Changing this value may disrupt components that monitor the kubelet health. The default **healthz port** value is **10251**, and the **healthz bindAddress** is **127.0.0.1**.

To ensure the collected data is not exploited, profiling endpoints are secured via RBAC (see cluster-debugger role). By default, the profiling endpoints are accessible only by users bound to **cluster-admin** or **cluster-debugger** role. Profiling can not be disabled.

To verify the configuration, run the following command:

Run the following command to check the **livenessProbe** configuration:

```
oc -n openshift-kube-scheduler get cm kube-scheduler-pod -o json | jq -r '.data."pod.yaml"' | jq '.spec.containers[].livenessProbe'
```

Verify the output path is set to **healthz**.

Run the following command to check the **readinessProbe** configuration:

```
oc -n openshift-kube-scheduler get cm kube-scheduler-pod -o json | jq -r '.data."pod.yaml"' | jq '.spec.containers[].readinessProbe'
```

Verify the output path is set to **healthz**.

Verify endpoints exist for the scheduler:

```
oc -n openshift-kube-scheduler describe endpoints
```

Validate that RBAC is enabled and protects controller endpoints. First, switch to the **openshift-kube-scheduler**:

```
oc project openshift-kube-scheduler
```

Next, get the schedule pod name and port:

```
export POD=$(oc get pods -l app=openshift-kube-scheduler -o  
jsonpath='{.items[0].metadata.name}')  
export PORT=$(oc get pod $POD -o  
jsonpath='{.spec.containers[0].livenessProbe.httpGet.port}')
```

Attempt to make an insecure **GET** request to the metrics endpoint:

```
oc rsh $POD curl https://localhost:$PORT/metrics -k
```

Verify that an **HTTP 403** is returned.

Create a test service account:

```
oc create sa permission-test-sa
```

Generate a service account token and attempt to access the metrics endpoint:

```
export SA_TOKEN=$(oc create token permission-test-sa)  
oc rsh $POD curl http://localhost:$PORT/metrics -H "Authorization: Bearer  
$SA_TOKEN" -k
```

Verify that an **HTTP 403** is returned.

Login as a cluster administrator and attempt to access the metrics endpoint:

```
CLUSTER_ADMIN_TOKEN=$(oc whoami -t)  
oc rsh $POD curl https://localhost:$PORT/metrics -H "Authorization: Bearer  
$CLUSTER_ADMIN_TOKEN" -k
```

Verify metrics output is returned. Unset environment variables used in the test and delete the test service account:

```
unset CLUSTER_ADMIN_TOKEN POD PORT SA_TOKEN  
oc delete sa permission-test-sa
```

Remediation:

None.

Default Value:

By default, profiling is enabled and protected by RBAC.

References:

1. <https://github.com/openshift/cluster-kube-scheduler-operator>

2. <https://github.com/openshift/cluster-kube-scheduler-operator/blob/release-4.5/bindata/v4.1.0/kube-scheduler/svc.yaml>
3. <https://github.com/openshift/cluster-kube-scheduler-operator/blob/release-4.5/bindata/v4.1.0/kube-scheduler/pod.yaml>
4. <https://github.com/openshift/cluster-kube-scheduler-operator/blob/release-4.5/bindata/v4.1.0/kube-scheduler/pod.yaml#L32-L37>
5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-scheduler/>
6. <https://github.com/kubernetes/community/blob/master/contributors/devel/profiling.md>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	6.8 Define and Maintain Role-Based Access Control Define and maintain role-based access control, through determining and documenting the access rights necessary for each role within the enterprise to successfully carry out its assigned duties. Perform access control reviews of enterprise assets to validate that all privileges are authorized, on a recurring schedule at a minimum annually, or more frequently.			●
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.	●	●	●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1609	TA0002	M1035

1.4.2 Verify that the scheduler API service is protected by RBAC (Manual)

Profile Applicability:

- Level 1

Description:

Do not bind the scheduler service to non-loopback insecure addresses.

Rationale:

The Scheduler API service which runs on port 10251/TCP by default is used for health and metrics information and is available without authentication or encryption. As such it should only be bound to a localhost interface, to minimize the cluster's attack surface

Impact:

None.

Audit:

In OpenShift 4, The Kubernetes Scheduler operator manages and updates the Kubernetes Scheduler deployed on top of OpenShift. By default, the operator exposes metrics via metrics service. The metrics are collected from the Kubernetes Scheduler operator. Profiling data is sent to **healthzPort**, the port of the localhost **healthz** endpoint. Changing this value may disrupt components that monitor the kubelet health. The default **healthz port** value is **10251**, and the **healthz bindAddress** is **127.0.0.1**

To ensure the collected data is not exploited, profiling endpoints are secured via RBAC (see cluster-debugger role). By default, the profiling endpoints are accessible only by users bound to **cluster-admin** or **cluster-debugger** role. Profiling can not be disabled.

The bind-address argument is not used. Both authentication and authorization are in place.

Run the following command to verify the schedule endpoints:

```
oc -n openshift-kube-scheduler describe endpoints
```

Verify the **bind-address** and **port** arguments are not used:

```
oc -n openshift-kube-scheduler get cm kube-scheduler-pod -o json | jq -r '.data."pod.yaml"' | jq '.spec.containers[]|select(.name=="kube-scheduler")|.args'
```

Verify the metrics endpoint is protected by RBAC.

First, find the schedule pod information:

```
oc project openshift-kube-scheduler
export POD=$(oc get pods -l app=openshift-kube-scheduler -o
jsonpath='{.items[0].metadata.name}')
export POD_IP=$(oc get pods -l app=openshift-kube-scheduler -o
jsonpath='{.items[0].status.podIP}')
export PORT=$(oc get pod $POD -o
jsonpath='{.spec.containers[0].livenessProbe.httpGet.port}')
```

Attempt to make an insecure **GET** request to the metrics endpoint:

```
oc rsh $POD curl https://$POD_IP:$PORT/metrics -k
```

Ensure an **HTTP 403** is returned.

Create a test service account:

```
oc create sa permission-test-sa
```

Generate a service account token and attempt to access the metrics endpoint:

```
export SA_TOKEN=$(oc create token permission-test-sa)
oc rsh $POD curl https://$POD_IP:$PORT/metrics -H "Authorization: Bearer
$SA_TOKEN" -k
```

Verify that an **HTTP 403** is returned.

Login as a cluster administrator and attempt to access the metrics endpoint:

```
export CLUSTER_ADMIN_TOKEN=$(oc whoami -t)
oc rsh $POD curl https://$POD_IP:$PORT/metrics -H "Authorization: Bearer
$CLUSTER_ADMIN_TOKEN" -k
```

Verify metrics output is returned. Unset environment variables used in the test and delete the test service account:

```
unset CLUSTER_ADMIN_TOKEN POD PORT SA_TOKEN POD_IP
oc delete sa permission-test-sa
```

Remediation:

None.

Default Value:





By default, the **--bind-address** parameter is not used and the metrics endpoint is protected by RBAC when using the pod IP address.

References:

1. <https://github.com/openshift/cluster-kube-scheduler-operator>
2. <https://github.com/openshift/cluster-kube-scheduler-operator/blob/release-4.5/bindata/v4.1.0/kube-scheduler/svc.yaml>
3. <https://github.com/openshift/cluster-kube-scheduler-operator/blob/release-4.5/bindata/v4.1.0/kube-scheduler/pod.yaml>
4. <https://github.com/openshift/cluster-kube-scheduler-operator/blob/release-4.5/bindata/v4.1.0/kube-scheduler/pod.yaml#L32-L37>

5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-scheduler/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	9.3 Maintain and Enforce Network-Based URL Filters Enforce and update network-based URL filters to limit an enterprise asset from connecting to potentially malicious or unapproved websites. Example implementations include category-based filtering, reputation-based filtering, or through the use of block lists. Enforce filters for all enterprise assets.			
v7	9.2 Ensure Only Approved Ports, Protocols and Services Are Running Ensure that only network ports, protocols, and services listening on a system with validated business needs, are running on each system.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1106	TA0002	M1035

2 etcd

This section covers recommendations for etcd configuration. The OpenShift cluster-etcd-operator (CEO) is an operator that handles the scaling of etcd during cluster bootstrap and regular operation. The operator also manages provisioning etcd dependencies such as TLS certificates.

OpenShift uses X.509 certificates to provide secure communication to etcd. OpenShift generates these files and sets the arguments appropriately. etcd certificates are used for encrypted communication between etcd member peers, as well as encrypted client traffic. The following certificates are generated and used by etcd and other processes that communicate with etcd:

- Peer certificates: Used for communication between etcd members.
- Client certificates: Used for encrypted server-client communication. Client certificates are currently used by the API server only, and no other service should connect to etcd directly except for the proxy. Client secrets (etcd-client, etcd-metric-client, etcd-metric-signer, and etcd-signer) are added to the openshift-config, openshift-monitoring, and openshift-kube-apiserver namespaces.
- Server certificates: Used by the etcd server for authenticating client requests.
- Metric certificates: All metric consumers connect to proxy with metric-client certificates.

2.1 Ensure that the `--cert-file` and `--key-file` arguments are set as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

Configure TLS encryption for the `etcd` service.

Rationale:

`etcd` is a highly-available key value store used by Kubernetes deployments for persistent storage of all of its REST API objects. These objects are sensitive in nature and should be encrypted in transit.

Impact:

Client connections only over TLS would be served.

Audit:

OpenShift uses X.509 certificates to provide secure communication to `etcd`. OpenShift generates these files and sets the arguments appropriately. OpenShift does not use the `etcd-certfile` or `etcd-keyfile` flags.

Keys and certificates for control plane components like `kube-apiserver`, `kube-controller-manager`, `kube-scheduler` and `etcd` are stored with their respective static pod configurations in the directory `/etc/kubernetes/static-pod-resources/*/secrets`.

Run the following command to check the value of the `--cert-file` parameter on all applicable nodes:

```
for i in $(oc get pods -oname -n openshift-etcd)
do
    oc exec -n openshift-etcd -c etcd $i -- \
        ps -o command= -C etcd | sed 's/.*\(--cert-file=[^ ]*\).*/\1/'
done
```

Run the following command to check the value of the `--key-file` parameter on all applicable nodes:

```
for i in $(oc get pods -oname -n openshift-etcd)
do
    oc exec -n openshift-etcd -c etcd $i -- \
        ps -o command= -C etcd | sed 's/.*\(--key-file=[^ ]*\).*/\1/'
done
```

Verify that cert-file and key-file values are returned for each etcd member. `--cert-file=/etc/kubernetes/static-pod-certs/secrets/etcd-all-serving/etcd-serving- $\{ETCD_DNS_NAME\}$.cert --key-file=/etc/kubernetes/static-pod-certs/secrets/etcd-all-serving/etcd-serving- $\{ETCD_DNS_NAME\}$.key`

For example:

`--cert-file=/etc/kubernetes/static-pod-certs/secrets/etcd-all-serving/etcd-serving-ip-10-0-165-75.us-east-2.compute.internal.crt --key-file=/etc/kubernetes/static-pod-certs/secrets/etcd-all-serving/etcd-serving-ip-10-0-165-75.us-east-2.compute.internal.key`

Remediation:

OpenShift does not use the `etcd-certfile` or `etcd-keyfile` flags. Certificates for `etcd` are managed by the `etcd` cluster operator.



Default Value:



By default, `etcd` communication is secured with X.509 certificates.

References:

1. https://docs.openshift.com/container-platform/latest/security/certificate_types_descriptions/etcd-certificates.html
2. <https://github.com/openshift/cluster-etcd-operator>
3. <https://github.com/openshift/cluster-etcd-operator/blob/master/bindata/etcd/pod.yaml#L154-L167>
4. <https://etcd.io/>
5. <https://kubernetes.io/docs/tasks/administer-cluster/configure-upgrade-etcd/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.11 <u>Encrypt Sensitive Data at Rest</u> Encrypt sensitive data at rest on servers, applications, and databases containing sensitive data. Storage-layer encryption, also known as server-side encryption, meets the minimum requirement of this Safeguard. Additional encryption methods may include application-layer encryption, also known as client-side encryption, where access to the data storage device(s) does not permit access to the plain-text data.			

Controls Version	Control	IG 1	IG 2	IG 3
v7	14.4 <u>Encrypt All Sensitive Information in Transit</u> Encrypt all sensitive information in transit.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

2.2 Ensure that the `--client-cert-auth` argument is set to `true` (Manual)

Profile Applicability:

- Level 1

Description:

Enable client authentication on etcd service.

Rationale:

etcd is a highly-available key value store used by Kubernetes deployments for persistent storage of all of its REST API objects. These objects are sensitive in nature and should not be available to unauthenticated clients. You should enable the client authentication via valid certificates to secure the access to the etcd service.

Impact:

All clients attempting to access the etcd server will require a valid client certificate.

Audit:

OpenShift uses X.509 certificates to provide secure communication to etcd. OpenShift installation generates these files and sets the arguments appropriately. The following certificates are generated and used by etcd and other processes that communicate with etcd:

- Client certificates: Client certificates are currently used by the API server only, and no other service should connect to etcd directly except for the proxy. Client secrets (`etcd-client`, `etcd-metric-client`, `etcd-metric-signer`, and `etcd-signer`) are added to the `openshift-config`, `openshift-monitoring`, and `openshift-kube-apiserver` namespaces.
- Server certificates: Used by the etcd server for authenticating client requests.

Run the following command on the etcd server node:

```
for i in $(oc get pods -oname -n openshift-etcd)
do
    oc exec -n openshift-etcd -c etcd $i -- \
        ps -o command= -C etcd | sed 's/.*\(--client-cert-auth=[^ ]*\)\.*/\1/'
done
```

Verify that the `--client-cert-auth` argument is set to `true` for each etcd member.

Remediation:

This setting is managed by the cluster etcd operator. No remediation required.




Default Value:

By default, `client-cert-auth` is set to `true`.

References:

1. https://docs.openshift.com/container-platform/latest/security/certificate_types_descriptions/etcd-certificates.html
2. <https://github.com/openshift/cluster-etcd-operator>
3. <https://github.com/openshift/cluster-etcd-operator/blob/release-4.5/bindata/etcd/pod.yaml#L154-L167>
4. <https://github.com/openshift/cluster-etcd-operator/blob/master/bindata/etcd/pod.yaml#L154-L167>
5. <https://etcd.io/>
6. <https://kubernetes.io/docs/tasks/administer-cluster/configure-upgrade-etcd/>
7. <https://etcd.io/#client-cert-auth>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.11 <u>Encrypt Sensitive Data at Rest</u> Encrypt sensitive data at rest on servers, applications, and databases containing sensitive data. Storage-layer encryption, also known as server-side encryption, meets the minimum requirement of this Safeguard. Additional encryption methods may include application-layer encryption, also known as client-side encryption, where access to the data storage device(s) does not permit access to the plain-text data.			
v7	14.8 <u>Encrypt Sensitive Information at Rest</u> Encrypt all sensitive information at rest using a tool that requires a secondary authentication mechanism not integrated into the operating system, in order to access the information.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

2.3 Ensure that the `--auto-tls` argument is not set to true (Manual)

Profile Applicability:

- Level 1

Description:

Do not use self-signed certificates for TLS.

Rationale:

etcd is a highly-available key value store used by Kubernetes deployments for persistent storage of all of its REST API objects. These objects are sensitive in nature and should not be available to unauthenticated clients. You should enable the client authentication via valid certificates to secure the access to the etcd service.

Impact:

Clients will not be able to use self-signed certificates for TLS.

Audit:

OpenShift configures etcd with secure communication. Openshift installs etcd as static pods on control plane nodes, and mounts the configuration files from `/etc/etcd/` on the host. The `etcd.conf` file includes `auto-tls` configurations as referenced in `/etc/etcd/etcd.conf`.

OpenShift 4 includes multiple CAs providing independent chains of trust, which ensure that a platform CA will never accidentally sign a certificate that can be used for the wrong purpose, increasing the security posture of the cluster.

These internal self-signing CAs enable automation because the key is known to the cluster. The certificates generated by each CA are used to identify a particular OpenShift platform component to another OpenShift platform component. The OpenShift CAs are managed by the cluster and are only used within the cluster.

- Each cluster CA can only issue certificates for its own purpose within its own cluster.
- CAs for one OpenShift cluster cannot influence CAs for a different OpenShift cluster, thus avoiding cross-cluster interference.
- Cluster CAs cannot be influenced by an external CA that the cluster does not control.

Run the following command to verify if auto-TLS is enabled:

```
for i in $(oc get pods -oname -n openshift-etcd)
do
    oc exec -n openshift-etcd -c etcd $i -- \
        ps -o command= -C etcd | grep -- '--auto-tls=true 2>&1>/dev/null' ; \
        echo $?
done
```

Verify that 1 is returned for each etcd member.

Remediation:

This setting is managed by the cluster etcd operator. No remediation required.

Default Value:

By default, OpenShift configures etcd to use a cluster CA which creates self-signed certificates. These internal self-signing CAs enable automation because the key is known to the cluster. The certificates generated by each CA are used to identify a particular OpenShift platform component to another OpenShift platform component. The OpenShift CAs are managed by the cluster and are only used within the cluster.





- Each cluster CA can only issue certificates for its own purpose within its own cluster.
- CAs for one OpenShift cluster cannot influence CAs for a different OpenShift cluster, thus avoiding cross-cluster interference.
- Cluster CAs cannot be influenced by an external CA that the cluster does not control.

This configuration cannot be changed.

References:

1. https://docs.openshift.com/container-platform/latest/security/certificate_types_descriptions/etcd-certificates.html
2. <https://github.com/openshift/cluster-etcd-operator>
3. <https://github.com/openshift/cluster-etcd-operator/blob/release-4.5/bindata/etcd/pod.yaml#L154-L167>
4. <https://github.com/openshift/cluster-etcd-operator/blob/master/bindata/etcd/pod.yaml#L154-L167>
5. <https://etcd.io/>
6. <https://kubernetes.io/docs/tasks/administer-cluster/configure-upgrade-etcd/>
7. <https://etcd.io/#auto-tls>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.10 <u>Encrypt Sensitive Data in Transit</u> Encrypt sensitive data in transit. Example implementations can include: Transport Layer Security (TLS) and Open Secure Shell (OpenSSH).			
v7	14.4 <u>Encrypt All Sensitive Information in Transit</u> Encrypt all sensitive information in transit.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

2.4 Ensure that the `--peer-cert-file` and `--peer-key-file` arguments are set as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

etcd should be configured to make use of TLS encryption for peer connections.

Rationale:

etcd is a highly-available key value store used by Kubernetes deployments for persistent storage of all of its REST API objects. These objects are sensitive in nature and should be encrypted in transit and also amongst peers in the etcd clusters.

Impact:

etcd cluster peers are set up TLS for their communication.

Audit:

OpenShift uses X.509 certificates to provide secure communication to etcd. OpenShift generates these files and sets the arguments appropriately. etcd certificates are used for encrypted communication between etcd member peers, as well as encrypted client traffic. Peer certificates are generated and used for communication between etcd members.

OpenShift installs etcd as static pods on control plane nodes, and mounts the configuration files from `/etc/etcd/` on the host. The `etcd.conf` file includes `peer-cert-file` and `peer-key-file` configurations as referenced in `/etc/etcd/etcd.conf`.

Run the following command to check the value of `--peer-cert-file`:

```
for i in $(oc get pods -oname -n openshift-etcd)
do
    oc exec -n openshift-etcd -c etcd $i -- \
        ps -o command= -C etcd | sed 's/.*\(--peer-cert-file=[^ ]*\).*/\1/'
done
```

Run the following command to check the value of `--peer-key-file`:

```
for i in $(oc get pods -oname -n openshift-etcd)
do
    oc exec -n openshift-etcd -c etcd $i -- \
        ps -o command= -C etcd | sed 's/.*\(--peer-key-file=[^ ]*\).*/\1/'
done
```

Verify that the following is returned for each etcd member. `--peer-cert-file=/etc/kubernetes/static-pod-certs/secrets/etcd-all-peer/etcd-peer-${ETCD_DNS_NAME}.cert --peer-key-file=/etc/kubernetes/static-pod-certs/secrets/etcd-all-peer/etcd-peer-${ETCD_DNS_NAME}.key`

For example `--peer-cert-file=/etc/kubernetes/static-pod-certs/secrets/etcd-all-peer/etcd-peer-ip-10-0-158-52.us-east-2.compute.internal.cert --peer-key-file=/etc/kubernetes/static-pod-certs/secrets/etcd-all-peer/etcd-peer-ip-10-0-158-52.us-east-2.compute.internal.key`

Remediation:

None. This configuration is managed by the etcd operator.





Default Value:

By default, peer communication over TLS is configured.

References:

1. https://docs.openshift.com/container-platform/latest/security/certificate_types_descriptions/etcd-certificates.html
2. <https://github.com/openshift/cluster-etcd-operator>
3. <https://github.com/openshift/cluster-etcd-operator/blob/release-4.5/bindata/etcd/pod.yaml#L154-L167>
4. <https://github.com/openshift/cluster-etcd-operator/blob/master/bindata/etcd/pod.yaml#L154-L167>
5. <https://etcd.io/>
6. <https://kubernetes.io/docs/tasks/administer-cluster/configure-upgrade-etcd/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.10 Encrypt Sensitive Data in Transit Encrypt sensitive data in transit. Example implementations can include: Transport Layer Security (TLS) and Open Secure Shell (OpenSSH).			
v7	14.4 Encrypt All Sensitive Information in Transit Encrypt all sensitive information in transit.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

2.5 Ensure that the `--peer-client-cert-auth` argument is set to `true` (Manual)

Profile Applicability:

- Level 1

Description:

etcd should be configured for peer authentication.

Rationale:

etcd is a highly-available key value store used by Kubernetes deployments for persistent storage of all of its REST API objects. These objects are sensitive in nature and should be accessible only by authenticated etcd peers in the etcd cluster.

Impact:

All peers attempting to communicate with the etcd server require a valid client certificate for authentication.

Audit:

OpenShift uses X.509 certificates to provide secure communication to etcd. OpenShift generates these files and sets the arguments appropriately. etcd certificates are used for encrypted communication between etcd member peers, as well as encrypted client traffic. Peer certificates are generated and used for communication between etcd members.

OpenShift installs etcd as static pods on control plane nodes, and mounts the configuration files from `/etc/etcd/` on the host. The `etcd.conf` file includes `peer-client-cert-auth` configurations as referenced in `/etc/etcd/etcd.conf`.

Run the following command:

```
for i in $(oc get pods -oname -n openshift-etcd)
do
    oc exec -n openshift-etcd -c etcd $i -- \
        ps -o command= -C etcd | sed 's/.*\(--peer-client-cert-auth=[^
]*\)*/\1/'
done
```

Verify that the `--peer-client-cert-auth` argument is set to `true` for each etcd member.

Remediation:

This setting is managed by the cluster etcd operator. No remediation required.





Default Value:

By default, `--peer-client-cert-auth` argument is set to `true`.

References:

1. https://docs.openshift.com/container-platform/latest/security/certificate_types_descriptions/etcd-certificates.html
2. <https://github.com/openshift/cluster-etcd-operator>
3. <https://github.com/openshift/cluster-etcd-operator/blob/release-4.5/bindata/etcd/pod.yaml#L154-L167>
4. <https://github.com/openshift/cluster-etcd-operator/blob/release-4.5/bindata/etcd/pod.yaml#L154-L167>
5. <https://etcd.io/>
6. <https://kubernetes.io/docs/tasks/administer-cluster/configure-upgrade-etcd/>
7. <https://etcd.io/#peer-client-cert-auth>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	6.8 Define and Maintain Role-Based Access Control Define and maintain role-based access control, through determining and documenting the access rights necessary for each role within the enterprise to successfully carry out its assigned duties. Perform access control reviews of enterprise assets to validate that all privileges are authorized, on a recurring schedule at a minimum annually, or more frequently.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

2.6 Ensure that the `--peer-auto-tls` argument is not set to true (Manual)

Profile Applicability:

- Level 1

Description:

Do not use automatically generated self-signed certificates for TLS connections between peers.

Rationale:

etcd is a highly-available key value store used by Kubernetes deployments for persistent storage of all of its REST API objects. These objects are sensitive in nature and should be accessible only by authenticated etcd peers in the etcd cluster. Hence, do not use self-signed certificates for authentication.

Impact:

All peers attempting to communicate with the etcd server require a valid client certificate for authentication.

Audit:

OpenShift does not use the `--peer-auto-tls` argument. OpenShift 4 includes multiple CAs providing independent chains of trust, which ensure that a platform CA will never accidentally sign a certificate that can be used for the wrong purpose, increasing the security posture of the cluster.

These internal self-signing CAs enable automation because the key is known to the cluster. The certificates generated by each CA are used to identify a particular OpenShift platform component to another OpenShift platform component. The OpenShift CAs are managed by the cluster and are only used within the cluster. This means that

- Each cluster CA can only issue certificates for its own purpose within its own cluster.
- CAs for one OpenShift cluster cannot influence CAs for a different OpenShift cluster, thus avoiding cross-cluster interference.

Cluster CAs cannot be influenced by an external CA that the cluster does not control.

Run the following command to check the value of `--peer-auto-tls`:

```
for i in $(oc get pods -oname -n openshift-etcd)
do
    oc exec -n openshift-etcd -c etcd $i -- \
        ps -o command= -C etcd | grep -- --peer-auto-tls=true 2>&1>/dev/null ;
\
    echo $?
done
```

Verify that **1** is returned for each etcd member.

Remediation:

This setting is managed by the cluster etcd operator. No remediation required.

Default Value:

OpenShift does not use the **--peer-auto-tls** argument. By default, OpenShift configures etcd to use a cluster CA which creates self-signed certificates. These internal self-signing CAs enable automation because the key is known to the cluster. The certificates generated by each CA are used to identify a particular OpenShift platform component to another OpenShift platform component. The OpenShift CAs are managed by the cluster and are only used within the cluster. This means that






- Each cluster CA can only issue certificates for its own purpose within its own cluster.
- CAs for one OpenShift cluster cannot influence CAs for a different OpenShift cluster, thus avoiding cross-cluster interference.
- Cluster CAs cannot be influenced by an external CA that the cluster does not control.

This configuration cannot be changed.

References:

1. https://docs.openshift.com/container-platform/latest/security/certificate_types_descriptions/etcd-certificates.html
2. <https://github.com/openshift/cluster-etcd-operator>
3. <https://github.com/openshift/cluster-etcd-operator/blob/release-4.5/bindata/etcd/pod.yaml#L154-L167>
4. <https://github.com/openshift/cluster-etcd-operator/blob/master/bindata/etcd/pod.yaml#L154-L167>
5. <https://etcd.io/>
6. <https://etcd.io/#peer-auto-tls>
7. <https://etcd.io/#peer-auto-tls>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.2 <u>Use Unique Passwords</u> Use unique passwords for all enterprise assets. Best practice implementation includes, at a minimum, an 8-character password for accounts using MFA and a 14-character password for accounts not using MFA.			
v7	4.4 <u>Use Unique Passwords</u> Where multi-factor authentication is not supported (such as local administrator, root, or service accounts), accounts will use passwords that are unique to that system.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

2.7 Ensure that a unique Certificate Authority is used for etcd (Manual)

Profile Applicability:

- Level 2

Description:

Use a different certificate authority for etcd from the one used for Kubernetes.

Rationale:

etcd is a highly available key-value store used by Kubernetes deployments for persistent storage of all of its REST API objects. Its access should be restricted to specifically designated clients and peers only.

Authentication to etcd is based on whether the certificate presented was issued by a trusted certificate authority. There is no checking of certificate attributes such as common name or subject alternative name. As such, if any attackers were able to gain access to any certificate issued by the trusted certificate authority, they would be able to gain full access to the etcd database.

Impact:

Additional management of the certificates and keys for the dedicated certificate authority will be required.

Audit:

OpenShift 4 includes multiple CAs providing independent chains of trust, which ensure that a platform CA will never accidentally sign a certificate that can be used for the wrong purpose, increasing the security posture of the cluster. OpenShift uses a separate CA for etcd.

These internal self-signing CAs enable automation because the key is known to the cluster. The certificates generated by each CA are used to identify a particular OpenShift platform component to another OpenShift platform component. The OpenShift CAs are managed by the cluster and are only used within the cluster. This means that

- Each cluster CA can only issue certificates for its own purpose within its own cluster.
- CAs for one OpenShift cluster cannot influence CAs for a different OpenShift cluster, thus avoiding cross-cluster interference.

Cluster CAs cannot be influenced by an external CA that the cluster does not control.

Run the following command to check the value of `--trusted-ca-file`:

```
for i in $(oc get pods -oname -n openshift-etcd)
do
    oc exec -n openshift-etcd -c etcd $i -- \
        ps -o command= -C etcd | sed 's/.*\(--trusted-ca-file=[^ ]*\).*/\1/'
done
```

Run the following command to check the value of **--peer-trusted-ca-file**:

```
for i in $(oc get pods -oname -n openshift-etcd)
do
    oc exec -n openshift-etcd -c etcd $i -- \
        ps -o command= -C etcd | sed 's/.*\(--peer-trusted-ca-file=[^ ]*\).*/\1/'
done
```

Verify that **--trusted-ca-file=/etc/kubernetes/static-pod-certs/configmaps/etcd-serving-ca/ca-bundle.crt** and **--peer-trusted-ca-file=/etc/kubernetes/static-pod-certs/configmaps/etcd-peer-client-ca/ca-bundle.crt** are returned for each member.

Remediation:

None required. Certificates for etcd are managed by the OpenShift cluster etcd operator.



Default Value:

By default, in OpenShift 4, communication with etcd is secured by the etcd serving CA.

References:

1. https://docs.openshift.com/container-platform/latest/security/certificate_types_descriptions/etcd-certificates.html
2. <https://github.com/openshift/cluster-etcd-operator>
3. <https://github.com/openshift/cluster-etcd-operator/blob/release-4.5/bindata/etcd/pod.yaml#L154-L167>
4. <https://github.com/openshift/cluster-etcd-operator/blob/master/bindata/etcd/pod.yaml#L154-L167>
5. <https://etcd.io/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.11 <u>Encrypt Sensitive Data at Rest</u> Encrypt sensitive data at rest on servers, applications, and databases containing sensitive data. Storage-layer encryption, also known as server-side encryption, meets the minimum requirement of this Safeguard. Additional encryption methods may include application-layer encryption, also known as client-side encryption, where access to the data storage device(s) does not permit access to the plain-text data.			

Controls Version	Control	IG 1	IG 2	IG 3
v7	14.8 Encrypt Sensitive Information at Rest Encrypt all sensitive information at rest using a tool that requires a secondary authentication mechanism not integrated into the operating system, in order to access the information.			●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1022

3 Control Plane Configuration

This section contains recommendations for cluster-wide areas, such as authentication and logging. Unlike section 1 these recommendations should apply to all deployments.

3.1 Authentication and Authorization

3.1.1 Client certificate authentication should not be used for users (Manual)

Profile Applicability:

- Level 2

Description:

Kubernetes provides the option to use client certificates for user authentication. However as there is no way to revoke these certificates when a user leaves an organization or loses their credential, they are not suitable for this purpose.

It is not possible to fully disable client certificate use within a cluster as it is used for component to component authentication.

Rationale:

With any authentication mechanism the ability to revoke credentials if they are compromised or no longer required, is a key control. Kubernetes client certificate authentication does not allow for this due to a lack of support for certificate revocation.

Impact:

External mechanisms for authentication generally require additional software to be deployed.

Audit:

For users to interact with OpenShift Container Platform, they must first authenticate to the cluster. The authentication layer identifies the user with requests to the OpenShift Container Platform API. The authorization layer then uses information about the requesting user to determine if the request is allowed. [Understanding authentication | Authentication | OpenShift Container Platform](#)

The OpenShift Container Platform includes a built-in OAuth server for token-based authentication. Developers and administrators obtain OAuth access tokens to authenticate themselves to the API. It is recommended for an administrator to configure OAuth to specify an identity provider after the cluster is installed. User access to the cluster is managed through the identity provider. [Understanding identity provider configuration | Authentication | OpenShift Container Platform](#)

First, verify user authentication is enabled:

```
oc describe authentication
```

Next, verify an identity provider is configured:

```
oc get oauth -o json | jq '.items[].spec.identityProviders'
```

Verify at least one identity provider is configured, and verify that the **kubeadmin** user does not exist.

Next, verify a **cluster-admin** user exists:

```
oc get clusterrolebindings -o='custom-  
columns=NAME:.metadata.name,ROLE:.roleRef.name,SUBJECT:.subjects[*].kind' |  
grep cluster-admin | grep User
```

Verify at least one user has the **cluster-admin** role.

Finally, verify that the **kubeadmin** secret is removed.

```
oc get secrets kubeadmin -n kube-system
```

No result is expected.

Remediation:

Configure an identity provider for the OpenShift cluster following the OpenShift [documentation](#). Once an identity provider has been defined, you can use RBAC to define and apply permissions.

After you define an identity provider and create a new **cluster-admin** user you can reduce the attack surface by removing the default **kubeadmin** user.

Default Value:

By default, only a **kubeadmin** user exists on your cluster. To specify an identity provider, you must create a Custom Resource (CR) that describes that identity provider and add it to the cluster.

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/understanding-identity-provider.html>
2. <https://docs.openshift.com/container-platform/latest/authentication/using-rbac.html#authorization-overview-using-rbac>
3. <https://docs.openshift.com/container-platform/latest/authentication/remove-kubeadmin.html>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.8 Uninstall or Disable Unnecessary Services on Enterprise Assets and Software Uninstall or disable unnecessary services on enterprise assets and software, such as an unused file sharing service, web application module, or service function.		●	●

Controls Version	Control	IG 1	IG 2	IG 3
v8	<p>16.11 Leverage Vetted Modules or Services for Application Security Components</p> <p>Leverage vetted modules or services for application security components, such as identity management, encryption, and auditing and logging. Using platform features in critical security functions will reduce developers' workload and minimize the likelihood of design or implementation errors. Modern operating systems provide effective mechanisms for identification, authentication, and authorization and make those mechanisms available to applications. Use only standardized, currently accepted, and extensively reviewed encryption algorithms. Operating systems also provide mechanisms to create and maintain secure audit logs.</p>		●	●
v7	<p>16.2 Configure Centralized Point of Authentication</p> <p>Configure access for all accounts through as few centralized points of authentication as possible, including network, security, and cloud systems.</p>		●	●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078, T1098	TA0003, TA0006	M1027, M1032

3.2 Logging

3.2.1 *Ensure that a minimal audit policy is created (Manual)*

Profile Applicability:

- Level 1

Description:

Kubernetes can audit the details of requests made to the API server.

Rationale:

Logging is an important detective control for all systems, to detect potential unauthorized access.

Impact:

Audit logs will be created on the master nodes, which will consume disk space. Care should be taken to avoid generating too large volumes of log information as this could impact the available of the cluster nodes.

Audit:

In OpenShift, auditing of the API Server is on by default. Audit provides a security-relevant chronological set of records documenting the sequence of activities that have affected the system by individual users, administrators, or other components of the system. Audit works at the API server level, logging all requests coming to the server. Each audit log contains two entries:

The request line containing:

A Unique ID allowing to match the response line (see #2)

- The source IP of the request
- The HTTP method being invoked
- The original user invoking the operation
- The impersonated user for the operation (self meaning himself)
- The impersonated group for the operation (lookup meaning user's group)
- The namespace of the request or
- The URI as requested

The response line containing:

- The unique ID from #1
- The response code

You can view logs for the OpenShift Container Platform API server or the Kubernetes API server for each master node. Follow the steps in documentation. [Viewing the audit log](#)

Use the following command to view the audit log profile:

```
oc get apiserver cluster -o json | jq .spec.audit.profile
```

Verify the result is not **None**, which means audit logging is disabled.

Review the audit log configuration for the OpenShift and Kubernetes API servers using the following commands:

```
oc get cm -n openshift-apiservert config -o json | jq -r '.data."config.yaml"' | jq .apiServerArguments
oc get cm -n openshift-kube-apiserver config -o json | jq -r '.data."config.yaml"' | jq .apiServerArguments
```

Review the audit policies for the OpenShift and Kubernetes API servers using the following commands:

```
oc get cm -n openshift-apiservert audit -o json | jq -r '.data."policy.yaml"'
oc get cm -n openshift-kube-apiserver kube-apiserver-audit-policies -o json | jq -r '.data."policy.yaml"'
```

Verify the returned configuration and ensure it aligns with data retention and storage requirements for the deployment.

Use the following command to view Kubernetes API server audit logs:

```
oc adm node-logs --role=master --path=kube-apiserver/
```

Verify logs are returned.

Use the following command to view OpenShift API server audit logs.

```
oc adm node-logs --role=master --path=openshift-apiserver/
```

Verify logs are returned.

Remediation:

None.

Default Value:







Auditing logging is enabled by default, using the **Default** audit profile.

Please reference the OpenShift audit logging [documentation](#) for more information on various profiles and configuration guidance.

References:

1. <https://docs.openshift.com/container-platform/latest/security/audit-log-policy-config.html>
2. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/master/bindata/v4.1.0/config/defaultconfig.yaml#L17-L31>
3. <https://kubernetes.io/docs/tasks/debug-application-cluster/audit/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	8.2 <u>Collect Audit Logs</u> Collect audit logs. Ensure that logging, per the enterprise's audit log management process, has been enabled across enterprise assets.			
v7	6.2 <u>Activate audit logging</u> Ensure that local logging has been enabled on all systems and networking devices.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1543	TA0003, TA0004	M1026

3.2.2 Ensure that the audit policy covers key security concerns (Manual)

Profile Applicability:

- Level 2

Description:

Ensure that the audit policy created for the cluster covers key security concerns.

Rationale:

Security audit logs should cover access and modification of key resources in the cluster, to enable them to form an effective part of a security environment.

Impact:

Increasing audit logging will consume resources on the nodes or other log destinations.

Audit:

Review the audit policy provided for the cluster and ensure that it covers the following areas:

- The use of sensitive resources like **Secrets**, **ConfigMaps**, and **TokenReviews** are logged at the **Metadata** level
- Modifications to **pods** and **deployments** are logged at the **Request** level
- The use of **pods/exec**, **pods/portforward**, **pods/proxy**, and **services/proxy** are at least logged at the **Metadata** level

For most requests, minimally logging at the **Metadata** level is recommended (the most basic level of logging).

You can configure the audit feature to set log level, retention policy, and the type of events to log. You can set the log level settings for an overall component or the API server to one of the following. The setting can be different for each setting.

Use the following command to view the audit policies for the Kubernetes API server:

```
oc get configmap -n openshift-kube-apiserver kube-apiserver-audit-policies -o json | jq -r '.data."policy.yaml"'
```

Use the following command to view the audit policies for the OpenShift API server:

```
oc get configmap -n openshift-apiserver audit -o json | jq -r '.data."policy.yaml"'
```

Remediation:

Update the audit log policy profile to use **WriteRequestBodies**.




Default Value:

Audit logging is configured by default using the **Default** audit policy, but you are advised to review the log retention settings and log levels to align with your cluster's security posture.

References:

1. <https://docs.openshift.com/container-platform/latest/security/audit-log-policy-config.html>
2. <https://docs.openshift.com/container-platform/latest/security/audit-log-view.html>
3. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.5/bindata/v4.1.0/config/defaultconfig.yaml#L47-L77>
4. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.6/bindata/v4.1.0/config/defaultconfig.yaml#L34-L78>
5. <https://github.com/k8scop/k8s-security-dashboard/blob/master/configs/kubernetes/adv-audit.yaml>
6. <https://kubernetes.io/docs/tasks/debug/debug-cluster/audit/>
7. https://github.com/falcosecurity/falco/blob/master/examples/k8s_audit_config/audit-policy.yaml
8. <https://github.com/kubernetes/kubernetes/blob/master/cluster/gce/gci/configure-helper.sh#L735>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	8.5 Collect Detailed Audit Logs Configure detailed audit logging for enterprise assets containing sensitive data. Include event source, date, username, timestamp, source addresses, destination addresses, and other useful elements that could assist in a forensic investigation.			
v7	14.9 Enforce Detail Logging for Access or Changes to Sensitive Data Enforce detailed audit logging for access to sensitive data or changes to sensitive data (utilizing tools such as File Integrity Monitoring or Security Information and Event Monitoring).			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1543	TA0003, TA0004	M1026

4 Worker Nodes

This section consists of security recommendations for the components that run on Kubernetes worker nodes.

Note that these components may also run on Kubernetes master nodes, so the recommendations in this section should be applied to master nodes as well as worker nodes where the master nodes make use of these components.

4.1 Worker Node Configuration Files

This section covers recommendations for configuration files on the worker nodes. As the same files exist on the master nodes, the same commands should be run on all nodes.

In OpenShift 4, node configuration files are managed by the Machine Config Operator.

4.1.1 Ensure that the kubelet service file permissions are set to 644 or more restrictive (Automated)

Profile Applicability:

- Level 1

Description:

Ensure that the kubelet service file has permissions of 644 or more restrictive.

Rationale:

The kubelet service file controls various parameters that set the behavior of the kubelet service in the worker node. You should restrict its file permissions to maintain the integrity of the file. The file should be writable by only the administrators on the system.

Impact:

None

Audit:

Kubelet is run as a **systemd** unit and its configuration file is created with **644** permissions.

Run the following command:

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}')
do
    oc debug -q node/${node} -- chroot /host stat -c %a
    /etc/systemd/system/kubelet.service
done
```

OR

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}')
do
    oc debug -q node/${node} -- chroot /host stat -c %a
    /host/usr/lib/systemd/system/kubelet.service
done
```

Note: It is possible to change the **oc debug** to **oc debug -q** in the above 2 statements to suppress or quiet the debug output

Verify that the permissions are **644** or more restrictive.

Remediation:

None.







Default Value:

By default, the kubelet service file has permissions of 644.

References:

1. https://docs.openshift.com/container-platform/4.5/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. https://docs.openshift.com/container-platform/4.5/scalability_and_performance/recommended-host-practices.html#create-a-kubeletconfig-crd-to-edit-kubelet-parameters
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>
4. <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/#44-joining-your-nodes>
5. <https://kubernetes.io/docs/reference/setup-tools/kubeadm/#kubelet-drop-in>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

4.1.2 Ensure that the kubelet service file ownership is set to root:root (Automated)

Profile Applicability:

- Level 1

Description:

Ensure that the kubelet service file ownership is set to **root:root**.

Rationale:

The kubelet service file controls various parameters that set the behavior of the kubelet service in the worker node. You should set its file ownership to maintain the integrity of the file. The file should be owned by **root:root**.

Impact:

None

Audit:

Run the following command to list the user and group for the **kubelet.service** file on each node:

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}');  
do  
oc debug node/${node} -- chroot /host stat -c %U:%G  
/etc/systemd/system/kubelet.service  
done
```

OR

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}');  
do  
oc debug node/${node} -- chroot /host stat -c %U:%G  
/host/usr/lib/systemd/system/kubelet.service  
done
```

Verify that the ownership is set to **root:root**.

Remediation:

None.







Default Value:

By default, OpenShift sets the default user and group for the **kubelet.service** file to **root:root**.

References:

1. https://docs.openshift.com/container-platform/4.5/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. https://docs.openshift.com/container-platform/4.5/scalability_and_performance/recommended-host-practices.html#create-a-kubeletconfig-crd-to-edit-kubelet-parameters
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>
4. <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/create-cluster-kubeadm/#44-joining-your-nodes>
5. <https://kubernetes.io/docs/reference/setup-tools/kubeadm/#kubelet-drop-in>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 Ensure the Use of Dedicated Administrative Accounts Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

4.1.3 If proxy kube proxy configuration file exists ensure permissions are set to 644 or more restrictive (Manual)

Profile Applicability:

- Level 1

Description:

If **kube-proxy** is running, and if it is using a file-based configuration file, ensure that the file has permissions of **644** or more restrictive.

Rationale:

The **kube-proxy** configuration file controls various parameters of the **kube-proxy** service in the worker node. You should restrict its file permissions to maintain the integrity of the file. The file should be writable by only the administrators on the system.

It is possible to run **kube-proxy** with the kubeconfig parameters configured as a Kubernetes ConfigMap instead of a file. In this case, there is no proxy kubeconfig file.

Impact:

None.

Audit:

In OpenShift 4, the **kube-proxy** runs within the **sdn** pods, which copies the **kubeconfig** from a configmap to the container at **/config/kube-proxy-config.yaml**, with **644** permissions.

Run the following command:

```
for i in $(oc get pods -n openshift-sdn -l app=sdn -oname)
do
    oc exec -n openshift-sdn $i -- \
        stat -Lc %a /config/kube-proxy-config.yaml
done
```

Verify that the **kube-proxy-config.yaml** file has permissions of **644**.

Remediation:

None.







Default Value:

By default, **kube-proxy** config file has permissions of **644**.

References:

1. https://docs.openshift.com/container-platform/latest/networking/openshift_sdn/configuring-kube-proxy.html
2. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-proxy/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 <u>Configure Data Access Control Lists</u> Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 <u>Protect Information through Access Control Lists</u> Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

4.1.4 If proxy kubeconfig file exists ensure ownership is set to root:root (Manual)

Profile Applicability:

- Level 1

Description:

If **kube-proxy** is running, ensure that the file ownership of its kubeconfig file is set to **root:root**.

Rationale:

The kubeconfig file for **kube-proxy** controls various parameters for the **kube-proxy** service in the worker node. You should set its file ownership to maintain the integrity of the file. The file should be owned by **root:root**.

Impact:

None

Audit:

In OpenShift 4, the **kube-proxy** runs within the **sdn** pods, which copies the **kubeconfig** from a configmap to the container at **/tmp/kubeconfig**, with root:root ownership.

Run the following command:

```
for i in $(oc get pods -n openshift-sdn -l app=sdn -oname)
do
    oc exec -n openshift-sdn $i -- \
        stat -Lc %U:%G /config/kube-proxy-config.yaml
done
```

Verify that the **kube-proxy-config.yaml** file has ownership root:root.

Remediation:

None required. The configuration is managed by OpenShift operators.







Default Value:

By default, **proxy** file ownership is set to **root:root**.

References:

1. https://docs.openshift.com/container-platform/4.5/networking/openshift_sdn/configuring-kube-proxy.html
2. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-proxy/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 <u>Restrict Administrator Privileges to Dedicated Administrator Accounts</u> Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 <u>Ensure the Use of Dedicated Administrative Accounts</u> Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

4.1.5 Ensure that the --kubeconfig kubelet.conf file permissions are set to 644 or more restrictive (Automated)

Profile Applicability:

- Level 1

Description:

Ensure that the `kubelet.conf` file has permissions of `644` or more restrictive.

Rationale:

The `kubelet.conf` file is the kubeconfig file for the node, and controls various parameters that set the behavior and identity of the worker node. You should restrict its file permissions to maintain the integrity of the file. The file should be writable by only the administrators on the system.

Impact:

None

Audit:

Run the following command to check the permissions of the `kubelet.conf` on each node:

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}'); do
oc debug node/${node} -- chroot /host stat -c %a /etc/kubernetes/kubelet.conf
done
```

OR

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}'); do
oc debug node/${node} -- chroot /host stat -c %a
/host/etc/kubernetes/kubelet.conf
done
```

Verify that the permissions are `644`.

Remediation:

None.







Default Value:

By default, OpenShift sets the default permissions for the `kubelet.conf` to `644`.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 <u>Configure Data Access Control Lists</u> Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 <u>Protect Information through Access Control Lists</u> Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

4.1.6 Ensure that the --kubeconfig kubelet.conf file ownership is set to root:root (Automated)

Profile Applicability:

- Level 1

Description:

Ensure that the **kubelet.conf** file ownership is set to **root:root**.

Rationale:

The **kubelet.conf** file is the Kubernetes configuration file for the node, and controls various parameters that set the behavior and identity of the worker node. You should set its file ownership to maintain the integrity of the file. The file should be owned by **root:root**.

Impact:

None.

Audit:

Run the following command to view the user and group ownership of the **kubelet.conf** file:

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}');  
do  
oc debug node/${node} -- chroot /host stat -c %U:%G  
/etc/kubernetes/kubelet.conf  
done
```

OR

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}');  
do  
oc debug node/${node} -- chroot /host stat -c %U:%G  
/host/etc/kubernetes/kubelet.conf  
done
```

Verify that the ownership is set to **root:root**.

Remediation:

None.







Default Value:

By default, **kubelet.conf** file ownership is set to **root:root**.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 Ensure the Use of Dedicated Administrative Accounts Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

4.1.7 Ensure that the certificate authorities file permissions are set to 644 or more restrictive (Automated)

Profile Applicability:

- Level 1

Description:

Ensure that the certificate authorities file has permissions of **644** or more restrictive.

Rationale:

The certificate authorities file controls the authorities used to validate API requests. You should restrict its file permissions to maintain the integrity of the file. The file should be writable by only the administrators on the system.

Impact:

None

Audit:

Use the following command to check the **clientCAFile** for each node in the cluster:

```
for node in $(oc get nodes -ojsonpath='{.items[*].metadata.name}');  
do  
oc get --raw /api/v1/nodes/$node/proxy/configz | jq  
'|.kubeletconfig.authentication.x509.clientCAFile'  
done
```

Which should result in output like the following: /etc/kubernetes/kubelet-ca.crt

Next, check the file permissions on each node:

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}');  
do  
oc debug node/${node} -- chroot /host stat -c %a /etc/kubernetes/kubelet-  
ca.crt  
done
```

OR

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}');  
do  
oc debug node/${node} -- chroot /host stat -c %a  
/host/etc/kubernetes/kubelet-ca.crt  
done
```

Verify that the permissions are **644**.

Remediation:

None.







Default Value:

By default, in OpenShift 4, the `kubelet-ca.crt` file has permissions set to `644`.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#about-machine-config-operator_control-plane
2. <https://kubernetes.io/docs/reference/access-authn-authz/authentication/#x509-client-certs>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

4.1.8 Ensure that the client certificate authorities file ownership is set to root:root (Automated)

Profile Applicability:

- Level 1

Description:

Ensure that the certificate authorities file ownership is set to **root:root**.

Rationale:

The certificate authorities file controls the authorities used to validate API requests. You should set its file ownership to maintain the integrity of the file. The file should be owned by **root:root**.

Impact:

None.

Audit:

The Client CA location for the **kubelet** is defined in **/etc/kubernetes/kubelet.conf** and is **/etc/kubernetes/kubelet-ca.crt** by default.

Run the following command to view the user and group ownership:

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}')
do
oc debug node/${node} -- chroot /host stat -c %U:%G /etc/kubernetes/kubelet-
ca.crt
done
```

OR

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}')
do
oc debug node/${node} -- chroot /host stat -c %U:%G
/host/etc/kubernetes/kubelet-ca.crt
done
```

Verify that the ownership is set to **root:root**.

Remediation:

None.







Default Value:

By default, in OpenShift 4, the **--client-ca-file** is set to **/etc/kubernetes/kubelet-ca.crt** with ownership **root:root**.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. <https://kubernetes.io/docs/reference/access-authn-authz/authentication/#x509-client-certs>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	<u>5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts</u> Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	<u>4.3 Ensure the Use of Dedicated Administrative Accounts</u> Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

4.1.9 Ensure that the kubelet --config configuration file has permissions set to 600 or more restrictive (Automated)

Profile Applicability:

- Level 1

Description:

Ensure that if the kubelet refers to a configuration file with the `--config` argument, that file has permissions of `600` or more restrictive.

Rationale:

The kubelet reads various parameters, including security settings, from a config file specified by the `--config` argument. If this file is specified you should restrict its file permissions to maintain the integrity of the file. The file should be writable by only the administrators on the system.

Impact:

None

Audit:

In OpenShift 4, the kubelet configuration file is managed by the Machine Config Operator

Run the following command to check the permission:

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}')
do
oc debug node/${node} -- chroot /host stat -c %a
/var/data/kubelet/config.json
done
```

OR

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}')
do
oc debug node/${node} -- chroot /host stat -c %a
/host/var/data/kubelet/config.json
done
```

Verify that the permissions are `600`.

Remediation:

None.






Default Value:

By default, the `/var/lib/kubelet/config.json` file has permissions of `600`.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. <https://kubernetes.io/docs/tasks/administer-cluster/kubelet-config-file/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	5.2 Maintain Secure Images Maintain secure images or templates for all systems in the enterprise based on the organization's approved configuration standards. Any new system deployment or existing system that becomes compromised should be imaged using one of those images or templates.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1022

4.1.10 Ensure that the kubelet configuration file ownership is set to root:root (Automated)

Profile Applicability:

- Level 1

Description:

Ensure that if the kubelet refers to a configuration file with the `--config` argument, that file is owned by `root:root`.

Rationale:

The kubelet reads various parameters, including security settings, from a config file specified by the `--config` argument. If this file is specified you should restrict its file permissions to maintain the integrity of the file. The file should be owned by `root:root`.

Impact:

None

Audit:

In OpenShift 4, the kubelet configuration file is managed by the Machine Config Operator and is found at `/var/lib/kubelet/config.json` or `/var/data/kubelet/config.json` with ownership set to `root:root`.

In OpenShift 4.13 and above

Run the following command to check the permission:

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}')
do
    oc debug node/${node} -- chroot /host stat -c %a
/var/data/kubelet/config.json
done
```

For Earlier Versions

Run the following command to check the permission:

```
for node in $(oc get nodes -o jsonpath='{.items[*].metadata.name}')
do
    oc debug node/${node} -- chroot /host stat -c %a
/var/lib/kubelet/config.json
done
```

Verify that the ownership is set to `root:root`.

Remediation:

None.







Default Value:

By default, `/var/lib/kubelet/config.json` file is owned by `root:root`.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. <https://kubernetes.io/docs/tasks/administer-cluster/kubelet-config-file/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	4.3 Ensure the Use of Dedicated Administrative Accounts Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1083, T1222	TA0005, TA0007	M1026

4.2 Kubelet

This section contains recommendations for kubelet configuration. In OpenShift 4, the kubelet is managed by the Machine Config Operator.

4.2.1 Activate Garbage collection in OpenShift Container Platform 4, as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

Configure garbage collection for containers and images as appropriate

Rationale:

Garbage collection is important to ensure sufficient resource availability and avoiding degraded performance and availability. In the worst case, the system might crash or just be unusable for a long period of time. Based on your system resources and tests, choose an appropriate threshold value to activate garbage collection.

Impact:

Incorrect configuration of garbage collection parameters can lead to system instability, degraded performance, and in worst cases, system crashes. Properly set parameters ensure efficient utilization of system resources."

Audit:

Two types of garbage collection are performed on an OpenShift Container Platform node:

- Container garbage collection: Removes terminated containers.
- Image garbage collection: Removes images not referenced by any running pods.

Container garbage collection can be performed using eviction thresholds. Image garbage collection relies on disk usage as reported by cAdvisor on the node to decide which images to remove from the node. The OpenShift administrator can configure how OpenShift Container Platform performs garbage collection by creating a `kubeletConfig` object for each Machine Config Pool using any combination of the following:

- soft eviction, which evicts containers based on eviction settings and a grace period
- hard eviction, which evicts containers based on eviction settings without a grace period
- eviction for images

To configure, follow the directions in [Freeing Node Resources Using Garbage Collection](#)

To verify settings, run the following command for each updated **configpool**

To verify, you can inspect the configuration of each node individually:

```
for node in $(oc get nodes -ojsonpath='{.items[*].metadata.name}')
do
    oc get --raw /api/v1/nodes/$node/proxy/configz | jq '.kubeletconfig'
done
```

You can verify the values of the **evictionHard** settings

Verify the values for the following are set as appropriate.

- **evictionHard**
- **evictionPressureTransitionPeriod**
- **imageMinimumGCAge**
- **imageGCHighThresholdPercent**
- **imageGCLowThresholdPercent**
- **evictionSoft** (if configured)
- **evictionSoftGracePeriod** (if configured)

Remediation:

To configure, follow the directions in [Garbage Collection Remediation](#)

Default Value:

The kubelet has the following default hard eviction thresholds:

```
{
  "imagefs.available": "15%",
  "memory.available": "100Mi",
  "nodefs.available": "10%",
  "nodefs.inodesFree": "5%"
}
```






Noted: These default values of hard eviction thresholds will only be set if none of the parameters is changed. If you changed the value of any parameter, then the values of other parameters will not be inherited as the default values and will be set to zero. In order to provide custom values, you should provide all the thresholds respectively.

References:

1. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#cluster-openshift-controller-manager-operator_cluster-operators-ref
2. https://docs.openshift.com/container-platform/latest/operators/operator-reference.html#kube-controller-manager-operator_cluster-operators-ref
3. <https://docs.openshift.com/container-platform/latest/nodes/nodes/nodes-nodes-garbage-collection.html>
4. <https://github.com/openshift/kubernetes-kubelet/blob/origin-4.5-kubernetes-1.18.3/config/v1beta1/types.go#L554-L604>

5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-controller-manager/>
6. <https://github.com/kubernetes/kubernetes/issues/28484>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	10.5 <u>Enable Anti-Exploitation Features</u> Enable anti-exploitation features on enterprise assets and software, where possible, such as Microsoft® Data Execution Prevention (DEP), Windows® Defender Exploit Guard (WDEG), or Apple® System Integrity Protection (SIP) and Gatekeeper™.			
v7	5.1 <u>Establish Secure Configurations</u> Maintain documented, standard security configuration standards for all authorized operating systems and software.			

4.2.2 Ensure that the `--anonymous-auth` argument is set to `false` (Automated)

Profile Applicability:

- Level 1

Description:

Disable anonymous requests to the Kubelet server.

Rationale:

When enabled, requests that are not rejected by other configured authentication methods are treated as anonymous requests. These requests are then served by the Kubelet server. You should rely on authentication to authorize access and disallow anonymous requests.

Impact:

Anonymous requests will be rejected.

Audit:

In OpenShift 4, the Kubernetes configuration file is managed by the Machine Config Operator and `anonymous-auth` is set to `false` by default.

Run the following command on each node to the configuration of anonymous authentication:

```
for node in $(oc get nodes -ojsonpath='{.items[*].metadata.name}'); do
  oc get --raw /api/v1/nodes/$node/proxy/configz | jq
  '.kubeletconfig.authentication.anonymous.enabled';
done
```

Verify that the configuration for each node returns `false`.

Remediation:

Create a `kubeletconfig` to explicitly disable anonymous authentication. Examples of how to do this can be found in the OpenShift [documentation](#).

Default Value:







By default, anonymous access is set to `false`.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane

2. https://docs.openshift.com/container-platform/latest/post_installation_configuration/machine-configuration-tasks.html#create-a-kubeletconfig-crd-to-edit-kubelet-parameters_post-install-machine-configuration-tasks
3. <https://github.com/openshift/machine-config-operator/blob/release-4.5/docs/KubeletConfigDesign.md>
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>
5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet-authentication-authorization/#kubelet-authentication>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 Configure Data Access Control Lists Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

4.2.3 Ensure that the `--authorization-mode` argument is not set to `AlwaysAllow` (Automated)

Profile Applicability:

- Level 1

Description:

Do not allow all requests. Enable explicit authorization.

Rationale:

Kubelets, by default, allow all authenticated requests (even anonymous ones) without needing explicit authorization checks from the apiserver. You should restrict this behavior and only allow explicitly authorized requests.

Impact:

Unauthorized requests will be denied.

Audit:

In OpenShift 4, the Kubernetes configuration file is managed by the Machine Config Operator. By default, OpenShift rejects unauthenticated and unauthorized users.

You can verify that each node in the cluster is configured to only accept authenticated users with the following command:

```
for node in $(oc get nodes -ojsonpath='{.items[*].metadata.name}'); do
  oc get --raw /api/v1/nodes/$node/proxy/configz | jq
  '.kubeletconfig.authorization.mode';
done
```

Verify none of the nodes return `AlwaysAllow` for the authorization mode.

Remediation:

None.

Default Value:






By default, OpenShift uses `Webhook` authorization.

References:

1. https://docs.openshift.com/container-platform/4.5/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. https://docs.openshift.com/container-platform/4.5/scalability_and_performance/recommended-host-practices.html#create-a-kubeletconfig-crd-to-edit-kubelet-parameters

3. <https://github.com/openshift/machine-config-operator/blob/release-4.5/docs/KubeletConfigDesign.md>
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>
5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet-authentication-authorization/#kubelet-authentication>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.3 <u>Configure Data Access Control Lists</u> Configure data access control lists based on a user's need to know. Apply data access control lists, also known as access permissions, to local and remote file systems, databases, and applications.			
v7	9.2 <u>Ensure Only Approved Ports, Protocols and Services Are Running</u> Ensure that only network ports, protocols, and services listening on a system with validated business needs, are running on each system.			

4.2.4 Ensure that the `--client-ca-file` argument is set as appropriate (Automated)

Profile Applicability:

- Level 1

Description:

Enable Kubelet authentication using certificates.

Rationale:

The connections from the apiserver to the kubelet are used for fetching logs for pods, attaching (through kubectl) to running pods, and using the kubelet's port-forwarding functionality. These connections terminate at the kubelet's HTTPS endpoint. By default, the apiserver does not verify the kubelet's serving certificate, which makes the connection subject to man-in-the-middle attacks, and unsafe to run over untrusted and/or public networks. Enabling Kubelet certificate authentication ensures that the apiserver could authenticate the Kubelet before submitting any requests.

Impact:

You require TLS to be configured on apiserver as well as kubelets.

Audit:

OpenShift provides integrated management of certificates for internal cluster components. OpenShift 4 includes multiple CAs providing independent chains of trust, which ensure that a platform CA will never accidentally sign a certificate that can be used for the wrong purpose, increasing the security posture of the cluster.

You can verify the client CA file with the following command:

```
for node in $(oc get nodes -ojsonpath='{.items[*].metadata.name}'); do
oc get --raw /api/v1/nodes/$node/proxy/configz | jq
'.kubeletconfig.authentication.x509.clientCAFile'
done
```

Verify all the nodes are using `/etc/kubernetes/kubelet-ca.crt` as the `clientCAFile` value.

Remediation:

None. Changing the `clientCAFile` value is unsupported.

Default Value:

By default, the `clientCAFile` is set to `/etc/kubernetes/kubelet-ca.crt`.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. https://docs.openshift.com/container-platform/latest/post_installation_configuration/machine-configuration-tasks.html#create-a-kubeletconfig-crd-to-edit-kubelet-parameters_post-install-machine-configuration-tasks
3. <https://github.com/openshift/cluster-kube-apiserver-operator/blob/release-4.5/bindata/v4.1.0/config/defaultconfig.yaml#L28-L29>
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>
5. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet-authentication-authorization/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.11 Encrypt Sensitive Data at Rest Encrypt sensitive data at rest on servers, applications, and databases containing sensitive data. Storage-layer encryption, also known as server-side encryption, meets the minimum requirement of this Safeguard. Additional encryption methods may include application-layer encryption, also known as client-side encryption, where access to the data storage device(s) does not permit access to the plain-text data.		●	●
v7	14.8 Encrypt Sensitive Information at Rest Encrypt all sensitive information at rest using a tool that requires a secondary authentication mechanism not integrated into the operating system, in order to access the information.			●

4.2.5 Verify that the read only port is not used or is set to 0 (Automated)

Profile Applicability:

- Level 1

Description:

Disable the read-only port.

Rationale:

The Kubelet process provides a read-only API in addition to the main Kubelet API. Unauthenticated access is provided to this read-only API which could possibly retrieve potentially sensitive information about the cluster.

Impact:

Removal of the read-only port will require that any service which made use of it will need to be re-configured to use the main Kubelet API.

Audit:

In OpenShift 4, the kubelet is managed by the Machine Config Operator. The kubelet config file is found at `/etc/kubernetes/kubelet.conf`. OpenShift disables the read-only port (10255) on all nodes by setting the `read-only-port` kubelet flag to 0 by default in OpenShift 4.6 and above.

Run the following command to verify the `kubelet-read-only-port` is set to 0 for the Kubernetes API server configuration map.

```
oc -n openshift-kube-apiserver get cm config -o json | jq -r '.data."config.yaml"' | yq '.apiServerArguments."kubelet-read-only-port"'
```

Verify the output is a list that contains 0, like the following:

```
[  
  "0"  
]
```

Remediation:

In earlier versions of OpenShift 4, the `read-only-port` argument is not used.

Follow the instructions in the [documentation](#) to create a `kubeletconfig` CRD and set the `kubelet-read-only-port` is set to 0.





Default Value:

By default, in OpenShift 4.5 and earlier, the `--read-only-port` is not used. In OpenShift 4.6 and above, the `kubelet-read-only-port` is set to 0.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. https://docs.openshift.com/container-platform/latest/post_installation_configuration/machine-configuration-tasks.html#create-a-kubeletconfig-crd-to-edit-kubelet-parameters_post-install-machine-configuration-tasks
3. <https://github.com/openshift/kubernetes-kubelet/blob/origin-4.5-kubernetes-1.18.3/config/v1beta1/types.go#L135-L141>
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.8 <u>Uninstall or Disable Unnecessary Services on Enterprise Assets and Software</u> Uninstall or disable unnecessary services on enterprise assets and software, such as an unused file sharing service, web application module, or service function.			
v7	9.2 <u>Ensure Only Approved Ports, Protocols and Services Are Running</u> Ensure that only network ports, protocols, and services listening on a system with validated business needs, are running on each system.			

4.2.6 Ensure that the `--streaming-connection-idle-timeout` argument is not set to 0 (Automated)

Profile Applicability:

- Level 1

Description:

Do not disable timeouts on streaming connections.

Rationale:

Setting idle timeouts ensures that you are protected against Denial-of-Service attacks, inactive connections and running out of ephemeral ports.

Note: By default, `--streaming-connection-idle-timeout` is set to 4 hours which might be too high for your environment. Setting this as appropriate would additionally ensure that such streaming connections are timed out after serving legitimate use cases.

Impact:

Long-lived connections could be interrupted.

Audit:

OpenShift uses the kubernetes default of 4 hours for the `streaming-connection-idle-timeout` argument. Unless the cluster administrator has added the value to the node configuration, the default will be used. The value is a timeout for HTTP streaming sessions going through a kubelet, like the port-forward, exec, or attach pod operations. The `streaming-connection-idle-timeout` should not be disabled by setting it to **zero**, but it can be lowered. Note that if the value is set too low, then users using those features may experience a service interruption due to the timeout.

The kubelet configuration is currently serialized as an ignition configuration, so it can be directly edited. However, there is also a new `kubelet-config-controller` added to the Machine Config Controller (MCC). This allows you to create a `KubeletConfig` custom resource (CR) to edit the kubelet parameters.

Run the following command to view the streaming connection timeout for each node:

```
for node in $(oc get nodes -ojsonpath='{.items[*].metadata.name}'); do
oc get --raw /api/v1/nodes/$node/proxy/configz | jq
'.kubeletconfig.streamingConnectionIdleTimeout'
done
```

Verify the values returned for each node are not **0**.

Remediation:

Follow the instructions in the [documentation](#) to create a `kubeletconfig` CRD and set the `streamingConnectionIdleTimeout` to the desired value. Do not set the value to `0`.





Default Value:

By default, `streamingConnectionIdleTimeout` is set to 4 hours.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. https://docs.openshift.com/container-platform/latest/post_installation_configuration/machine-configuration-tasks.html#create-a-kubeletconfig-crd-to-edit-kubelet-parameters_post-install-machine-configuration-tasks
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>
4. <https://github.com/kubernetes/kubernetes/pull/18552>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.8 <u>Uninstall or Disable Unnecessary Services on Enterprise Assets and Software</u> Uninstall or disable unnecessary services on enterprise assets and software, such as an unused file sharing service, web application module, or service function.			
v7	9.2 <u>Ensure Only Approved Ports, Protocols and Services Are Running</u> Ensure that only network ports, protocols, and services listening on a system with validated business needs, are running on each system.			

4.2.7 Ensure that the `--make-iptables-util-chains` argument is set to `true` (Manual)

Profile Applicability:

- Level 1

Description:

Allow Kubelet to manage iptables.

Rationale:

Kubelets can automatically manage the required changes to iptables based on how you choose your networking options for the pods. It is recommended to let kubelets manage the changes to iptables. This ensures that the iptables configuration remains in sync with pods networking configuration. Manually configuring iptables with dynamic pod network configuration changes might hamper the communication between pods/containers and to the outside world. You might have iptables rules too restrictive or too open.

Impact:

Kubelet would manage the iptables on the system and keep it in sync. If you are using any other iptables management solution, then there might be some conflicts.

Audit:

Use the following command to ensure each node sets `makeIPTablesUtilChains` to `true`.

```
for node in $(oc get nodes -ojsonpath='{.items[*].metadata.name}'); do
oc get --raw /api/v1/nodes/$node/proxy/configz | jq
'.kubeletconfig.makeIPTablesUtilChains'
done
```

Verify the output returned for each node is `true`.

Remediation:

None.

Default Value:





By default, the `makeIPTablesUtilChains` argument is set to `true`.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane

2. <https://github.com/openshift/kubernetes-kubelet/blob/origin-4.5-kubernetes-1.18.3/config/v1beta1/types.go#L618-L626>
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.4 Implement and Manage a Firewall on Servers Implement and manage a firewall on servers, where supported. Example implementations include a virtual firewall, operating system firewall, or a third-party firewall agent.			
v7	14.7 Enforce Access Control to Data through Automated Tools Use an automated tool, such as host-based Data Loss Prevention, to enforce access controls to data even when data is copied off a system.			

4.2.8 Ensure that the kubeAPIQPS [--event-qps] argument is set to a level which ensures appropriate event capture (Manual)

Profile Applicability:

- Level 2

Description:

Security relevant information should be captured. The `--event-qps` flag on the Kubelet can be used to limit the rate at which events are gathered. Setting this too low could result in relevant events not being logged, however the unlimited setting of `0` could result in a denial of service on the kubelet.

Rationale:

It is important to capture all events and not restrict event creation. Events are an important source of security information and analytics that ensure that your environment is consistently monitored using the event data.

Impact:

Setting this parameter to `0` could result in a denial of service condition due to excessive events being created. The cluster's event processing and storage systems should be scaled to handle expected event loads.

- The kubelet's `--event-qps` (surfaced in OpenShift via `kubeAPIQPS` / `eventRecordQPS`) controls the rate at which events are created/sent.
- Setting it to `0` disables rate limiting, meaning event creation is unlimited.
- Multiple OpenShift/Kubernetes security guides explicitly state that:
 - Setting `--event-qps` too low risks missing important events, but
 - Setting it to the unlimited value `0` “could result in a denial of service on the kubelet” due to excessive events overwhelming event processing and storage.

Audit:

OpenShift uses the `kubeAPIQPS` argument and sets it to a default value of `50`. When this value is set to `> 0`, event creations per second are limited to the value set. If this value is set to `0`, event creations per second are unlimited.

Use the following command to validate the `kubeAPIQPS` configuration:

```
for node in $(oc get nodes -ojsonpath='{.items[*].metadata.name}'); do
oc get --raw /api/v1/nodes/$node/proxy/configz | jq
'.kubeletconfig.kubeAPIQPS'
done
```

Review the value set for the **kubeAPIQPS** argument and determine whether this has been set to an appropriate level for the cluster. If this value is set to **0**, event creations per second are unlimited.

Remediation:

None by default.

Follow the [documentation](#) to edit **kubeletconfig** parameters.





Default Value:

By default, the **kubeAPIQPS** argument is set to **50**.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. https://docs.openshift.com/container-platform/latest/post_installation_configuration/machine-configuration-tasks.html#create-a-kubeletconfig-crd-to-edit-kubelet-parameters_post-install-machine-configuration-tasks
3. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>
4. <https://github.com/kubernetes/kubernetes/blob/master/pkg/kubelet/apis/kubeletconfig/v1beta1/types.go>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	8.5 Collect Detailed Audit Logs Configure detailed audit logging for enterprise assets containing sensitive data. Include event source, date, username, timestamp, source addresses, destination addresses, and other useful elements that could assist in a forensic investigation.			
v7	8.8 Enable Command-line Audit Logging Enable command-line audit logging for command shells, such as Microsoft Powershell and Bash.			

4.2.9 Ensure that the `--tls-cert-file` and `--tls-private-key-file` arguments are set as appropriate (Manual)

Profile Applicability:

- Level 1

Description:

Setup TLS connection on the Kubelets.

Rationale:

The connections from the **apiserver** to the kubelet are used for fetching logs for pods, attaching (through kubectl) to running pods, and using the kubelet's port-forwarding functionality. These connections terminate at the kubelet's HTTPS endpoint. By default, the **apiserver** does not verify the kubelet's serving certificate, which makes the connection subject to man-in-the-middle attacks, and unsafe to run over untrusted and/or public networks.

Impact:

TLS and client certificate authentication must be configured for your Kubernetes cluster deployment.

Audit:

By default, OpenShift uses X.509 certificates to provide secure connections between the API server and **node/kubelet**. OpenShift Container Platform monitors certificates for proper validity, for the cluster certificates it issues and manages. The OpenShift Container Platform manages certificate rotation and the alerting framework has rules to help identify when a certificate issue is about to occur.

Use the following command to check the kubelet client certificate:

```
oc get configmap config -n openshift-kube-apiserver -ojson | jq -r  
'data["config.yaml"]' | jq '.apiServerArguments."kubelet-client-  
certificate"'
```

Verify the certificate path contains **/etc/kubernetes/static-pod-certs/secrets/kubelet-client/tls.crt**.

Use the following command to check the kubelet client key:

```
oc get configmap config -n openshift-kube-apiserver -ojson | jq -r  
'data["config.yaml"]' | jq '.apiServerArguments."kubelet-client-key"'
```

Verify the key path contains **/etc/kubernetes/static-pod-certs/secrets/kubelet-client/tls.key**.

Remediation:

OpenShift automatically manages TLS authentication for the API server communication with the `node/kubelet`. This is not configurable.





Default Value:

By default, OpenShift uses X.509 certificates to provide secure connections between the API server and `node/kubelet`. OpenShift does not use values assigned to the `tls-cert-file` or `tls-private-key-file` flags.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane
2. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>
3. <https://rootsquash.com/2016/05/10/securing-the-kubernetes-api/>
4. <https://github.com/kelseyhightower/docker-kubernetes-tls-guide>
5. <https://jvns.ca/blog/2017/08/05/how-kubernetes-certificates-work/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.10 <u>Encrypt Sensitive Data in Transit</u> Encrypt sensitive data in transit. Example implementations can include: Transport Layer Security (TLS) and Open Secure Shell (OpenSSH).			
v7	14.4 <u>Encrypt All Sensitive Information in Transit</u> Encrypt all sensitive information in transit.			

4.2.10 Ensure that the `--rotate-certificates` argument is not set to `false` (Manual)

Profile Applicability:

- Level 2

Description:

Enable kubelet client certificate rotation.

Rationale:

The `--rotate-certificates` setting causes the kubelet to rotate its client certificates by creating new CSRs as its existing credentials expire. This automated periodic rotation ensures that there is no downtime due to expired certificates and thus addressing availability in the CIA security triad.

Note: This recommendation only applies if you let kubelets get their certificates from the API server. In case your kubelet certificates come from an outside authority/tool (e.g. Vault) then you need to take care of rotation yourself.

Impact:

None.

Audit:

This feature also requires the `RotateKubeletClientCertificate` feature gate to be enabled. The feature gate is enabled by default.

Run the following command to check that certificate rotation is enabled for each node:

```
for node in $(oc get nodes -ojsonpath='{.items[*].metadata.name}'); do
oc get --raw /api/v1/nodes/$node/proxy/configz | jq
'.kubeletconfig.rotateCertificates'
done
```

Verify all the nodes return `true`.

Remediation:

None.

Default Value:





By default, in OpenShift 4, kubelet client certificate rotation is enabled.

References:

1. https://docs.openshift.com/container-platform/latest/architecture/control-plane.html#understanding-machine-config-operator_control-plane

2. <https://github.com/openshift/kubernetes-kubelet/blob/origin-4.5/kubernetes-1.18.3/config/v1beta1/types.go#L172-L181>
3. <https://github.com/openshift/machine-config-operator/blob/release-4.5/templates/master/01-master-kubelet/ base/files/kubelet.yaml>
4. <https://github.com/openshift/machine-config-operator/blob/release-4.5/templates/worker/01-worker-kubelet/ base/files/kubelet.yaml>
5. <https://github.com/kubernetes/kubernetes/pull/41912>
6. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet-tls-bootstrapping/#kubelet-configuration>
7. <https://kubernetes.io/docs/imported/release/notes/>
8. <https://kubernetes.io/docs/reference/command-line-tools-reference/feature-gates/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.10 <u>Encrypt Sensitive Data in Transit</u> Encrypt sensitive data in transit. Example implementations can include: Transport Layer Security (TLS) and Open Secure Shell (OpenSSH).			
v7	14.4 <u>Encrypt All Sensitive Information in Transit</u> Encrypt all sensitive information in transit.			

4.2.11 Verify that the *RotateKubeletServerCertificate* argument is set to true (Manual)

Profile Applicability:

- Level 2

Description:

Enable kubelet server certificate rotation.

Rationale:

RotateKubeletServerCertificate causes the kubelet to both request a serving certificate after bootstrapping its client credentials and rotate the certificate as its existing credentials expire. This automated periodic rotation ensures that there are no downtimes due to expired certificates and thus addressing availability in the CIA security triad.

Note: This recommendation only applies if you let kubelets get their certificates from the API server. In case your kubelet certificates come from an outside authority/tool (e.g. Vault) then you need to take care of rotation yourself.

Impact:

None.

Audit:

Run the following command to check if the **RotateKubeletServerCertificate** feature gate is enabled on each node:

```
for node in $(oc get nodes -ojsonpath='{.items[*].metadata.name}'); do
oc get --raw /api/v1/nodes/$node/proxy/configz | jq
'.kubeletconfig.featureGates.RotateKubeletServerCertificate'
done
```

Verify that all nodes return **true**.

Run the following command to check that each node is configured to rotate certificates:

```
for node in $(oc get nodes -ojsonpath='{.items[*].metadata.name}'); do
oc get --raw /api/v1/nodes/$node/proxy/configz | jq
'.kubeletconfig.rotateCertificates'
done
```

Verify that all nodes return **true**.

Remediation:

None.





Default Value:

By default, `RotateKubeletServerCertificate` is enabled.

References:

1. <https://github.com/openshift/machine-config-operator/blob/release-4.5/templates/master/01-master-kubelet/base/files/kubelet.yaml>
2. <https://github.com/openshift/machine-config-operator/blob/release-4.5/templates/worker/01-worker-kubelet/base/files/kubelet.yaml>
3. <https://github.com/kubernetes/kubernetes/pull/45059>
4. <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet-tls-bootstrapping/#kubelet-configuration>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.10 Encrypt Sensitive Data in Transit Encrypt sensitive data in transit. Example implementations can include: Transport Layer Security (TLS) and Open Secure Shell (OpenSSH).			
v7	14.4 Encrypt All Sensitive Information in Transit Encrypt all sensitive information in transit.			

4.2.12 Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers (Manual)

Profile Applicability:

- Level 1

Description:

Ensure that the Kubelet is configured to only use strong cryptographic ciphers.

Rationale:

TLS ciphers have had a number of known vulnerabilities and weaknesses, which can reduce the protection provided by them. By default Kubernetes supports a number of TLS ciphersuites including some that have security concerns, weakening the protection provided.

Impact:

Kubelet clients that cannot support modern cryptographic ciphers will not be able to make connections to the Kubelet API.

Audit:

The set of cryptographic ciphers currently considered secure is the following:

```
TLS_ECDHE_ECDSA_WITH_AES_128_GCM_SHA256
TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256
TLS_ECDHE_ECDSA_WITH_CHACHA20_POLY1305
TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384
TLS_ECDHE_RSA_WITH_CHACHA20_POLY1305
TLS_ECDHE_ECDSA_WITH_AES_256_GCM_SHA384
TLS_RSA_WITH_AES_256_GCM_SHA384
TLS_RSA_WITH_AES_128_GCM_SHA256
```

Ciphers for the API servers, authentication and the ingress controller can be configured using the **tlsSecurityProfile** parameter as of OpenShift 4.3. The ingress controller provides external access to the API server. There are four TLS security profile types:

- Old
- Intermediate
- Modern
- Custom

Only the Old, Intermediate and Custom profiles are supported at this time for the Ingress controller. Custom provides the ability to specify individual TLS security profile parameters. Follow the steps in the documentation to configure the cipher suite for Ingress, API server and Authentication. <https://docs.openshift.com/container-platform/4.5/networking/ingress-operator.html#nw-ingress-controller-configuration-parameters> [configuring-ingress](#)

Run the following commands to verify the cipher suite and minTLSversion for the ingress operator, authentication operator, **cliconfig**, OpenShift **APIserver** and Kube APIServer.

Use the following command to verify the available ciphers:

```
oc get --namespace=openshift-ingress-operator ingresscontroller/default -ojson | jq .status.tlsProfile.ciphers
```

Verify the ciphers used by the Kubernetes API server:

```
oc get kubeapiservers.operator.openshift.io cluster -ojson | jq .spec.observedConfig.servingInfo.cipherSuites
```

Verify the ciphers used by the OpenShift API server:

```
oc get openshiftapiservers.operator.openshift.io cluster -ojson | jq .spec.observedConfig.servingInfo.cipherSuites
```

Verify the ciphers used by OpenShift authentication:

```
oc get cm -n openshift-authentication v4-0-config-system-cliconfig -o jsonpath='{.data.v4\0-config-system-cliconfig}' | jq .servingInfo.cipherSuites
```

Verify **tlsSecurityProfile** is using the default value:

```
oc get kubeapiservers.operator.openshift.io cluster -ojson | jq .spec.tlsSecurityProfile
```

Verify that the cipher suites are appropriate.

Verify that the **tlsSecurityProfile** is set to the value you chose, or using the default of **Intermediate**.

Note: The HAProxy Ingress controller image does not support TLS 1.3 and because the **Modern** profile requires TLS 1.3, it is not supported. The Ingress Operator converts the **Modern** profile to **Intermediate**. The Ingress Operator also converts the TLS 1.0 of an **Old** or **Custom** profile to 1.1, and TLS 1.3 of a **Custom** profile to 1.2.

Remediation:








Follow the directions above and in the OpenShift documentation to configure the **tlsSecurityProfile**. [Configuring Ingress](#).

Please reference the OpenShift TLS security profile [documentation](#) for more detail on each profile.

Default Value:

By default the Kubernetes API server supports a wide range of TLS ciphers.

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.2 <u>Use Unique Passwords</u> Use unique passwords for all enterprise assets. Best practice implementation includes, at a minimum, an 8-character password for accounts using MFA and a 14-character password for accounts not using MFA.			
v7	1.8 <u>Utilize Client Certificates to Authenticate Hardware Assets</u> Use client certificates to authenticate hardware assets connecting to the organization's trusted network.			
v7	2.6 <u>Address unapproved software</u> Ensure that unauthorized software is either removed or the inventory is updated in a timely manner			

5 Policies

This section contains recommendations for various Kubernetes policies which are important to the security of the environment.

5.1 RBAC and Service Accounts

5.1.1 Ensure that the cluster-admin role is only used where required (Manual)

Profile Applicability:

- Level 1

Description:

The RBAC role `cluster-admin` provides wide-ranging powers over the environment and should be used only where and when needed.

Rationale:

Kubernetes provides a set of default roles where RBAC is used. Some of these roles such as `cluster-admin` provide wide-ranging privileges which should only be applied where absolutely necessary. Roles such as `cluster-admin` allow super-user access to perform any action on any resource. When used in a `ClusterRoleBinding`, it gives full control over every resource in the cluster and in all namespaces. When used in a `RoleBinding`, it gives full control over every resource in the rolebinding's namespace, including the namespace itself.

Impact:

Care should be taken before removing any `clusterrolebindings` from the environment to ensure they were not required for operation of the cluster. Specifically, modifications should not be made to `clusterrolebindings` with the `system:` prefix as they are required for the operation of system components.

Audit:

OpenShift provides a set of default cluster roles that you can bind to users and groups cluster-wide or locally (per project namespace). Be mindful of the difference between local and cluster bindings. For example, if you bind the cluster-admin role to a user by using a local role binding, it might appear that this user has the privileges of a cluster administrator. This is not the case. Binding the cluster-admin to a user in a project grants super administrator privileges for only that project to the user. You can use the oc CLI to view cluster roles and bindings by using the oc describe command. For more information, see [Default Cluster Roles](#)

Some of these roles such as cluster-admin provide wide-ranging privileges which should only be applied where absolutely necessary. Roles such as cluster-admin allow super-user access to perform any action on any resource. When used in a `ClusterRoleBinding`, it gives full control over every resource in the cluster and in all namespaces. When used in a `RoleBinding`, it gives full control over every resource in the rolebinding's namespace, including the namespace itself.

Review users and groups bound to cluster-admin and decide whether they require such access. Consider creating least-privilege roles for users and service accounts.

Obtain a list of the principals who have access to the cluster-admin role by reviewing the **clusterrolebinding** output for each role binding that has access to the cluster-admin role.

```
# needs verification

# To get a list of users and service accounts with the cluster-admin role
oc get clusterrolebindings -o=custom-
columns=NAME:.metadata.name,ROLE:.roleRef.name,SUBJECT:.subjects[*].kind |
grep cluster-admin

# To verify that kubeadmin is removed, no results should be returned
oc get secrets kubeadmin -n kube-system
```

Review each principal listed and ensure that cluster-admin privilege is required for it.

Verify that the kubeadmin user no longer exists.

Remediation:

Identify all **clusterrolebindings** to the cluster-admin role. Check if they are used and if they need this role or if they could use a role with fewer privileges.

Where possible, first bind users to a lower privileged role and then remove the **clusterrolebinding** to the cluster-admin role :

```
oc delete clusterrolebinding [name]
```

Default Value:




By default a single **clusterrolebinding** called **cluster-admin** is provided with the **system:masters** group as its principal.




The principal for cluster-admin also includes system:cluster-admins (Group) and system:admin (User).

References:

1. <https://kubernetes.io/docs/reference/access-authn-authz/rbac/#user-facing-roles>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			

Controls Version	Control	IG 1	IG 2	IG 3
v7	4.3 <u>Ensure the Use of Dedicated Administrative Accounts</u> Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078, T1078.002	TA0001, TA0004	M1026

5.1.2 Minimize access to secrets (Manual)

Profile Applicability:

- Level 1

Description:

The Kubernetes API stores secrets, which may be service account tokens for the Kubernetes API or credentials used by workloads in the cluster. Access to these secrets should be restricted to the smallest possible group of users to reduce the risk of privilege escalation.

Rationale:

Inappropriate access to secrets stored within the Kubernetes cluster can allow for an attacker to gain additional access to the Kubernetes cluster or external resources whose credentials are stored as secrets.

Impact:

Care should be taken not to remove access to secrets to system components which require this for their operation

Audit:

Review the users who have **get**, **list** or **watch** access to **secrets** objects in the Kubernetes API.

Executing the command below will return a list of users and groups who are allowed to get secrets:

```
oc adm policy who-can get secrets
```

The following command returns the users and groups who are allowed to list secrets:

```
oc adm policy who-can list secrets
```

The following command returns the users and groups who are allowed to watch secrets:

```
oc adm policy who-can watch secrets
```

Remediation:

Where possible, remove **get**, **list** and **watch** access to **secret** objects in the cluster.

Default Value:

By default in a OpenShift cluster the following list of principals have get privileges on secret objects

```





for i in $(oc get clusterroles -o jsonpath='{.items[*].metadata.name}'); do
oc describe clusterrole ${i}; done

# The following default cluster roles have get privileges on secret objects

admin
cloud-credential-operator-role
cluster-admin
cluster-image-registry-operator
cluster-monitoring-operator
cluster-node-tuning-operator
edit
kube-state-metrics
machine-config-controller
marketplace-operator
openshift-ingress-operator
prometheus-operator
registry-admin
registry-editor
system:aggregate-to-edit
system:controller:expand-controller
system:controller:generic-garbage-collector
system:controller:namespace-controller
system:controller:operator-lifecycle-manager
system:controller:persistent-volume-binder
system:kube-controller-manager
system:master
system:node
system:openshift:controller:build-controller
system:openshift:controller:cluster-quota-reconciliation-controller
system:openshift:controller:ingress-to-route-controller
system:openshift:controller:service-ca
system:openshift:controller:service-serving-cert-controller
system:openshift:controller:serviceaccount-pull-secrets-controller
system:openshift:controller:template-service-broker

```

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	<u>3.12 Segment Data Processing and Storage Based on Sensitivity</u> Segment data processing and storage based on the sensitivity of the data. Do not process sensitive data on enterprise assets intended for lower sensitivity data.			
v7	<u>13.4 Only Allow Access to Authorized Cloud Storage or Email Providers</u> Only allow access to authorized cloud storage or email providers.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1026

5.1.3 Minimize wildcard use in Roles and ClusterRoles (Manual)

Profile Applicability:

- Level 1

Description:

Kubernetes Roles and ClusterRoles provide access to resources based on sets of objects and actions that can be taken on those objects. It is possible to set either of these to be the wildcard "*" which matches all items.

Use of wildcards is not optimal from a security perspective as it may allow for inadvertent access to be granted when new resources are added to the Kubernetes API either as CRDs or in later versions of the product.

Rationale:

The principle of least privilege recommends that users are provided only the access required for their role and nothing more. The use of wildcard rights grants is likely to provide excessive rights to the Kubernetes API.

Audit:

Run the command below to describe each cluster role and inspect it for wildcard usage:

```
oc describe clusterrole
```


Run the command below to describe each role and inspect it for wildcard usage:






```
oc describe role -A
```

Remediation:

Where possible replace any use of wildcards in clusterroles and roles with specific objects or actions.

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	6.8 Define and Maintain Role-Based Access Control Define and maintain role-based access control, through determining and documenting the access rights necessary for each role within the enterprise to successfully carry out its assigned duties. Perform access control reviews of enterprise assets to validate that all privileges are authorized, on a recurring schedule at a minimum annually, or more frequently.			

Controls Version	Control	IG 1	IG 2	IG 3
v7	4.4 Use Unique Passwords Where multi-factor authentication is not supported (such as local administrator, root, or service accounts), accounts will use passwords that are unique to that system.			
v7	14.6 Protect Information through Access Control Lists Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078, T1078.002	TA0001, TA0004	M1026

5.1.4 Minimize access to create pods (Manual)

Profile Applicability:

- Level 1

Description:

The ability to create pods in a namespace can provide a number of opportunities for privilege escalation, such as assigning privileged service accounts to these pods or mounting hostPaths with access to sensitive data (unless Pod Security Policies are implemented to restrict this access)

As such, access to create new pods should be restricted to the smallest possible group of users.

Rationale:

The ability to create pods in a cluster opens up possibilities for privilege escalation and should be restricted, where possible.

Impact:

Care should be taken not to remove access to pods to system components which require this for their operation

Audit:

Review the users who have create access to pod objects in the Kubernetes API with the following command:

```
oc adm policy who-can create pod
```

Remediation:

Where possible, remove **create** access to **pod** objects in the cluster.

Default Value:

By default in a kubeadm cluster the following list of principals have **create** privileges on **pod** objects

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	2.7 Allowlist Authorized Scripts Use technical controls, such as digital signatures and version control, to ensure that only authorized scripts, such as specific .ps1, .py, etc., files, are allowed to execute. Block unauthorized scripts from executing. Reassess bi-annually, or more frequently.			●
v7	5.2 Maintain Secure Images Maintain secure images or templates for all systems in the enterprise based on the organization's approved configuration standards. Any new system deployment or existing system that becomes compromised should be imaged using one of those images or templates.		●	●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078, T1078.002	TA0001, TA0004	M1026

5.1.5 Ensure that default service accounts are not actively used. (Manual)

Profile Applicability:

- Level 1

Description:

The **default** service account should not be used to ensure that rights granted to applications can be more easily audited and reviewed.

Rationale:

Kubernetes provides a **default** service account which is used by cluster workloads where no specific service account is assigned to the pod.

Where access to the Kubernetes API from a pod is required, a specific service account should be created for that pod, and rights granted to that service account.

The default service account should be configured such that it does not provide a service account token and does not have any explicit rights assignments.

Impact:

All workloads which require access to the Kubernetes API will require an explicit service account to be created.

Audit:

Every OpenShift project has its own service accounts. Every service account has an associated user name that can be granted roles, just like a regular user. The user name for each service account is derived from its project and the name of the service account. Service accounts are required in each project to run builds, deployments, and other pods. The default service accounts that are automatically created for each project are isolated by the project namespace.

Remediation:

None required.







Default Value:

By default, in OpenShift 4 every project has its own service accounts. Every service account has an associated user name that can be granted roles, just like a regular user. The user name for each service account is derived from its project and the name of the service account. Service accounts are required in each project to run builds, deployments, and other pods. The default service accounts that are automatically created for each project are isolated by the project namespace.

References:

1. <https://kubernetes.io/docs/tasks/configure-pod-container/configure-service-account/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.3 <u>Disable Dormant Accounts</u> Delete or disable any dormant accounts after a period of 45 days of inactivity, where supported.			
v7	16.9 <u>Disable Dormant Accounts</u> Automatically disable dormant accounts after a set period of inactivity.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1028

5.1.6 Ensure that Service Account Tokens are only mounted where necessary (Manual)

Profile Applicability:

- Level 1

Description:

Service accounts tokens should not be mounted in pods except where the workload running in the pod explicitly needs to communicate with the API server

Rationale:

Mounting service account tokens inside pods can provide an avenue for privilege escalation attacks where an attacker is able to compromise a single pod in the cluster.

Avoiding mounting these tokens removes this attack avenue.

Impact:

Pods mounted without service account tokens will not be able to communicate with the API server, except where the resource is available to unauthenticated principals.

Audit:

Review pod and service account objects in the cluster and ensure automatically mounting the service account token is disabled (**automountServiceAccountToken: false**), unless the resource explicitly requires this access.

Find all pods that automatically mount service account tokens:

```
oc get pods -A -o json | jq '.items[] | select(.spec.automountServiceAccountToken) | .metadata.name'
```

Find all service accounts that automatically mount service tokens:

```
oc get serviceaccounts -A -o json | jq '.items[] | select(.automountServiceAccountToken) | .metadata.name'
```

Remediation:

Modify the definition of pods and service accounts which do not need to mount service account tokens to disable it.





Default Value:

By default, all pods get a service account token mounted in them.

References:

1. <https://kubernetes.io/docs/tasks/configure-pod-container/configure-service-account/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.8 <u>Uninstall or Disable Unnecessary Services on Enterprise Assets and Software</u> Uninstall or disable unnecessary services on enterprise assets and software, such as an unused file sharing service, web application module, or service function.			
v7	13.4 <u>Only Allow Access to Authorized Cloud Storage or Email Providers</u> Only allow access to authorized cloud storage or email providers.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1528, T1555	TA0006	M1026

5.2 Security Context Constraints

A Security Context Constraint (SCC) is a cluster-level resource that controls security settings for pods. OpenShift uses the Security Context Constraints (SCC) admission controller plugin instead of Pod Security Policies (PSP). The SCC plugin cannot be disabled. The PSP plugin cannot be enabled.

Similar to the way that RBAC resources control user access, administrators can use Security Context Constraints (SCCs) to control permissions for pods. These permissions include actions that a pod, a collection of containers, can perform and what resources it can access. You can use SCCs to define a set of conditions that a pod must run with in order to be accepted into the system.

Make sure that the use of SCCs complies with site policy.

By default, OpenShift 4 is configured with multiple SCCs. You can query SCCs with the following command:

```
oc get scc
```

5.2.1 Minimize the admission of privileged containers (Manual)

Profile Applicability:

- Level 1

Description:

Do not generally permit containers to be run with the `securityContext.privileged` flag set to `true`.

Rationale:

Privileged containers have access to all Linux Kernel capabilities and devices. A container running with full privileges can do almost everything that the host can do. This flag exists to allow special use-cases, like manipulating the network stack and accessing devices.

There should be at least one Security Context Constraint (SCC) defined which does not permit privileged containers.

If you need to run privileged containers, this should be defined in a separate SCC and you should carefully check RBAC controls to ensure that only limited service accounts and users are given permission to access that SCC.

Impact:

Pods defined with `spec.containers[].securityContext.privileged: true` will not be permitted.

Audit:

The set of SCCs that admission uses to authorize a pod are determined by the user identity and groups that the user belongs to. Additionally, if the pod specifies a service account, the set of allowable SCCs includes any constraints accessible to the service account.

Admission uses the following approach to create the final security context for the pod:

- Retrieve all SCCs available for use.
- Generate field values for security context settings that were not specified on the request.
- Validate the final settings against the available constraints.

If a matching set of constraints is found, then the pod is accepted. If the request cannot be matched to an SCC, the pod is rejected.

A pod must validate every field against the SCC.

You can use the following command to list all SCCs that do not allow privileged containers:

```
oc get scc -o json | jq '.items[] | select(.allowPrivilegedContainer==false) | .metadata.name'
```

Verify that at least one SCC is returned.

Remediation:

Create an SCC that sets **allowPrivilegedContainer** to **false** and take it into use by assigning it to applicable users and groups.

Default Value:





By default, the following SCCs do not allow users to create privileged containers:




```
"anyuid"  
"hostaccess"  
"hostmount-anyuid"  
"hostnetwork"  
"hostnetwork-v2"  
"machine-api-termination-handler"  
"nonroot"  
"nonroot-v2"  
"restricted"  
"restricted-v2"
```

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/managing-security-context-constraints.html>
2. <https://kubernetes.io/docs/concepts/policy/pod-security-policy/#enabling-pod-security-policies>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v8	12.8 Establish and Maintain Dedicated Computing Resources for All Administrative Work Establish and maintain dedicated computing resources, either physically or logically separated, for all administrative tasks or tasks requiring administrative access. The computing resources should be segmented from the enterprise's primary network and not be allowed internet access.			

Controls Version	Control	IG 1	IG 2	IG 3
v7	4.3 <u>Ensure the Use of Dedicated Administrative Accounts</u> Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1611	TA0004	M1048

5.2.2 Minimize the admission of containers wishing to share the host process ID namespace (Manual)

Profile Applicability:

- Level 1

Description:

Do not generally permit containers to be run with the **hostPID** flag set to true.

Rationale:

A container running in the host's PID namespace can inspect processes running outside the container. If the container also has access to ptrace capabilities this can be used to escalate privileges outside of the container.

There should be at least one Security Context Constraint (SCC) defined which does not permit containers to share the host PID namespace.

If you need to run containers which require hostPID, this should be defined in a separate SCC and you should carefully check RBAC controls to ensure that only limited service accounts and users are given permission to access that SCC.

Impact:

Pods defined with **Allow Host PID: true** will not be permitted unless they are run under a specific SCC.

Audit:

Use the following command to list all SCCs with **allowHostPID** set to **true**:

```
oc get scc -o json | jq '.items[] | select(.allowHostPID) | .metadata.name'
```

Verify that at least one SCC is returned.

Remediation:

Create an SCC that sets **allowHostPID** to **false** and take it into use by assigning it to applicable users and groups.

Default Value:






By default, the following SCCs do not allow users to run within the host process namespace:


```
"anyuid"  
"hostmount-anyuid"  
"hostnetwork"  
"hostnetwork-v2"  
"machine-api-termination-handler"  
"nonroot"  
"nonroot-v2"  
"restricted"  
"restricted-v2"
```

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/managing-security-context-constraints.html>
2. <https://kubernetes.io/docs/concepts/policy/pod-security-policy/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v7	5.2 Maintain Secure Images Maintain secure images or templates for all systems in the enterprise based on the organization's approved configuration standards. Any new system deployment or existing system that becomes compromised should be imaged using one of those images or templates.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078, T1078.002	TA0001, TA0004	M1026

5.2.3 Minimize the admission of containers wishing to share the host IPC namespace (Manual)

Profile Applicability:

- Level 1

Description:

Do not generally permit containers to be run with the **hostIPC** flag set to true.

Rationale:

A container running in the host's IPC namespace can use IPC to interact with processes outside the container.

There should be at least one Security Context Constraint (SCC) defined which does not permit containers to share the host IPC namespace.

If you have a requirement to containers which require hostIPC, this should be defined in a separate SCC and you should carefully check RBAC controls to ensure that only limited service accounts and users are given permission to access that SCC.

Impact:

Pods defined with **Allow Host IPC: true** will not be permitted unless they are run under a specific SCC.

Audit:

Use the following command to list all SCCs with **allowHostIPC** set to **false**:

```
oc get scc -o json | jq '.items[] | select(.allowHostIPC==false) | .metadata.name'
```

Verify at least one SCC is returned.

Remediation:

Create an SCC that sets **allowHostIPC** to **false** and take it into use by assigning it to applicable users and groups.

Default Value:

By default, the following SCCs do not allow users to run within the host IPC namespace:

```
"anyuid"  
"hostmount-anyuid"  
"hostnetwork"  
"hostnetwork-v2"  
"machine-api-termination-handler"  
"node-exporter"  
"nonroot"  
"nonroot-v2"  
"restricted"  
"restricted-v2"
```

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/managing-security-context-constraints.html>
2. <https://kubernetes.io/docs/concepts/policy/pod-security-policy/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	13.10 Perform Application Layer Filtering Perform application layer filtering. Example implementations include a filtering proxy, application layer firewall, or gateway.			●
v7	12.9 Deploy Application Layer Filtering Proxy Server Ensure that all network traffic to or from the Internet passes through an authenticated application layer proxy that is configured to filter unauthorized connections.			●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1078, T1078.002	TA0001, TA0004	M1026

5.2.4 Minimize the admission of containers wishing to share the host network namespace (Manual)

Profile Applicability:

- Level 1

Description:

Do not generally permit containers to be run with the **hostNetwork** flag set to true.

Rationale:

A container running in the host's network namespace could access the local loopback device, and could access network traffic to and from other pods.

There should be at least one Security Context Constraint (SCC) defined which does not permit containers to share the host network namespace.

If you have need to run containers which require hostNetwork, this should be defined in a separate SCC and you should carefully check RBAC controls to ensure that only limited service accounts and users are given permission to access that SCC.

Impact:

Pods defined with **Allow Host Network: true** will not be permitted unless they are run under a specific SCC.

Audit:

Use the following command to list all SCCs with **allowHostNetwork** set to **false**:

```
oc get scc -A -o json | jq '.items[] | select(.allowHostNetwork==false) | .metadata.name'
```

Verify at least one SCC is returned.

Remediation:

Create an SCC that sets **allowHostNetwork** to **false** and take it into use by assigning it to applicable users and groups.

Default Value:





By default, the following SCCs do not allow access to the host network:

```
"anyuid"  
"hostmount-anyuid"  
"nonroot"  
"nonroot-v2"  
"restricted"  
"restricted-v2"
```

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/managing-security-context-constraints.html>
2. <https://kubernetes.io/docs/concepts/policy/pod-security-policy/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.12 Segment Data Processing and Storage Based on Sensitivity Segment data processing and storage based on the sensitivity of the data. Do not process sensitive data on enterprise assets intended for lower sensitivity data.			
v7	14.1 Segment the Network Based on Sensitivity Segment the network based on the label or classification level of the information stored on the servers, locate all sensitive information on separated Virtual Local Area Networks (VLANs).			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1098	TA0003	M1030

5.2.5 Minimize the admission of containers with `allowPrivilegeEscalation` (Manual)

Profile Applicability:

- Level 1

Description:

Do not generally permit containers to be run with the `allowPrivilegeEscalation` flag set to `true`.

Rationale:

A container running with the `allowPrivilegeEscalation` flag set to `true` may have processes that can gain more privileges than their parent.

There should be at least one Security Context Constraint (SCC) defined which does not permit containers to allow privilege escalation. The option exists (and is defaulted to `true`) to permit `setuid` binaries to run.

If you have need to run containers which use `setuid` binaries or require privilege escalation, this should be defined in a separate SCC and you should carefully check RBAC controls to ensure that only limited service accounts and users are given permission to access that SCC.

Impact:

Pods defined with `Allow Privilege Escalation: true` will not be permitted unless they are run under a specific SCC.

Audit:

Use the following command to list all SCCs with `allowPrivilegeEscalation` set to `false`:

```
oc get scc -A -o json | jq '.items[] | select(.allowPrivilegeEscalation==false) | .metadata.name'
```

Verify that there is at least one SCC is returned.

Remediation:

Create an SCC that sets `allowPrivilegeEscalation` to `false` and take it into use by assigning it to applicable users and groups.

Default Value:








By default, the following SCCs do not allow privilege escalation:

```
"hostnetwork-v2"  
"nonroot-v2"  
"restricted-v2"
```

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/managing-security-context-constraints.html>
2. <https://kubernetes.io/docs/concepts/policy/pod-security-policy/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 <u>Restrict Administrator Privileges to Dedicated Administrator Accounts</u> Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v8	12.8 <u>Establish and Maintain Dedicated Computing Resources for All Administrative Work</u> Establish and maintain dedicated computing resources, either physically or logically separated, for all administrative tasks or tasks requiring administrative access. The computing resources should be segmented from the enterprise's primary network and not be allowed internet access.			
v7	4.3 <u>Ensure the Use of Dedicated Administrative Accounts</u> Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1611	TA0004	M1038, M1048

5.2.6 Minimize the admission of root containers (Manual)

Profile Applicability:

- Level 2

Description:

Do not generally permit containers to be run as the root user.

Rationale:

Containers may run as any Linux user. Containers which run as the root user, whilst constrained by Container Runtime security features still have an escalated likelihood of container breakout.

Ideally, all containers should run as a defined non-UID 0 user.

There should be at least one Security Context Constraint (SCC) defined which does not permit root users in a container.

If you need to run root containers, this should be defined in a separate SCC and you should carefully check RBAC controls to ensure that only limited service accounts and users are given permission to access that SCC.

Impact:

Pods with containers which run as the root user will not be permitted.

Audit:

Use the following command to list all SCCs that restrict the ability to run the container as root:

```
oc get scc -A -o json | jq '.items[] | select(.runAsUser["type"] == "MustRunAsNonRoot") | .metadata.name'
```

Verify at least one SCC is returned.

You can perform additional validation by using the following command to list the UID range for each SCC:

```
for i in `oc get scc --template '{{range .items}}{{.metadata.name}}{{"\n"}}{{end}}'`; do echo "$i"; oc describe scc $i | grep "\sUID"; done
```

Verify there is at least one SCC that doesn't contain 0 in the UID range.

Remediation:

None required. By default, OpenShift includes the **nonroot** and **nonroot-v2** SCCs that restrict the ability to run as nonroot. If additional SCCs are appropriate, follow the OpenShift documentation to create custom SCCs.

Default Value:








By default, the following SCCs restrict the ability to run as non-root:

```
"nonroot"  
"nonroot-v2"
```

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/managing-security-context-constraints.html>
2. <https://kubernetes.io/docs/concepts/policy/pod-security-policy/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	5.4 Restrict Administrator Privileges to Dedicated Administrator Accounts Restrict administrator privileges to dedicated administrator accounts on enterprise assets. Conduct general computing activities, such as internet browsing, email, and productivity suite use, from the user's primary, non-privileged account.			
v8	12.8 Establish and Maintain Dedicated Computing Resources for All Administrative Work Establish and maintain dedicated computing resources, either physically or logically separated, for all administrative tasks or tasks requiring administrative access. The computing resources should be segmented from the enterprise's primary network and not be allowed internet access.			
v7	4.3 Ensure the Use of Dedicated Administrative Accounts Ensure that all users with administrative account access use a dedicated or secondary account for elevated activities. This account should only be used for administrative activities and not internet browsing, email, or similar activities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1098	TA0003	M1026

5.2.7 Minimize the admission of containers with the NET_RAW capability (Manual)

Profile Applicability:

- Level 1

Description:

Do not generally permit containers with the potentially dangerous NET_RAW capability.

Rationale:

Containers run with a default set of capabilities as assigned by the Container Runtime. By default this can include potentially dangerous capabilities. With Docker as the container runtime the NET_RAW capability is enabled which may be misused by malicious containers.

Ideally, all containers should drop this capability.

There should be at least one Security Context Constraint (SCC) defined which prevents containers with the NET_RAW capability from launching.

If you need to run containers with this capability, this should be defined in a separate SCC and you should carefully check RBAC controls to ensure that only limited service accounts and users are given permission to access that SCC.

Impact:

Pods with containers which run with the NET_RAW capability will not be permitted.

Audit:

Use the following command to list all SCCs that drop all capabilities:

```
oc get scc -A -o json | jq '.items[] | select(.requiredDropCapabilities[]?|any(. == "ALL"; .)) | .metadata.name'
```

Verify at least one SCC is returned.

Remediation:

Create an SCC that sets requiredDropCapabilities to include ALL or at least NET_RAW and take it into use by assigning it to applicable users and groups.

Default Value:

By default, the following SCCs drop all capabilities:

```
"hostnetwork-v2"  
"nonroot-v2"  
"restricted-v2"
```

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/managing-security-context-constraints.html>
2. <https://kubernetes.io/docs/concepts/policy/pod-security-policy/#enabling-pod-security-policies>
3. <https://www.nccgroup.trust/uk/our-research/abusing-privileged-and-unprivileged-linux-containers/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.8 <u>Uninstall or Disable Unnecessary Services on Enterprise Assets and Software</u> Uninstall or disable unnecessary services on enterprise assets and software, such as an unused file sharing service, web application module, or service function.		●	●
v7	5.2 <u>Maintain Secure Images</u> Maintain secure images or templates for all systems in the enterprise based on the organization's approved configuration standards. Any new system deployment or existing system that becomes compromised should be imaged using one of those images or templates.		●	●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1068	TA0004	M1050

5.2.8 Minimize the admission of containers with added capabilities (Manual)

Profile Applicability:

- Level 1

Description:

Do not generally permit containers with capabilities assigned beyond the default set.

Rationale:

Containers run with a default set of capabilities as assigned by the Container Runtime. Capabilities outside this set can be added to containers which could expose them to risks of container breakout attacks.

There should be at least one Security Context Constraint (SCC) defined which prevents containers with capabilities beyond the default set from launching.

If you need to run containers with additional capabilities, this should be defined in a separate SCC and you should carefully check RBAC controls to ensure that only limited service accounts and users are given permission to access that SCC.

Impact:

Pods with containers which require capabilities outside the default set will not be permitted.

Audit:

Use the following command to list all SCCs that prohibit users from defining container capabilities:

```
oc get scc -A -o json | jq '.items[] | select(.allowedCapabilities==null) | .metadata.name'
```

Verify at least one SCC is returned.

Additionally, use the following command to list all SCCs that do not set default container capabilities:

```
oc get scc -A -o json | jq '.items[] | select(.defaultAddCapabilities==null) | .metadata.name'
```

Verify at least one SCC is returned.

Remediation:

Utilize the restricted-v2 SCC or create an SCC that sets **allowedCapabilities** and **defaultAddCapabilities** to an empty list and take it into use by assigning it to applicable users and groups.





Default Value:

By default authenticated users are allowed to use the restricted-v2 SCC, which drops all container capabilities.

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/managing-security-context-constraints.html>
2. <https://kubernetes.io/docs/concepts/policy/pod-security-policy/#enabling-pod-security-policies>
3. <https://www.nccgroup.com/uk/our-research/abusing-privileged-and-unprivileged-linux-containers/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.8 <u>Uninstall or Disable Unnecessary Services on Enterprise Assets and Software</u> Uninstall or disable unnecessary services on enterprise assets and software, such as an unused file sharing service, web application module, or service function.			
v7	5.2 <u>Maintain Secure Images</u> Maintain secure images or templates for all systems in the enterprise based on the organization's approved configuration standards. Any new system deployment or existing system that becomes compromised should be imaged using one of those images or templates.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1204	TA0002, TA0003	M1047

5.2.9 Minimize the admission of containers with capabilities assigned (Manual)

Profile Applicability:

- Level 2

Description:

Do not generally permit containers with capabilities.

Rationale:

Containers run with a default set of capabilities as assigned by the Container Runtime. Capabilities are parts of the rights generally granted on a Linux system to the root user.

In many cases applications running in containers do not require any capabilities to operate, so from the perspective of the principal of least privilege use of capabilities should be minimized.

Impact:

Pods with containers which require capabilities to operate will not be permitted.

Audit:

Use the following command to list SCCs that drop all capabilities from containers:

```
oc get scc -A -o json | jq '.items[] | select(.requiredDropCapabilities[]?|any(. == "ALL"; .)) | .metadata.name'
```

Verify at least one SCC is returned.

Remediation:

Review the use of capabilities in applications running on your cluster. Where a namespace contains applications which do not require any Linux capabilities to operate, consider adding a SCC which forbids the admission of containers which do not drop all capabilities.

Default Value:

By default, OpenShift includes three SCCs that drop all container capabilities:

```
"hostnetwork-v2"  
"nonroot-v2"  
"restricted-v2"
```

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/managing-security-context-constraints.html>

2. <https://kubernetes.io/docs/concepts/policy/pod-security-policy/#enabling-pod-security-policies>
3. <https://www.nccgroup.com/uk/our-research/abusing-privileged-and-unprivileged-linux-containers/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.8 <u>Uninstall or Disable Unnecessary Services on Enterprise Assets and Software</u> Uninstall or disable unnecessary services on enterprise assets and software, such as an unused file sharing service, web application module, or service function.		●	●
v7	5.2 <u>Maintain Secure Images</u> Maintain secure images or templates for all systems in the enterprise based on the organization's approved configuration standards. Any new system deployment or existing system that becomes compromised should be imaged using one of those images or templates.		●	●

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1204	TA0002, TA0003	M1045, M1047

5.2.10 Minimize access to privileged Security Context Constraints (Manual)

Profile Applicability:

- Level 2

Description:

OpenShift has the concept of Security Context Constraints (SCCs) that supplement the Pod Security Admission controller.

SCCs allow you to group elevated container capabilities and assign those capabilities to users and groups. For example, you can have an SCC that restricts the ability to launch privileged containers and assign that SCC to all authenticated users. As a result, users requesting a pod that contains a privileged container will be rejected.

You can find more information on SCCs in the OpenShift [documentation](#).

Rationale:

SCCs that contain the ability to permit privileged or elevated container action should be carefully managed. Users with access to such an SCC can leverage the privileged functionality granted by that SCC, increasing the risk of compromising the container or host.

Impact:

Users should only have access to SCCs that allow them to perform functions required by their roles, and no more, following the principle of least privilege.

Audit:

Find all users and groups with access to SCCs that include privileged or elevated capabilities:

```
oc get scc -ojson | jq '.items[]|select(.allowHostIPC or .allowHostPID or .allowHostPorts or .allowHostNetwork or .allowHostDirVolumePlugin or .allowPrivilegedContainer or .runAsUser.type != "MustRunAsRange")|.metadata.name,{"Group":.groups},{"User":.users}'
```

Review the returned users and groups and verify they actually need access to those SCCs.

Remediation:

Remove any users and groups who do not need access to an SCC, following the principle of least privilege.

You can remove users and groups from an SCC using the `oc edit scc $NAME` command.

Additionally, you can create your own SCCs that contain the container functionality you need for a particular use case and assign that SCC to users and groups if the default SCCs are not appropriate for your use case.

Default Value:

OpenShift provides the following SCCs by default:

```
"anyuid"
"hostaccess"
"hostmount-anyuid"
"hostnetwork"
"hostnetwork-v2"
"machine-api-termination-handler"
"node-exporter"
"nonroot"
"nonroot-v2"
"privileged"
"restricted"
"restricted-v2"
```

These default SCCs attempt to group similar privileged container functionality into a single SCC that fits particular use cases.

Please refer to the OpenShift [documentation](#) for a complete list of capabilities associated with each default SCC.

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	6.8 Define and Maintain Role-Based Access Control Define and maintain role-based access control, through determining and documenting the access rights necessary for each role within the enterprise to successfully carry out its assigned duties. Perform access control reviews of enterprise assets to validate that all privileges are authorized, on a recurring schedule at a minimum annually, or more frequently.			●
v7	4.7 Limit Access to Script Tools Limit access to scripting tools (such as Microsoft PowerShell and Python) to only administrative or development users with the need to access those capabilities.		●	●

5.3 Network Policies and CNI

5.3.1 Ensure that the CNI in use supports Network Policies (Manual)

Profile Applicability:

- Level 1

Description:

There are a variety of CNI plugins available for Kubernetes. If the CNI in use does not support Network Policies it may not be possible to effectively restrict traffic in the cluster.

Rationale:

Kubernetes network policies are enforced by the CNI plugin in use. As such it is important to ensure that the CNI plugin supports both Ingress and Egress network policies.

Impact:

None

Audit:

Review the documentation of the CNI plugin (OpenShift SDN or OVN-Kubernetes) in use by the cluster, and confirm that it supports ingress and egress network policies.

OpenShift Container Platform uses a software-defined networking (SDN) approach to offer a unified cluster network that enables communication between Pods across the OpenShift Container Platform cluster. This Pod network can be established and maintained by either the OpenShift SDN plugin, which configures an overlay network using Open vSwitch (OVS), or by the OVN-Kubernetes plugin.

OVN-Kubernetes, based on Open Virtual Network (OVN), offers an overlay-based networking implementation. A cluster using the OVN-Kubernetes plugin also runs Open vSwitch (OVS) on each node. OVN configures OVS on each node to implement the declared network configuration.

As of OpenShift 4.10, both these plugins provide all Kubernetes Network Policy features, including ingress, egress, and **ipBlock**. OpenShift provides several options for controlling the traffic leaving the cluster, the supporting matrix can be found at.

Please refer to the OpenShift [documentation](#) for a complete list of features available for securing cluster networks.

Remediation:

None required.





Default Value:

This will depend on the CNI plugin in use.

References:

1. <https://docs.openshift.com/container-platform/4.5/networking/openshift-sdn/about-openshift-sdn.html>
2. <https://kubernetes.io/docs/concepts/extend-kubernetes/compute-storage-net/network-plugins/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.4 Implement and Manage a Firewall on Servers Implement and manage a firewall on servers, where supported. Example implementations include a virtual firewall, operating system firewall, or a third-party firewall agent.			
v7	9.5 Implement Application Firewalls Place application firewalls in front of any critical servers to verify and validate the traffic going to the server. Any unauthorized traffic should be blocked and logged.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1046	TA0007	M1030, M1042

5.3.2 Ensure that all Namespaces have Network Policies defined (Manual)

Profile Applicability:

- Level 2

Description:

Use network policies to isolate traffic in your cluster network.

Rationale:

Running different applications on the same Kubernetes cluster creates a risk of one compromised application attacking a neighboring application. Network segmentation is important to ensure that containers can communicate only with those they are supposed to. A network policy is a specification of how selections of pods are allowed to communicate with each other and other network endpoints.

Once there is any Network Policy in a namespace selecting a particular pod, that pod will reject any connections that are not allowed by any Network Policy. Other pods in the namespace that are not selected by any Network Policy will continue to accept all traffic

Impact:

Once there is any Network Policy in a namespace selecting a particular pod, that pod will reject any connections that are not allowed by any Network Policy. Other pods in the namespace that are not selected by any Network Policy will continue to accept all traffic"

Audit:

The OpenShift 4 CNI plugin uses network policies and by default all Pods in a project are accessible from other Pods and network endpoints. To isolate one or more Pods in a project, you create NetworkPolicy objects in that project to indicate the allowed incoming connections. Project administrators can create and delete NetworkPolicy objects within their own project. For more information see:

Run the following command and review the **NetworkPolicy** objects created in the cluster.

```
oc -n all get networkpolicy
```

Ensure that each namespace defined in the cluster has at least one Network Policy.

Remediation:

Follow the documentation and create **NetworkPolicy** objects as you need them.






Default Value:

By default, all Pods in a project are accessible from other Pods and network endpoints; network policies are not created.

References:

1. https://docs.openshift.com/container-platform/latest/networking/network_policy/about-network-policy.html
2. https://docs.openshift.com/container-platform/latest/networking/network_policy/creating-network-policy.html
3. https://docs.openshift.com/container-platform/latest/networking/network_policy/multitenant-network-policy.html
4. https://docs.openshift.com/container-platform/latest/networking/network_policy/default-network-policy.html
5. <https://kubernetes.io/docs/concepts/services-networking/network-policies/>
6. <https://octetz.com/docs/2019/2019-04-22-netpol-api-k8s/>
7. <https://kubernetes.io/docs/tasks/administer-cluster/declare-network-policy/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.4 Implement and Manage a Firewall on Servers Implement and manage a firewall on servers, where supported. Example implementations include a virtual firewall, operating system firewall, or a third-party firewall agent.			
v7	14.2 Enable Firewall Filtering Between VLANs Enable firewall filtering between VLANs to ensure that only authorized systems are able to communicate with other systems necessary to fulfill their specific responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1046	TA0007	M1030, M1042

5.4 Secrets Management

5.4.1 Prefer using secrets as files over secrets as environment variables (Manual)

Profile Applicability:

- Level 2

Description:

Kubernetes supports mounting secrets as data volumes or as environment variables. Minimize the use of environment variable secrets.

Rationale:

It is reasonably common for application code to log out its environment (particularly in the event of an error). This will include any secret values passed in as environment variables, so secrets can easily be exposed to any user or entity who has access to the logs.

Impact:

Application code which expects to read secrets in the form of environment variables would need modification

Audit:

Information about ways to provide sensitive data to pods is included in the documentation. [Providing sensitive data to pods](#) Run the following command to find references to objects which use environment variables defined from secrets.

```
oc get all -o jsonpath='{range .items[?(@..secretKeyRef)]} {.kind}
{.metadata.name} {"\n"}{end}' -A
```

Remediation:

If possible, rewrite application code to read secrets from mounted secret files, rather than from environment variables.

Default Value:

By default, application secrets are not defined.





In a default OpenShift 4 cluster, the following platform objects are returned


```
Pod aws-ebs-csi-driver-controller-699844b6d7-2pl8k
Pod aws-ebs-csi-driver-controller-6dcc794464-gnl6l
Pod aws-ebs-csi-driver-controller-6dcc794464-kkhr9
Deployment aws-ebs-csi-driver-controller
Deployment aws-ebs-csi-driver-controller
ReplicaSet aws-ebs-csi-driver-controller-699844b6d7
ReplicaSet aws-ebs-csi-driver-controller-777d8fbb87
ReplicaSet aws-ebs-csi-driver-controller-658754b8c8
ReplicaSet aws-ebs-csi-driver-controller-6dcc794464
```

References:

1. <https://kubernetes.io/docs/concepts/configuration/secret/#using-secrets>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.10 <u>Encrypt Sensitive Data in Transit</u> Encrypt sensitive data in transit. Example implementations can include: Transport Layer Security (TLS) and Open Secure Shell (OpenSSH).			
v7	14.4 <u>Encrypt All Sensitive Information in Transit</u> Encrypt all sensitive information in transit.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1026

5.4.2 Consider external secret storage (Manual)

Profile Applicability:

- Level 2

Description:

Consider the use of an external secrets storage and management system, instead of using Kubernetes Secrets directly, if you have more complex secret management needs. Ensure the solution requires authentication to access secrets, has auditing of access to and use of secrets, and encrypts secrets. Some solutions also make it easier to rotate secrets.

Rationale:

Kubernetes supports secrets as first-class objects, but care needs to be taken to ensure that access to secrets is carefully limited. Using an external secrets provider can ease the management of access to secrets, especially where secrets are used across both Kubernetes and non-Kubernetes environments.

Impact:

None

Audit:

OpenShift supports a broad ecosystem of security partners many of whom provide integration with enterprise secret vaults.

Review your secrets management implementation.



Remediation:




Refer to the secrets management options offered by your cloud provider or a third-party secrets management solution.

Default Value:

By default, no external secret management is configured.

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	3.12 Segment Data Processing and Storage Based on Sensitivity Segment data processing and storage based on the sensitivity of the data. Do not process sensitive data on enterprise assets intended for lower sensitivity data.			

Controls Version	Control	IG 1	IG 2	IG 3
v7	14.6 <u>Protect Information through Access Control Lists</u> Protect all information stored on systems with file system, network share, claims, application, or database specific access control lists. These controls will enforce the principle that only authorized individuals should have access to the information based on their need to access the information as a part of their responsibilities.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1552	TA0006	M1026

5.5 Extensible Admission Control

5.5.1 Configure Image Provenance using image controller configuration parameters (Manual)

Profile Applicability:

- Level 2

Description:

Configure Image Provenance for your deployment.

Rationale:

Kubernetes supports plugging in provenance rules to accept or reject the images in your deployments. You could configure such rules to ensure that only approved images are deployed in the cluster.

You can control which images can be imported, tagged, and run in a cluster using the image controller. For additional information on the image controller, see [Image configuration resources](#)

Impact:

You need to regularly maintain your provenance configuration based on container image updates.

Audit:

Review the image controller parameters in your cluster and verify that image provenance is configured as appropriate.

Run the following command to return all registry sources:

```
oc get image.config.openshift.io/cluster -o json | jq .spec.registrySources
```

If nothing is returned, this is a finding and you should refer to the [OpenShift image guide](#) for configuring trusted registries.

Remediation:

Follow the OpenShift documentation: [Image configuration resources](#)

Default Value:






By default, image provenance is not set.

References:

1. <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/#imagepolicywebhook>

2. <https://github.com/kubernetes/community/blob/master/contributors/design-proposals/image-provenance.md>
3. <https://hub.docker.com/r/dnurmi/anchore-toolbox/>
4. <https://github.com/kubernetes/kubernetes/issues/22888>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.1 <u>Establish and Maintain a Secure Configuration Process</u> Establish and maintain a secure configuration process for enterprise assets (end-user devices, including portable and mobile, non-computing/IoT devices, and servers) and software (operating systems and applications). Review and update documentation annually, or when significant enterprise changes occur that could impact this Safeguard.			
v7	5.4 <u>Deploy System Configuration Management Tools</u> Deploy system configuration management tools that will automatically enforce and redeploy configuration settings to systems at regularly scheduled intervals.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1133, T1195	TA0001, TA0003	M1016, M1042

5.7 General Policies

These policies relate to general cluster management topics, like namespace best practices and policies applied to pod objects in the cluster.

5.7.1 Create administrative boundaries between resources using namespaces (Manual)

Profile Applicability:

- Level 1

Description:

Use namespaces to isolate your Kubernetes objects.

Rationale:

Limiting the scope of user permissions can reduce the impact of mistakes or malicious activities. A Kubernetes namespace allows you to partition created resources into logically named groups. Resources created in one namespace can be hidden from other namespaces. By default, each resource created by a user in Kubernetes cluster runs in a default namespace, called **default**. You can create additional namespaces and attach resources and users to them. You can use Kubernetes Authorization plugins to create policies that segregate access to namespace resources between different users.

Impact:

You need to switch between namespaces for administration.

Audit:

OpenShift Projects wrap Kubernetes namespaces and are used by default in OpenShift 4.

Run the following command to obtain a list of all non-default OpenShift and Kubernetes namespaces in the cluster.

```
oc get namespaces -o json | jq '.items[] | select(.metadata.name|test("(?!default|kube-|.openshift.)^.*")) | .metadata.name'
```

Ensure that these namespaces are the ones you need and are adequately administered as per your requirements.

Remediation:

Follow the documentation and create namespaces for objects in your deployment as you need them.

Default Value:

By default, Kubernetes starts with two initial namespaces:





1. **default** - The default namespace for objects with no other namespace

2. **kube-system** - The namespace for objects created by the Kubernetes system
3. **openshift** -
4. **openshift-*** - The namespace for objects created by OpenShift

References:

1. <https://kubernetes.io/docs/concepts/overview/working-with-objects/namespaces/>
2. <https://kubernetes.io/blog/2016/08/security-best-practices-kubernetes-deployment/>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.4 Implement and Manage a Firewall on Servers Implement and manage a firewall on servers, where supported. Example implementations include a virtual firewall, operating system firewall, or a third-party firewall agent.			
v7	9.5 Implement Application Firewalls Place application firewalls in front of any critical servers to verify and validate the traffic going to the server. Any unauthorized traffic should be blocked and logged.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1106, T1609	TA0002, TA0008	M1038

5.7.2 Ensure that the seccomp profile is set to docker/default in your pod definitions (Manual)

Profile Applicability:

- Level 2

Description:

Enable **default** seccomp profile in your pod definitions.

Rationale:

Seccomp (secure computing mode) is used to restrict the set of system calls applications can make, allowing cluster administrators greater control over the security of workloads running in the cluster. Kubernetes disables seccomp profiles by default for historical reasons. You should enable it to ensure that the workloads have restricted actions available within the container.

Impact:

If the **default** seccomp profile is too restrictive for you, you will need to create and manage your own seccomp profiles, which can be done using [OpenShift Security Context Constraints](#) and [custom seccomp profiles](#).

Audit:

In OpenShift 4, CRI-O is the supported runtime. CRI-O runs unconfined by default in order to meet CRI conformance criteria.

On Red Hat CoreOS, the default seccomp policy is associated with CRI-O and stored in **/etc/crio/seccomp.json**. The default profile is applied when the user asks for the **RuntimeDefault** profile via annotation to the pod and when the associated SCC allows use of the specified seccomp profile.

Use the following command to find all non-default pods that do not have a seccomp profile set:

```
oc get pods -A -o json | jq '.items[] | select( (.metadata.namespace | test("^kube*|^openshift*") | not) and .spec.securityContext.seccompProfile.type==null) | (.metadata.namespace + "/" + .metadata.name) '
```

The output will return a list of namespace pod pairs that do not have a seccomp profile set.

Remediation:

For any non-privileged pods or containers that do not have seccomp profiles, consider using the **RuntimeDefault** or creating a custom seccomp profile specifically for the workload.

Please refer to the OpenShift [documentation](#) for working with custom seccomp profiles.






Default Value:

By default, OpenShift applies the **restricted-v2** SCC to all pods, which uses the **RuntimeDefault** seccomp profile.

References:

1. <https://docs.openshift.com/container-platform/latest/authentication/managing-security-context-constraints.html>
2. <https://github.com/kubernetes/kubernetes/issues/39845>
3. <https://github.com/kubernetes/kubernetes/pull/21790>
4. <https://github.com/kubernetes/community/blob/master/contributors/design-proposals/seccomp.md#examples>
5. <https://docs.docker.com/engine/security/seccomp/>
6. <https://docs.openshift.com/container-platform/latest/security/seccomp-profiles.html>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.1 Establish and Maintain a Secure Configuration Process Establish and maintain a secure configuration process for enterprise assets (end-user devices, including portable and mobile, non-computing/IoT devices, and servers) and software (operating systems and applications). Review and update documentation annually, or when significant enterprise changes occur that could impact this Safeguard.			
v7	5.2 Maintain Secure Images Maintain secure images or templates for all systems in the enterprise based on the organization's approved configuration standards. Any new system deployment or existing system that becomes compromised should be imaged using one of those images or templates.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1611	TA0004	M1048

5.7.3 Apply Security Context to Your Pods and Containers (Manual)

Profile Applicability:

- Level 2

Description:

Apply Security Context to Your Pods and Containers

Rationale:

A security context defines the operating system security settings (uid, gid, capabilities, SELinux role, etc..) applied to a container. When designing your containers and pods, make sure that you configure the security context for your pods, containers, and volumes. A security context is a property defined in the deployment yaml. It controls the security parameters that will be assigned to the pod/container/volume. There are two levels of security context: pod level security context, and container level security context.

Impact:

If you incorrectly apply security contexts, you may have trouble running the pods.

Audit:

Review the pod definitions in your cluster and verify that you have security contexts defined as appropriate.

OpenShift's Security Context Constraint feature is on by default in OpenShift 4 and applied to all pods deployed. SCC selection is determined by a combination of the values in the securityContext and the rolebindings for the account deploying the pod.

Run the following command to obtain a list of pods that are using **privileged** security context constraints:

```
oc get pods -A -o json | jq '.items[] | select(.metadata.annotations."openshift.io/scc"|test("privileged?")) | .metadata.name'
```

Review each pod to ensure it requires the use of **privileged** security context constraints, which is the most relaxed security context constraint available in OpenShift by default.

Run the following command to obtain a list of pods that are not using security context constraints at all:

```
oc get pods -A -o json | jq '.items[] | select(.metadata.annotations."openshift.io/scc" == null) | .metadata.name'
```

Review each pod and determine if there is an existing security context constraint that it should be using.

Remediation:

Follow the Kubernetes documentation and apply security contexts to your pods. For a suggested list of security contexts, you may refer to the CIS Security Benchmark for Docker Containers.






Default Value:

By default, no security contexts are automatically applied to pods.

References:

1. <https://kubernetes.io/docs/concepts/policy/security-context/>
2. <https://learn.cisecurity.org/benchmarks>

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	7.3 <u>Perform Automated Operating System Patch Management</u> Perform operating system updates on enterprise assets through automated patch management on a monthly, or more frequent, basis.			
v7	8.3 <u>Enable Operating System Anti-Exploitation Features/ Deploy Anti-Exploit Technologies</u> Enable anti-exploitation features such as Data Execution Prevention (DEP) or Address Space Layout Randomization (ASLR) that are available in an operating system or deploy appropriate toolkits that can be configured to apply protection to a broader set of applications and executables.			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1556, T1611	TA0004, TA0006	M1048

5.7.4 The default namespace should not be used (Manual)

Profile Applicability:

- Level 2

Description:

Kubernetes provides a default namespace, where objects are placed if no namespace is specified for them. Placing objects in this namespace makes application of RBAC and other controls more difficult.

Rationale:

Resources in a Kubernetes cluster should be segregated by namespace, to allow for security controls to be applied at that level and to make it easier to manage resources.

Impact:

None

Audit:

In OpenShift, projects (namespaces) are used to group and isolate related objects. When a request is made to create a new project using the web console or **oc new-project** command, an endpoint in OpenShift Container Platform is used to provision the project according to a template, which can be customized.

The cluster administrator can allow and configure how developers and service accounts can create, or self-provision, their own projects. Regular users do not have access to the **default** project.

Projects starting with **openshift-** and **kube-** host cluster components that run as Pods and other infrastructure components. As such, OpenShift does not allow you to create Projects starting with **openshift-** or **kube-** using the **oc new-project** command.

For more information, see [Working with projects](#) and [Configuring project creation](#)

Run the following command to list all resources in the **default** namespace, besides the **kubernetes** and **openshift** services, which are expected to be in the **default** namespace:

```
oc get all -n default -o json | jq '.items[] | select((.kind|test("Service")) and (.metadata.name|test("openshift|kubernetes"))? | not) | (.kind + "/" + .metadata.name) '
```

There are 15 separate resource categories returned in the json dump from **oc get all**; they all should be empty unless related to a system resource.

Carefully review the list of returned resources and consider moving them to another namespace.



Remediation:

Ensure that namespaces are created to allow for appropriate segregation of Kubernetes resources and that all new resources are created in a specific namespace.

Default Value:

Unless a namespace is specific on object creation, the **default** namespace will be used

CIS Controls:

Controls Version	Control	IG 1	IG 2	IG 3
v8	4.9 Configure Trusted DNS Servers on Enterprise Assets Configure trusted DNS servers on enterprise assets. Example implementations include: configuring assets to use enterprise-controlled DNS servers and/or reputable externally accessible DNS servers.			
v7	5 Secure Configuration for Hardware and Software on Mobile Devices, Laptops, Workstations and Servers Secure Configuration for Hardware and Software on Mobile Devices, Laptops, Workstations and Servers			

MITRE ATT&CK Mappings:

Techniques / Sub-techniques	Tactics	Mitigations
T1578	TA0005	M1018

Appendix: Summary Table

CIS Benchmark Recommendation		Set Correctly	
		Yes	No
1	Control Plane Components		
1.1	Master Node Configuration Files		
1.1.1	Ensure that the API server pod specification file permissions are set to 600 or more restrictive (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.2	Ensure that the API server pod specification file ownership is set to root:root (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.3	Ensure that the controller manager pod specification file permissions are set to 600 or more restrictive (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.4	Ensure that the controller manager pod specification file ownership is set to root:root (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.5	Ensure that the scheduler pod specification file permissions are set to 600 or more restrictive (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.6	Ensure that the scheduler pod specification file ownership is set to root:root (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.7	Ensure that the etcd pod specification file permissions are set to 600 or more restrictive (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.8	Ensure that the etcd pod specification file ownership is set to root:root (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.9	Ensure that the Container Network Interface file permissions are set to 600 or more restrictive (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.10	Ensure that the Container Network Interface file ownership is set to root:root (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.11	Ensure that the etcd data directory permissions are set to 700 or more restrictive (Manual)	<input type="checkbox"/>	<input type="checkbox"/>

CIS Benchmark Recommendation		Set Correctly	
		Yes	No
1.1.12	Ensure that the etcd data directory ownership is set to root:root (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.13	Ensure that the kubeconfig file permissions are set to 600 or more restrictive (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.14	Ensure that the kubeconfig file ownership is set to root:root (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.15	Ensure that the Scheduler kubeconfig file permissions are set to 600 or more restrictive (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.16	Ensure that the Scheduler kubeconfig file ownership is set to root:root (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.17	Ensure that the Controller Manager kubeconfig file permissions are set to 600 or more restrictive (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.18	Ensure that the Controller Manager kubeconfig file ownership is set to root:root (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.19	Ensure that the OpenShift PKI directory and file ownership is set to root:root (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.20	Ensure that the OpenShift PKI certificate file permissions are set to 600 or more restrictive (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.1.21	Ensure that the OpenShift PKI key file permissions are set to 600 (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2	API Server		
1.2.1	Ensure that anonymous requests are authorized (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.2	Use https for kubelet connections (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.3	Ensure that the kubelet uses certificates to authenticate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.4	Verify that the kubelet certificate authority is set as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>

CIS Benchmark Recommendation		Set Correctly	
		Yes	No
1.2.5	Ensure that the --authorization-mode argument is not set to AlwaysAllow (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.6	Verify that RBAC is enabled (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.7	Ensure that the APIPriorityAndFairness feature gate is enabled (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.8	Ensure that the admission control plugin AlwaysAdmit is not set (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.9	Ensure that the admission control plugin AlwaysPullImages is not set (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.10	Ensure that the admission control plugin ServiceAccount is set (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.11	Ensure that the admission control plugin NamespaceLifecycle is set (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.12	Ensure that the admission control plugin SecurityContextConstraint is set (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.13	Ensure that the admission control plugin NodeRestriction is set (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.14	Ensure that the --insecure-bind-address argument is not set (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.15	Ensure that the --insecure-port argument is set to 0 (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.16	Ensure that the --secure-port argument is not set to 0 (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.17	Ensure that the healthz endpoint is protected by RBAC (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.18	Ensure that the --audit-log-path argument is set (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.19	Ensure that the audit logs are forwarded off the cluster for retention (Manual)	<input type="checkbox"/>	<input type="checkbox"/>

CIS Benchmark Recommendation		Set Correctly	
		Yes	No
1.2.20	Ensure that the maximumRetainedFiles argument is set to 10 or as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.21	Configure Kubernetes API Server Maximum Audit Log Size (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.22	Ensure that the --request-timeout argument is set (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.23	Ensure that the --service-account-lookup argument is set to true (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.24	Ensure that the --service-account-key-file argument is set as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.25	Ensure that the --etcd-certfile and --etcd-keyfile arguments are set as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.26	Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.27	Ensure that the --client-ca-file argument is set as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.28	Ensure that the --etcd-cafile argument is set as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.29	Ensure that encryption providers are appropriately configured (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.30	Ensure that the API Server only makes use of Strong Cryptographic Ciphers (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.2.31	Ensure unsupported configuration overrides are not used (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.3	Controller Manager		
1.3.1	Ensure that controller manager healthz endpoints are protected by RBAC (Manual)	<input type="checkbox"/>	<input type="checkbox"/>

CIS Benchmark Recommendation		Set Correctly	
		Yes	No
1.3.2	Ensure that the --use-service-account-credentials argument is set to true (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.3.3	Ensure that the --service-account-private-key-file argument is set as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.3.4	Ensure that the --root-ca-file argument is set as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.4	Scheduler		
1.4.1	Ensure that the healthz endpoints for the scheduler are protected by RBAC (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
1.4.2	Verify that the scheduler API service is protected by RBAC (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
2	etcd		
2.1	Ensure that the --cert-file and --key-file arguments are set as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
2.2	Ensure that the --client-cert-auth argument is set to true (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
2.3	Ensure that the --auto-tls argument is not set to true (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
2.4	Ensure that the --peer-cert-file and --peer-key-file arguments are set as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
2.5	Ensure that the --peer-client-cert-auth argument is set to true (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
2.6	Ensure that the --peer-auto-tls argument is not set to true (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
2.7	Ensure that a unique Certificate Authority is used for etcd (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
3	Control Plane Configuration		

CIS Benchmark Recommendation		Set Correctly	
		Yes	No
3.1	Authentication and Authorization		
3.1.1	Client certificate authentication should not be used for users (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
3.2	Logging		
3.2.1	Ensure that a minimal audit policy is created (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
3.2.2	Ensure that the audit policy covers key security concerns (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
4	Worker Nodes		
4.1	Worker Node Configuration Files		
4.1.1	Ensure that the kubelet service file permissions are set to 644 or more restrictive (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.1.2	Ensure that the kubelet service file ownership is set to root:root (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.1.3	If proxy kube proxy configuration file exists ensure permissions are set to 644 or more restrictive (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
4.1.4	If proxy kubeconfig file exists ensure ownership is set to root:root (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
4.1.5	Ensure that the --kubeconfig kubelet.conf file permissions are set to 644 or more restrictive (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.1.6	Ensure that the --kubeconfig kubelet.conf file ownership is set to root:root (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.1.7	Ensure that the certificate authorities file permissions are set to 644 or more restrictive (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.1.8	Ensure that the client certificate authorities file ownership is set to root:root (Automated)	<input type="checkbox"/>	<input type="checkbox"/>

CIS Benchmark Recommendation		Set Correctly	
		Yes	No
4.1.9	Ensure that the kubelet --config configuration file has permissions set to 600 or more restrictive (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.1.10	Ensure that the kubelet configuration file ownership is set to root:root (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.2	Kubelet		
4.2.1	Activate Garbage collection in OpenShift Container Platform 4, as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
4.2.2	Ensure that the --anonymous-auth argument is set to false (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.2.3	Ensure that the --authorization-mode argument is not set to AlwaysAllow (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.2.4	Ensure that the --client-ca-file argument is set as appropriate (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.2.5	Verify that the read only port is not used or is set to 0 (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.2.6	Ensure that the --streaming-connection-idle-timeout argument is not set to 0 (Automated)	<input type="checkbox"/>	<input type="checkbox"/>
4.2.7	Ensure that the --make-iptables-util-chains argument is set to true (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
4.2.8	Ensure that the kubeAPIQPS [--event-qps] argument is set to a level which ensures appropriate event capture (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
4.2.9	Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
4.2.10	Ensure that the --rotate-certificates argument is not set to false (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
4.2.11	Verify that the RotateKubeletServerCertificate argument is set to true (Manual)	<input type="checkbox"/>	<input type="checkbox"/>

CIS Benchmark Recommendation		Set Correctly	
		Yes	No
4.2.12	Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5	Policies		
5.1	RBAC and Service Accounts		
5.1.1	Ensure that the cluster-admin role is only used where required (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.1.2	Minimize access to secrets (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.1.3	Minimize wildcard use in Roles and ClusterRoles (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.1.4	Minimize access to create pods (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.1.5	Ensure that default service accounts are not actively used. (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.1.6	Ensure that Service Account Tokens are only mounted where necessary (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.2	Security Context Constraints		
5.2.1	Minimize the admission of privileged containers (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.2.2	Minimize the admission of containers wishing to share the host process ID namespace (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.2.3	Minimize the admission of containers wishing to share the host IPC namespace (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.2.4	Minimize the admission of containers wishing to share the host network namespace (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.2.5	Minimize the admission of containers with allowPrivilegeEscalation (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.2.6	Minimize the admission of root containers (Manual)	<input type="checkbox"/>	<input type="checkbox"/>

CIS Benchmark Recommendation		Set Correctly	
		Yes	No
5.2.7	Minimize the admission of containers with the NET_RAW capability (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.2.8	Minimize the admission of containers with added capabilities (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.2.9	Minimize the admission of containers with capabilities assigned (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.2.10	Minimize access to privileged Security Context Constraints (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.3	Network Policies and CNI		
5.3.1	Ensure that the CNI in use supports Network Policies (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.3.2	Ensure that all Namespaces have Network Policies defined (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.4	Secrets Management		
5.4.1	Prefer using secrets as files over secrets as environment variables (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.4.2	Consider external secret storage (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.5	Extensible Admission Control		
5.5.1	Configure Image Provenance using image controller configuration parameters (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.7	General Policies		
5.7.1	Create administrative boundaries between resources using namespaces (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.7.2	Ensure that the seccomp profile is set to docker/default in your pod definitions (Manual)	<input type="checkbox"/>	<input type="checkbox"/>
5.7.3	Apply Security Context to Your Pods and Containers (Manual)	<input type="checkbox"/>	<input type="checkbox"/>

CIS Benchmark Recommendation		Set Correctly	
		Yes	No
5.7.4	The default namespace should not be used (Manual)	<input type="checkbox"/>	<input type="checkbox"/>

Appendix: CIS Controls v7 IG 1 Mapped Recommendations

Recommendation		Set Correctly	
		Yes	No
1.1.1	Ensure that the API server pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.2	Ensure that the API server pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.3	Ensure that the controller manager pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.4	Ensure that the controller manager pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.5	Ensure that the scheduler pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.6	Ensure that the scheduler pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.7	Ensure that the etcd pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.8	Ensure that the etcd pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.9	Ensure that the Container Network Interface file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.10	Ensure that the Container Network Interface file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.11	Ensure that the etcd data directory permissions are set to 700 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.12	Ensure that the etcd data directory ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.13	Ensure that the kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.14	Ensure that the kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.15	Ensure that the Scheduler kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.1.16	Ensure that the Scheduler kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.17	Ensure that the Controller Manager kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.18	Ensure that the Controller Manager kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.19	Ensure that the OpenShift PKI directory and file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.20	Ensure that the OpenShift PKI certificate file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.21	Ensure that the OpenShift PKI key file permissions are set to 600	<input type="checkbox"/>	<input type="checkbox"/>
1.2.1	Ensure that anonymous requests are authorized	<input type="checkbox"/>	<input type="checkbox"/>
1.2.6	Verify that RBAC is enabled	<input type="checkbox"/>	<input type="checkbox"/>
1.2.8	Ensure that the admission control plugin AlwaysAdmit is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.9	Ensure that the admission control plugin AlwaysPullImages is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.10	Ensure that the admission control plugin ServiceAccount is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.11	Ensure that the admission control plugin NamespaceLifecycle is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.12	Ensure that the admission control plugin SecurityContextConstraint is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.13	Ensure that the admission control plugin NodeRestriction is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.15	Ensure that the --insecure-port argument is set to 0	<input type="checkbox"/>	<input type="checkbox"/>
1.2.17	Ensure that the healthz endpoint is protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
1.2.18	Ensure that the --audit-log-path argument is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.22	Ensure that the --request-timeout argument is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.31	Ensure unsupported configuration overrides are not used	<input type="checkbox"/>	<input type="checkbox"/>
1.3.1	Ensure that controller manager healthz endpoints are protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
1.4.1	Ensure that the healthz endpoints for the scheduler are protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
2.5	Ensure that the --peer-client-cert-auth argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
3.2.1	Ensure that a minimal audit policy is created	<input type="checkbox"/>	<input type="checkbox"/>
4.1.1	Ensure that the kubelet service file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.2	Ensure that the kubelet service file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.3	If proxy kube proxy configuration file exists ensure permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.4	If proxy kubeconfig file exists ensure ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.5	Ensure that the --kubeconfig kubelet.conf file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.6	Ensure that the --kubeconfig kubelet.conf file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.7	Ensure that the certificate authorities file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.8	Ensure that the client certificate authorities file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.10	Ensure that the kubelet configuration file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.2.1	Activate Garbage collection in OpenShift Container Platform 4, as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.2	Ensure that the --anonymous-auth argument is set to false	<input type="checkbox"/>	<input type="checkbox"/>
4.2.12	Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers	<input type="checkbox"/>	<input type="checkbox"/>
5.1.1	Ensure that the cluster-admin role is only used where required	<input type="checkbox"/>	<input type="checkbox"/>
5.1.3	Minimize wildcard use in Roles and ClusterRoles	<input type="checkbox"/>	<input type="checkbox"/>
5.1.5	Ensure that default service accounts are not actively used.	<input type="checkbox"/>	<input type="checkbox"/>
5.2.1	Minimize the admission of privileged containers	<input type="checkbox"/>	<input type="checkbox"/>
5.2.5	Minimize the admission of containers with allowPrivilegeEscalation	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
5.2.6	Minimize the admission of root containers	<input type="checkbox"/>	<input type="checkbox"/>
5.4.2	Consider external secret storage	<input type="checkbox"/>	<input type="checkbox"/>

Appendix: CIS Controls v7 IG 2 Mapped Recommendations

Recommendation		Set Correctly	
		Yes	No
1.1.1	Ensure that the API server pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.2	Ensure that the API server pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.3	Ensure that the controller manager pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.4	Ensure that the controller manager pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.5	Ensure that the scheduler pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.6	Ensure that the scheduler pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.7	Ensure that the etcd pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.8	Ensure that the etcd pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.9	Ensure that the Container Network Interface file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.10	Ensure that the Container Network Interface file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.11	Ensure that the etcd data directory permissions are set to 700 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.12	Ensure that the etcd data directory ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.13	Ensure that the kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.14	Ensure that the kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.15	Ensure that the Scheduler kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.1.16	Ensure that the Scheduler kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.17	Ensure that the Controller Manager kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.18	Ensure that the Controller Manager kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.19	Ensure that the OpenShift PKI directory and file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.20	Ensure that the OpenShift PKI certificate file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.21	Ensure that the OpenShift PKI key file permissions are set to 600	<input type="checkbox"/>	<input type="checkbox"/>
1.2.1	Ensure that anonymous requests are authorized	<input type="checkbox"/>	<input type="checkbox"/>
1.2.2	Use https for kubelet connections	<input type="checkbox"/>	<input type="checkbox"/>
1.2.5	Ensure that the --authorization-mode argument is not set to AlwaysAllow	<input type="checkbox"/>	<input type="checkbox"/>
1.2.6	Verify that RBAC is enabled	<input type="checkbox"/>	<input type="checkbox"/>
1.2.7	Ensure that the APIPriorityAndFairness feature gate is enabled	<input type="checkbox"/>	<input type="checkbox"/>
1.2.8	Ensure that the admission control plugin AlwaysAdmit is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.9	Ensure that the admission control plugin AlwaysPullImages is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.10	Ensure that the admission control plugin ServiceAccount is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.11	Ensure that the admission control plugin NamespaceLifecycle is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.12	Ensure that the admission control plugin SecurityContextConstraint is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.13	Ensure that the admission control plugin NodeRestriction is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.14	Ensure that the --insecure-bind-address argument is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.15	Ensure that the --insecure-port argument is set to 0	<input type="checkbox"/>	<input type="checkbox"/>
1.2.16	Ensure that the --secure-port argument is not set to 0	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.2.17	Ensure that the healthz endpoint is protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
1.2.18	Ensure that the --audit-log-path argument is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.19	Ensure that the audit logs are forwarded off the cluster for retention	<input type="checkbox"/>	<input type="checkbox"/>
1.2.20	Ensure that the maximumRetainedFiles argument is set to 10 or as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.21	Configure Kubernetes API Server Maximum Audit Log Size	<input type="checkbox"/>	<input type="checkbox"/>
1.2.22	Ensure that the --request-timeout argument is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.24	Ensure that the --service-account-key-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.26	Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.27	Ensure that the --client-ca-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.28	Ensure that the --etcd-cafile argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.31	Ensure unsupported configuration overrides are not used	<input type="checkbox"/>	<input type="checkbox"/>
1.3.1	Ensure that controller manager healthz endpoints are protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
1.3.2	Ensure that the --use-service-account-credentials argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
1.3.3	Ensure that the --service-account-private-key-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.3.4	Ensure that the --root-ca-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.4.1	Ensure that the healthz endpoints for the scheduler are protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
1.4.2	Verify that the scheduler API service is protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
2.1	Ensure that the --cert-file and --key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
2.3	Ensure that the --auto-tls argument is not set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.4	Ensure that the --peer-cert-file and --peer-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
2.5	Ensure that the --peer-client-cert-auth argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.6	Ensure that the --peer-auto-tls argument is not set to true	<input type="checkbox"/>	<input type="checkbox"/>
3.1.1	Client certificate authentication should not be used for users	<input type="checkbox"/>	<input type="checkbox"/>
3.2.1	Ensure that a minimal audit policy is created	<input type="checkbox"/>	<input type="checkbox"/>
4.1.1	Ensure that the kubelet service file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.2	Ensure that the kubelet service file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.3	If proxy kube proxy configuration file exists ensure permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.4	If proxy kubeconfig file exists ensure ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.5	Ensure that the --kubeconfig kubelet.conf file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.6	Ensure that the --kubeconfig kubelet.conf file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.7	Ensure that the certificate authorities file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.8	Ensure that the client certificate authorities file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.9	Ensure that the kubelet --config configuration file has permissions set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.10	Ensure that the kubelet configuration file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.2.1	Activate Garbage collection in OpenShift Container Platform 4, as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.2	Ensure that the --anonymous-auth argument is set to false	<input type="checkbox"/>	<input type="checkbox"/>
4.2.3	Ensure that the --authorization-mode argument is not set to AlwaysAllow	<input type="checkbox"/>	<input type="checkbox"/>
4.2.5	Verify that the read only port is not used or is set to 0	<input type="checkbox"/>	<input type="checkbox"/>
4.2.6	Ensure that the --streaming-connection-idle-timeout argument is not set to 0	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
4.2.8	Ensure that the kubeAPIQPS [--event-qps] argument is set to a level which ensures appropriate event capture	<input type="checkbox"/>	<input type="checkbox"/>
4.2.9	Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.10	Ensure that the --rotate-certificates argument is not set to false	<input type="checkbox"/>	<input type="checkbox"/>
4.2.11	Verify that the RotateKubeletServerCertificate argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
4.2.12	Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers	<input type="checkbox"/>	<input type="checkbox"/>
5.1.1	Ensure that the cluster-admin role is only used where required	<input type="checkbox"/>	<input type="checkbox"/>
5.1.2	Minimize access to secrets	<input type="checkbox"/>	<input type="checkbox"/>
5.1.3	Minimize wildcard use in Roles and ClusterRoles	<input type="checkbox"/>	<input type="checkbox"/>
5.1.4	Minimize access to create pods	<input type="checkbox"/>	<input type="checkbox"/>
5.1.5	Ensure that default service accounts are not actively used.	<input type="checkbox"/>	<input type="checkbox"/>
5.1.6	Ensure that Service Account Tokens are only mounted where necessary	<input type="checkbox"/>	<input type="checkbox"/>
5.2.1	Minimize the admission of privileged containers	<input type="checkbox"/>	<input type="checkbox"/>
5.2.2	Minimize the admission of containers wishing to share the host process ID namespace	<input type="checkbox"/>	<input type="checkbox"/>
5.2.4	Minimize the admission of containers wishing to share the host network namespace	<input type="checkbox"/>	<input type="checkbox"/>
5.2.5	Minimize the admission of containers with allowPrivilegeEscalation	<input type="checkbox"/>	<input type="checkbox"/>
5.2.6	Minimize the admission of root containers	<input type="checkbox"/>	<input type="checkbox"/>
5.2.7	Minimize the admission of containers with the NET_RAW capability	<input type="checkbox"/>	<input type="checkbox"/>
5.2.8	Minimize the admission of containers with added capabilities	<input type="checkbox"/>	<input type="checkbox"/>
5.2.9	Minimize the admission of containers with capabilities assigned	<input type="checkbox"/>	<input type="checkbox"/>
5.2.10	Minimize access to privileged Security Context Constraints	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
5.3.2	Ensure that all Namespaces have Network Policies defined	<input type="checkbox"/>	<input type="checkbox"/>
5.4.1	Prefer using secrets as files over secrets as environment variables	<input type="checkbox"/>	<input type="checkbox"/>
5.4.2	Consider external secret storage	<input type="checkbox"/>	<input type="checkbox"/>
5.5.1	Configure Image Provenance using image controller configuration parameters	<input type="checkbox"/>	<input type="checkbox"/>
5.7.2	Ensure that the seccomp profile is set to docker/default in your pod definitions	<input type="checkbox"/>	<input type="checkbox"/>
5.7.3	Apply Security Context to Your Pods and Containers	<input type="checkbox"/>	<input type="checkbox"/>

Appendix: CIS Controls v7 IG 3 Mapped Recommendations

Recommendation		Set Correctly	
		Yes	No
1.1.1	Ensure that the API server pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.2	Ensure that the API server pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.3	Ensure that the controller manager pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.4	Ensure that the controller manager pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.5	Ensure that the scheduler pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.6	Ensure that the scheduler pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.7	Ensure that the etcd pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.8	Ensure that the etcd pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.9	Ensure that the Container Network Interface file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.10	Ensure that the Container Network Interface file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.11	Ensure that the etcd data directory permissions are set to 700 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.12	Ensure that the etcd data directory ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.13	Ensure that the kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.14	Ensure that the kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.15	Ensure that the Scheduler kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.1.16	Ensure that the Scheduler kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.17	Ensure that the Controller Manager kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.18	Ensure that the Controller Manager kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.19	Ensure that the OpenShift PKI directory and file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.20	Ensure that the OpenShift PKI certificate file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.21	Ensure that the OpenShift PKI key file permissions are set to 600	<input type="checkbox"/>	<input type="checkbox"/>
1.2.1	Ensure that anonymous requests are authorized	<input type="checkbox"/>	<input type="checkbox"/>
1.2.2	Use https for kubelet connections	<input type="checkbox"/>	<input type="checkbox"/>
1.2.3	Ensure that the kubelet uses certificates to authenticate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.4	Verify that the kubelet certificate authority is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.5	Ensure that the --authorization-mode argument is not set to AlwaysAllow	<input type="checkbox"/>	<input type="checkbox"/>
1.2.6	Verify that RBAC is enabled	<input type="checkbox"/>	<input type="checkbox"/>
1.2.7	Ensure that the APIPriorityAndFairness feature gate is enabled	<input type="checkbox"/>	<input type="checkbox"/>
1.2.8	Ensure that the admission control plugin AlwaysAdmit is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.9	Ensure that the admission control plugin AlwaysPullImages is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.10	Ensure that the admission control plugin ServiceAccount is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.11	Ensure that the admission control plugin NamespaceLifecycle is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.12	Ensure that the admission control plugin SecurityContextConstraint is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.13	Ensure that the admission control plugin NodeRestriction is set	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.2.14	Ensure that the --insecure-bind-address argument is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.15	Ensure that the --insecure-port argument is set to 0	<input type="checkbox"/>	<input type="checkbox"/>
1.2.16	Ensure that the --secure-port argument is not set to 0	<input type="checkbox"/>	<input type="checkbox"/>
1.2.17	Ensure that the healthz endpoint is protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
1.2.18	Ensure that the --audit-log-path argument is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.19	Ensure that the audit logs are forwarded off the cluster for retention	<input type="checkbox"/>	<input type="checkbox"/>
1.2.20	Ensure that the maximumRetainedFiles argument is set to 10 or as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.21	Configure Kubernetes API Server Maximum Audit Log Size	<input type="checkbox"/>	<input type="checkbox"/>
1.2.22	Ensure that the --request-timeout argument is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.23	Ensure that the --service-account-lookup argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
1.2.24	Ensure that the --service-account-key-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.25	Ensure that the --etcd-certfile and --etcd-keyfile arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.26	Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.27	Ensure that the --client-ca-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.28	Ensure that the --etcd-cafile argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.29	Ensure that encryption providers are appropriately configured	<input type="checkbox"/>	<input type="checkbox"/>
1.2.30	Ensure that the API Server only makes use of Strong Cryptographic Ciphers	<input type="checkbox"/>	<input type="checkbox"/>
1.2.31	Ensure unsupported configuration overrides are not used	<input type="checkbox"/>	<input type="checkbox"/>
1.3.1	Ensure that controller manager healthz endpoints are protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
1.3.2	Ensure that the --use-service-account-credentials argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.3.3	Ensure that the --service-account-private-key-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.3.4	Ensure that the --root-ca-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.4.1	Ensure that the healthz endpoints for the scheduler are protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
1.4.2	Verify that the scheduler API service is protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
2.1	Ensure that the --cert-file and --key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
2.2	Ensure that the --client-cert-auth argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.3	Ensure that the --auto-tls argument is not set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.4	Ensure that the --peer-cert-file and --peer-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
2.5	Ensure that the --peer-client-cert-auth argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.6	Ensure that the --peer-auto-tls argument is not set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.7	Ensure that a unique Certificate Authority is used for etcd	<input type="checkbox"/>	<input type="checkbox"/>
3.1.1	Client certificate authentication should not be used for users	<input type="checkbox"/>	<input type="checkbox"/>
3.2.1	Ensure that a minimal audit policy is created	<input type="checkbox"/>	<input type="checkbox"/>
3.2.2	Ensure that the audit policy covers key security concerns	<input type="checkbox"/>	<input type="checkbox"/>
4.1.1	Ensure that the kubelet service file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.2	Ensure that the kubelet service file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.3	If proxy kube proxy configuration file exists ensure permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.4	If proxy kubeconfig file exists ensure ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.5	Ensure that the --kubeconfig kubelet.conf file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.6	Ensure that the --kubeconfig kubelet.conf file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
4.1.7	Ensure that the certificate authorities file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.8	Ensure that the client certificate authorities file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.9	Ensure that the kubelet --config configuration file has permissions set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.10	Ensure that the kubelet configuration file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.2.1	Activate Garbage collection in OpenShift Container Platform 4, as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.2	Ensure that the --anonymous-auth argument is set to false	<input type="checkbox"/>	<input type="checkbox"/>
4.2.3	Ensure that the --authorization-mode argument is not set to AlwaysAllow	<input type="checkbox"/>	<input type="checkbox"/>
4.2.4	Ensure that the --client-ca-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.5	Verify that the read only port is not used or is set to 0	<input type="checkbox"/>	<input type="checkbox"/>
4.2.6	Ensure that the --streaming-connection-idle-timeout argument is not set to 0	<input type="checkbox"/>	<input type="checkbox"/>
4.2.7	Ensure that the --make-iptables-util-chains argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
4.2.8	Ensure that the kubeAPIQPS [--event-qps] argument is set to a level which ensures appropriate event capture	<input type="checkbox"/>	<input type="checkbox"/>
4.2.9	Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.10	Ensure that the --rotate-certificates argument is not set to false	<input type="checkbox"/>	<input type="checkbox"/>
4.2.11	Verify that the RotateKubeletServerCertificate argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
4.2.12	Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers	<input type="checkbox"/>	<input type="checkbox"/>
5.1.1	Ensure that the cluster-admin role is only used where required	<input type="checkbox"/>	<input type="checkbox"/>
5.1.2	Minimize access to secrets	<input type="checkbox"/>	<input type="checkbox"/>
5.1.3	Minimize wildcard use in Roles and ClusterRoles	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
5.1.4	Minimize access to create pods	<input type="checkbox"/>	<input type="checkbox"/>
5.1.5	Ensure that default service accounts are not actively used.	<input type="checkbox"/>	<input type="checkbox"/>
5.1.6	Ensure that Service Account Tokens are only mounted where necessary	<input type="checkbox"/>	<input type="checkbox"/>
5.2.1	Minimize the admission of privileged containers	<input type="checkbox"/>	<input type="checkbox"/>
5.2.2	Minimize the admission of containers wishing to share the host process ID namespace	<input type="checkbox"/>	<input type="checkbox"/>
5.2.3	Minimize the admission of containers wishing to share the host IPC namespace	<input type="checkbox"/>	<input type="checkbox"/>
5.2.4	Minimize the admission of containers wishing to share the host network namespace	<input type="checkbox"/>	<input type="checkbox"/>
5.2.5	Minimize the admission of containers with allowPrivilegeEscalation	<input type="checkbox"/>	<input type="checkbox"/>
5.2.6	Minimize the admission of root containers	<input type="checkbox"/>	<input type="checkbox"/>
5.2.7	Minimize the admission of containers with the NET_RAW capability	<input type="checkbox"/>	<input type="checkbox"/>
5.2.8	Minimize the admission of containers with added capabilities	<input type="checkbox"/>	<input type="checkbox"/>
5.2.9	Minimize the admission of containers with capabilities assigned	<input type="checkbox"/>	<input type="checkbox"/>
5.2.10	Minimize access to privileged Security Context Constraints	<input type="checkbox"/>	<input type="checkbox"/>
5.3.1	Ensure that the CNI in use supports Network Policies	<input type="checkbox"/>	<input type="checkbox"/>
5.3.2	Ensure that all Namespaces have Network Policies defined	<input type="checkbox"/>	<input type="checkbox"/>
5.4.1	Prefer using secrets as files over secrets as environment variables	<input type="checkbox"/>	<input type="checkbox"/>
5.4.2	Consider external secret storage	<input type="checkbox"/>	<input type="checkbox"/>
5.5.1	Configure Image Provenance using image controller configuration parameters	<input type="checkbox"/>	<input type="checkbox"/>
5.7.1	Create administrative boundaries between resources using namespaces	<input type="checkbox"/>	<input type="checkbox"/>
5.7.2	Ensure that the seccomp profile is set to docker/default in your pod definitions	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
5.7.3	Apply Security Context to Your Pods and Containers	<input type="checkbox"/>	<input type="checkbox"/>

Appendix: CIS Controls v7 Unmapped Recommendations

Recommendation		Set Correctly	
		Yes	No
	No unmapped recommendations to CIS Controls v7	<input type="checkbox"/>	<input type="checkbox"/>

Appendix: CIS Controls v8 IG 1 Mapped Recommendations

Recommendation		Set Correctly	
		Yes	No
1.1.1	Ensure that the API server pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.2	Ensure that the API server pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.3	Ensure that the controller manager pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.4	Ensure that the controller manager pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.5	Ensure that the scheduler pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.6	Ensure that the scheduler pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.7	Ensure that the etcd pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.8	Ensure that the etcd pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.9	Ensure that the Container Network Interface file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.10	Ensure that the Container Network Interface file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.11	Ensure that the etcd data directory permissions are set to 700 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.12	Ensure that the etcd data directory ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.13	Ensure that the kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.14	Ensure that the kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.15	Ensure that the Scheduler kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.1.16	Ensure that the Scheduler kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.17	Ensure that the Controller Manager kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.18	Ensure that the Controller Manager kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.19	Ensure that the OpenShift PKI directory and file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.20	Ensure that the OpenShift PKI certificate file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.21	Ensure that the OpenShift PKI key file permissions are set to 600	<input type="checkbox"/>	<input type="checkbox"/>
1.2.1	Ensure that anonymous requests are authorized	<input type="checkbox"/>	<input type="checkbox"/>
1.2.5	Ensure that the --authorization-mode argument is not set to AlwaysAllow	<input type="checkbox"/>	<input type="checkbox"/>
1.2.8	Ensure that the admission control plugin AlwaysAdmit is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.9	Ensure that the admission control plugin AlwaysPullImages is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.10	Ensure that the admission control plugin ServiceAccount is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.11	Ensure that the admission control plugin NamespaceLifecycle is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.19	Ensure that the audit logs are forwarded off the cluster for retention	<input type="checkbox"/>	<input type="checkbox"/>
1.2.20	Ensure that the maximumRetainedFiles argument is set to 10 or as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.21	Configure Kubernetes API Server Maximum Audit Log Size	<input type="checkbox"/>	<input type="checkbox"/>
1.2.22	Ensure that the --request-timeout argument is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.23	Ensure that the --service-account-lookup argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
1.2.24	Ensure that the --service-account-key-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.31	Ensure unsupported configuration overrides are not used	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.3.2	Ensure that the --use-service-account-credentials argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
1.3.3	Ensure that the --service-account-private-key-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
2.6	Ensure that the --peer-auto-tls argument is not set to true	<input type="checkbox"/>	<input type="checkbox"/>
3.2.1	Ensure that a minimal audit policy is created	<input type="checkbox"/>	<input type="checkbox"/>
4.1.1	Ensure that the kubelet service file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.2	Ensure that the kubelet service file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.3	If proxy kube proxy configuration file exists ensure permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.4	If proxy kubeconfig file exists ensure ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.5	Ensure that the --kubeconfig kubelet.conf file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.6	Ensure that the --kubeconfig kubelet.conf file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.7	Ensure that the certificate authorities file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.8	Ensure that the client certificate authorities file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.9	Ensure that the kubelet --config configuration file has permissions set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.10	Ensure that the kubelet configuration file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.2.2	Ensure that the --anonymous-auth argument is set to false	<input type="checkbox"/>	<input type="checkbox"/>
4.2.3	Ensure that the --authorization-mode argument is not set to AlwaysAllow	<input type="checkbox"/>	<input type="checkbox"/>
4.2.7	Ensure that the --make-iptables-util-chains argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
4.2.12	Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers	<input type="checkbox"/>	<input type="checkbox"/>
5.1.1	Ensure that the cluster-admin role is only used where required	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
5.1.5	Ensure that default service accounts are not actively used.	<input type="checkbox"/>	<input type="checkbox"/>
5.2.1	Minimize the admission of privileged containers	<input type="checkbox"/>	<input type="checkbox"/>
5.2.2	Minimize the admission of containers wishing to share the host process ID namespace	<input type="checkbox"/>	<input type="checkbox"/>
5.2.5	Minimize the admission of containers with allowPrivilegeEscalation	<input type="checkbox"/>	<input type="checkbox"/>
5.2.6	Minimize the admission of root containers	<input type="checkbox"/>	<input type="checkbox"/>
5.3.1	Ensure that the CNI in use supports Network Policies	<input type="checkbox"/>	<input type="checkbox"/>
5.3.2	Ensure that all Namespaces have Network Policies defined	<input type="checkbox"/>	<input type="checkbox"/>
5.5.1	Configure Image Provenance using image controller configuration parameters	<input type="checkbox"/>	<input type="checkbox"/>
5.7.1	Create administrative boundaries between resources using namespaces	<input type="checkbox"/>	<input type="checkbox"/>
5.7.2	Ensure that the seccomp profile is set to docker/default in your pod definitions	<input type="checkbox"/>	<input type="checkbox"/>
5.7.3	Apply Security Context to Your Pods and Containers	<input type="checkbox"/>	<input type="checkbox"/>

Appendix: CIS Controls v8 IG 2 Mapped Recommendations

Recommendation		Set Correctly	
		Yes	No
1.1.1	Ensure that the API server pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.2	Ensure that the API server pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.3	Ensure that the controller manager pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.4	Ensure that the controller manager pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.5	Ensure that the scheduler pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.6	Ensure that the scheduler pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.7	Ensure that the etcd pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.8	Ensure that the etcd pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.9	Ensure that the Container Network Interface file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.10	Ensure that the Container Network Interface file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.11	Ensure that the etcd data directory permissions are set to 700 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.12	Ensure that the etcd data directory ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.13	Ensure that the kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.14	Ensure that the kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.15	Ensure that the Scheduler kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.1.16	Ensure that the Scheduler kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.17	Ensure that the Controller Manager kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.18	Ensure that the Controller Manager kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.19	Ensure that the OpenShift PKI directory and file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.20	Ensure that the OpenShift PKI certificate file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.21	Ensure that the OpenShift PKI key file permissions are set to 600	<input type="checkbox"/>	<input type="checkbox"/>
1.2.1	Ensure that anonymous requests are authorized	<input type="checkbox"/>	<input type="checkbox"/>
1.2.2	Use https for kubelet connections	<input type="checkbox"/>	<input type="checkbox"/>
1.2.3	Ensure that the kubelet uses certificates to authenticate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.4	Verify that the kubelet certificate authority is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.5	Ensure that the --authorization-mode argument is not set to AlwaysAllow	<input type="checkbox"/>	<input type="checkbox"/>
1.2.7	Ensure that the APIPriorityAndFairness feature gate is enabled	<input type="checkbox"/>	<input type="checkbox"/>
1.2.8	Ensure that the admission control plugin AlwaysAdmit is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.9	Ensure that the admission control plugin AlwaysPullImages is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.10	Ensure that the admission control plugin ServiceAccount is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.11	Ensure that the admission control plugin NamespaceLifecycle is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.12	Ensure that the admission control plugin SecurityContextConstraint is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.13	Ensure that the admission control plugin NodeRestriction is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.14	Ensure that the --insecure-bind-address argument is not set	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.2.15	Ensure that the --insecure-port argument is set to 0	<input type="checkbox"/>	<input type="checkbox"/>
1.2.16	Ensure that the --secure-port argument is not set to 0	<input type="checkbox"/>	<input type="checkbox"/>
1.2.18	Ensure that the --audit-log-path argument is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.19	Ensure that the audit logs are forwarded off the cluster for retention	<input type="checkbox"/>	<input type="checkbox"/>
1.2.20	Ensure that the maximumRetainedFiles argument is set to 10 or as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.21	Configure Kubernetes API Server Maximum Audit Log Size	<input type="checkbox"/>	<input type="checkbox"/>
1.2.22	Ensure that the --request-timeout argument is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.23	Ensure that the --service-account-lookup argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
1.2.24	Ensure that the --service-account-key-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.25	Ensure that the --etcd-certfile and --etcd-keyfile arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.26	Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.27	Ensure that the --client-ca-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.28	Ensure that the --etcd-cafile argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.29	Ensure that encryption providers are appropriately configured	<input type="checkbox"/>	<input type="checkbox"/>
1.2.30	Ensure that the API Server only makes use of Strong Cryptographic Ciphers	<input type="checkbox"/>	<input type="checkbox"/>
1.2.31	Ensure unsupported configuration overrides are not used	<input type="checkbox"/>	<input type="checkbox"/>
1.3.2	Ensure that the --use-service-account-credentials argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
1.3.3	Ensure that the --service-account-private-key-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.4.2	Verify that the scheduler API service is protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
2.1	Ensure that the --cert-file and --key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
2.2	Ensure that the --client-cert-auth argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.3	Ensure that the --auto-tls argument is not set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.4	Ensure that the --peer-cert-file and --peer-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
2.6	Ensure that the --peer-auto-tls argument is not set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.7	Ensure that a unique Certificate Authority is used for etcd	<input type="checkbox"/>	<input type="checkbox"/>
3.1.1	Client certificate authentication should not be used for users	<input type="checkbox"/>	<input type="checkbox"/>
3.2.1	Ensure that a minimal audit policy is created	<input type="checkbox"/>	<input type="checkbox"/>
3.2.2	Ensure that the audit policy covers key security concerns	<input type="checkbox"/>	<input type="checkbox"/>
4.1.1	Ensure that the kubelet service file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.2	Ensure that the kubelet service file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.3	If proxy kube proxy configuration file exists ensure permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.4	If proxy kubeconfig file exists ensure ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.5	Ensure that the --kubeconfig kubelet.conf file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.6	Ensure that the --kubeconfig kubelet.conf file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.7	Ensure that the certificate authorities file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.8	Ensure that the client certificate authorities file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.9	Ensure that the kubelet --config configuration file has permissions set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.10	Ensure that the kubelet configuration file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.2.1	Activate Garbage collection in OpenShift Container Platform 4, as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.2	Ensure that the --anonymous-auth argument is set to false	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
4.2.3	Ensure that the --authorization-mode argument is not set to AlwaysAllow	<input type="checkbox"/>	<input type="checkbox"/>
4.2.4	Ensure that the --client-ca-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.5	Verify that the read only port is not used or is set to 0	<input type="checkbox"/>	<input type="checkbox"/>
4.2.6	Ensure that the --streaming-connection-idle-timeout argument is not set to 0	<input type="checkbox"/>	<input type="checkbox"/>
4.2.7	Ensure that the --make-iptables-util-chains argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
4.2.8	Ensure that the kubeAPIQPS [--event-qps] argument is set to a level which ensures appropriate event capture	<input type="checkbox"/>	<input type="checkbox"/>
4.2.9	Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.10	Ensure that the --rotate-certificates argument is not set to false	<input type="checkbox"/>	<input type="checkbox"/>
4.2.11	Verify that the RotateKubeletServerCertificate argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
4.2.12	Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers	<input type="checkbox"/>	<input type="checkbox"/>
5.1.1	Ensure that the cluster-admin role is only used where required	<input type="checkbox"/>	<input type="checkbox"/>
5.1.2	Minimize access to secrets	<input type="checkbox"/>	<input type="checkbox"/>
5.1.5	Ensure that default service accounts are not actively used.	<input type="checkbox"/>	<input type="checkbox"/>
5.1.6	Ensure that Service Account Tokens are only mounted where necessary	<input type="checkbox"/>	<input type="checkbox"/>
5.2.1	Minimize the admission of privileged containers	<input type="checkbox"/>	<input type="checkbox"/>
5.2.2	Minimize the admission of containers wishing to share the host process ID namespace	<input type="checkbox"/>	<input type="checkbox"/>
5.2.4	Minimize the admission of containers wishing to share the host network namespace	<input type="checkbox"/>	<input type="checkbox"/>
5.2.5	Minimize the admission of containers with allowPrivilegeEscalation	<input type="checkbox"/>	<input type="checkbox"/>
5.2.6	Minimize the admission of root containers	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
5.2.7	Minimize the admission of containers with the NET_RAW capability	<input type="checkbox"/>	<input type="checkbox"/>
5.2.8	Minimize the admission of containers with added capabilities	<input type="checkbox"/>	<input type="checkbox"/>
5.2.9	Minimize the admission of containers with capabilities assigned	<input type="checkbox"/>	<input type="checkbox"/>
5.3.1	Ensure that the CNI in use supports Network Policies	<input type="checkbox"/>	<input type="checkbox"/>
5.3.2	Ensure that all Namespaces have Network Policies defined	<input type="checkbox"/>	<input type="checkbox"/>
5.4.1	Prefer using secrets as files over secrets as environment variables	<input type="checkbox"/>	<input type="checkbox"/>
5.4.2	Consider external secret storage	<input type="checkbox"/>	<input type="checkbox"/>
5.5.1	Configure Image Provenance using image controller configuration parameters	<input type="checkbox"/>	<input type="checkbox"/>
5.7.1	Create administrative boundaries between resources using namespaces	<input type="checkbox"/>	<input type="checkbox"/>
5.7.2	Ensure that the seccomp profile is set to docker/default in your pod definitions	<input type="checkbox"/>	<input type="checkbox"/>
5.7.3	Apply Security Context to Your Pods and Containers	<input type="checkbox"/>	<input type="checkbox"/>
5.7.4	The default namespace should not be used	<input type="checkbox"/>	<input type="checkbox"/>

Appendix: CIS Controls v8 IG 3 Mapped Recommendations

Recommendation		Set Correctly	
		Yes	No
1.1.1	Ensure that the API server pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.2	Ensure that the API server pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.3	Ensure that the controller manager pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.4	Ensure that the controller manager pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.5	Ensure that the scheduler pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.6	Ensure that the scheduler pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.7	Ensure that the etcd pod specification file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.8	Ensure that the etcd pod specification file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.9	Ensure that the Container Network Interface file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.10	Ensure that the Container Network Interface file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.11	Ensure that the etcd data directory permissions are set to 700 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.12	Ensure that the etcd data directory ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.13	Ensure that the kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.14	Ensure that the kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.15	Ensure that the Scheduler kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.1.16	Ensure that the Scheduler kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.17	Ensure that the Controller Manager kubeconfig file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.18	Ensure that the Controller Manager kubeconfig file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.19	Ensure that the OpenShift PKI directory and file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
1.1.20	Ensure that the OpenShift PKI certificate file permissions are set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
1.1.21	Ensure that the OpenShift PKI key file permissions are set to 600	<input type="checkbox"/>	<input type="checkbox"/>
1.2.1	Ensure that anonymous requests are authorized	<input type="checkbox"/>	<input type="checkbox"/>
1.2.2	Use https for kubelet connections	<input type="checkbox"/>	<input type="checkbox"/>
1.2.3	Ensure that the kubelet uses certificates to authenticate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.4	Verify that the kubelet certificate authority is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.5	Ensure that the --authorization-mode argument is not set to AlwaysAllow	<input type="checkbox"/>	<input type="checkbox"/>
1.2.6	Verify that RBAC is enabled	<input type="checkbox"/>	<input type="checkbox"/>
1.2.7	Ensure that the APIPriorityAndFairness feature gate is enabled	<input type="checkbox"/>	<input type="checkbox"/>
1.2.8	Ensure that the admission control plugin AlwaysAdmit is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.9	Ensure that the admission control plugin AlwaysPullImages is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.10	Ensure that the admission control plugin ServiceAccount is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.11	Ensure that the admission control plugin NamespaceLifecycle is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.12	Ensure that the admission control plugin SecurityContextConstraint is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.13	Ensure that the admission control plugin NodeRestriction is set	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.2.14	Ensure that the --insecure-bind-address argument is not set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.15	Ensure that the --insecure-port argument is set to 0	<input type="checkbox"/>	<input type="checkbox"/>
1.2.16	Ensure that the --secure-port argument is not set to 0	<input type="checkbox"/>	<input type="checkbox"/>
1.2.17	Ensure that the healthz endpoint is protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
1.2.18	Ensure that the --audit-log-path argument is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.19	Ensure that the audit logs are forwarded off the cluster for retention	<input type="checkbox"/>	<input type="checkbox"/>
1.2.20	Ensure that the maximumRetainedFiles argument is set to 10 or as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.21	Configure Kubernetes API Server Maximum Audit Log Size	<input type="checkbox"/>	<input type="checkbox"/>
1.2.22	Ensure that the --request-timeout argument is set	<input type="checkbox"/>	<input type="checkbox"/>
1.2.23	Ensure that the --service-account-lookup argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
1.2.24	Ensure that the --service-account-key-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.25	Ensure that the --etcd-certfile and --etcd-keyfile arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.26	Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.27	Ensure that the --client-ca-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.28	Ensure that the --etcd-cafile argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.2.29	Ensure that encryption providers are appropriately configured	<input type="checkbox"/>	<input type="checkbox"/>
1.2.30	Ensure that the API Server only makes use of Strong Cryptographic Ciphers	<input type="checkbox"/>	<input type="checkbox"/>
1.2.31	Ensure unsupported configuration overrides are not used	<input type="checkbox"/>	<input type="checkbox"/>
1.3.1	Ensure that controller manager healthz endpoints are protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
1.3.2	Ensure that the --use-service-account-credentials argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
1.3.3	Ensure that the --service-account-private-key-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.3.4	Ensure that the --root-ca-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
1.4.1	Ensure that the healthz endpoints for the scheduler are protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
1.4.2	Verify that the scheduler API service is protected by RBAC	<input type="checkbox"/>	<input type="checkbox"/>
2.1	Ensure that the --cert-file and --key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
2.2	Ensure that the --client-cert-auth argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.3	Ensure that the --auto-tls argument is not set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.4	Ensure that the --peer-cert-file and --peer-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
2.5	Ensure that the --peer-client-cert-auth argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.6	Ensure that the --peer-auto-tls argument is not set to true	<input type="checkbox"/>	<input type="checkbox"/>
2.7	Ensure that a unique Certificate Authority is used for etcd	<input type="checkbox"/>	<input type="checkbox"/>
3.1.1	Client certificate authentication should not be used for users	<input type="checkbox"/>	<input type="checkbox"/>
3.2.1	Ensure that a minimal audit policy is created	<input type="checkbox"/>	<input type="checkbox"/>
3.2.2	Ensure that the audit policy covers key security concerns	<input type="checkbox"/>	<input type="checkbox"/>
4.1.1	Ensure that the kubelet service file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.2	Ensure that the kubelet service file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.3	If proxy kube proxy configuration file exists ensure permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.4	If proxy kubeconfig file exists ensure ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.5	Ensure that the --kubeconfig kubelet.conf file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.6	Ensure that the --kubeconfig kubelet.conf file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
4.1.7	Ensure that the certificate authorities file permissions are set to 644 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.8	Ensure that the client certificate authorities file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.1.9	Ensure that the kubelet --config configuration file has permissions set to 600 or more restrictive	<input type="checkbox"/>	<input type="checkbox"/>
4.1.10	Ensure that the kubelet configuration file ownership is set to root:root	<input type="checkbox"/>	<input type="checkbox"/>
4.2.1	Activate Garbage collection in OpenShift Container Platform 4, as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.2	Ensure that the --anonymous-auth argument is set to false	<input type="checkbox"/>	<input type="checkbox"/>
4.2.3	Ensure that the --authorization-mode argument is not set to AlwaysAllow	<input type="checkbox"/>	<input type="checkbox"/>
4.2.4	Ensure that the --client-ca-file argument is set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.5	Verify that the read only port is not used or is set to 0	<input type="checkbox"/>	<input type="checkbox"/>
4.2.6	Ensure that the --streaming-connection-idle-timeout argument is not set to 0	<input type="checkbox"/>	<input type="checkbox"/>
4.2.7	Ensure that the --make-iptables-util-chains argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
4.2.8	Ensure that the kubeAPIQPS [--event-qps] argument is set to a level which ensures appropriate event capture	<input type="checkbox"/>	<input type="checkbox"/>
4.2.9	Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate	<input type="checkbox"/>	<input type="checkbox"/>
4.2.10	Ensure that the --rotate-certificates argument is not set to false	<input type="checkbox"/>	<input type="checkbox"/>
4.2.11	Verify that the RotateKubeletServerCertificate argument is set to true	<input type="checkbox"/>	<input type="checkbox"/>
4.2.12	Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers	<input type="checkbox"/>	<input type="checkbox"/>
5.1.1	Ensure that the cluster-admin role is only used where required	<input type="checkbox"/>	<input type="checkbox"/>
5.1.2	Minimize access to secrets	<input type="checkbox"/>	<input type="checkbox"/>
5.1.3	Minimize wildcard use in Roles and ClusterRoles	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
5.1.4	Minimize access to create pods	<input type="checkbox"/>	<input type="checkbox"/>
5.1.5	Ensure that default service accounts are not actively used.	<input type="checkbox"/>	<input type="checkbox"/>
5.1.6	Ensure that Service Account Tokens are only mounted where necessary	<input type="checkbox"/>	<input type="checkbox"/>
5.2.1	Minimize the admission of privileged containers	<input type="checkbox"/>	<input type="checkbox"/>
5.2.2	Minimize the admission of containers wishing to share the host process ID namespace	<input type="checkbox"/>	<input type="checkbox"/>
5.2.3	Minimize the admission of containers wishing to share the host IPC namespace	<input type="checkbox"/>	<input type="checkbox"/>
5.2.4	Minimize the admission of containers wishing to share the host network namespace	<input type="checkbox"/>	<input type="checkbox"/>
5.2.5	Minimize the admission of containers with allowPrivilegeEscalation	<input type="checkbox"/>	<input type="checkbox"/>
5.2.6	Minimize the admission of root containers	<input type="checkbox"/>	<input type="checkbox"/>
5.2.7	Minimize the admission of containers with the NET_RAW capability	<input type="checkbox"/>	<input type="checkbox"/>
5.2.8	Minimize the admission of containers with added capabilities	<input type="checkbox"/>	<input type="checkbox"/>
5.2.9	Minimize the admission of containers with capabilities assigned	<input type="checkbox"/>	<input type="checkbox"/>
5.2.10	Minimize access to privileged Security Context Constraints	<input type="checkbox"/>	<input type="checkbox"/>
5.3.1	Ensure that the CNI in use supports Network Policies	<input type="checkbox"/>	<input type="checkbox"/>
5.3.2	Ensure that all Namespaces have Network Policies defined	<input type="checkbox"/>	<input type="checkbox"/>
5.4.1	Prefer using secrets as files over secrets as environment variables	<input type="checkbox"/>	<input type="checkbox"/>
5.4.2	Consider external secret storage	<input type="checkbox"/>	<input type="checkbox"/>
5.5.1	Configure Image Provenance using image controller configuration parameters	<input type="checkbox"/>	<input type="checkbox"/>
5.7.1	Create administrative boundaries between resources using namespaces	<input type="checkbox"/>	<input type="checkbox"/>
5.7.2	Ensure that the seccomp profile is set to docker/default in your pod definitions	<input type="checkbox"/>	<input type="checkbox"/>

Recommendation		Set Correctly	
		Yes	No
5.7.3	Apply Security Context to Your Pods and Containers	<input type="checkbox"/>	<input type="checkbox"/>
5.7.4	The default namespace should not be used	<input type="checkbox"/>	<input type="checkbox"/>

Appendix: CIS Controls v8 Unmapped Recommendations

Recommendation		Set Correctly	
		Yes	No
	No unmapped recommendations to CIS Controls v8	<input type="checkbox"/>	<input type="checkbox"/>

Appendix: Change History

Date	Version	Changes for this version
12/13/2025	V1.9.0	Tested all AAC against the latest OpenShift cluster version 4.20
12/13/2025	V1.9.0	Updated recommendation 1.2.4 and cert path Ticket # 24083
12/11/2025	V1.9.0	Edited Guidance for Section 5.2 Ticket #24084
12/03/2025	V1.9.0	Edited impact statement for recommendation 4.2.8 on Ticket # 25808
12/03/2025	V1.9.0	Edited audit procedure for recommendation 1.1.12 on Ticket # 24544
11/23/2025	V1.9.0	Edited audit procedure for recommendation 1.2.29 on Ticket # 24003
11/23/2025	V1.9.0	Edited ownership requirement for recommendation 5.7.4 on Ticket # 24108
6/19/2025	V1.8.0	Tested all AAC against the latest OpenShift cluster version 4.18
6/16/2025	V1.8.0	Modified recommendation 1.1.12 audit to reflect correct ownership values.
6/15/2025	V1.8.0	Updated recommendation reference Links Ticket # 21290

Date	Version	Changes for this version
3/1/2025	V1.8.0	Added OVN reference Ticket #18799
12/1/2024	V1.7.0	Tested all AAC against the latest OpenShift cluster version 4.16.27
6/1/2023	V1.6.0	Tested all AAC against the latest OpenShift cluster version 4.15
5/14/2024	V1.6.0	Updated and Edited recommendation 1.2.9 Ticket #21723
5/12/2024	V1.6.0	Edited assessment content for recommendation 4.2.4 Ticket #20799
4/05/2024	V1.6.0	Reviewed and modified assessment methods and sce for recommendation 4.1.4 Ticket #20798
11/30/2023	V1.5.0	Tested all AAC against the latest OpenShift cluster version 4.14
10/12/2023	V1.5.0	Updated file permission guidelines
9/22/2023	V1.5.0	Recommendations with multiple levels set were corrected to have either level 1 or level 2
6/15/2023	V1.4.0	Made Audit Procedure Clearer for recommendation 1.2.32 on Ticket # 18870
6/15/2023	V1.4.0	Renamed Section 5.2 Pod Security Section "Security Context Constraints" on Ticket #19012

Date	Version	Changes for this version
6/15/2023	V1.4.0	Updated Audit Command to reflect the correct path for recommendation 1.2.20 on Ticket #18943
6/15/2023	V1.4.0	Fixed literal formatting in the audit procedure for recommendation 1.4.2 on Ticket #18986
6/15/2023	V1.4.0	Updated the audit procedure and reference links to the OpenShift documentation to reflect latest version. Ticket # 18987
6/15/2023	V1.4.0	Revised audit procedure for recommendation 1.1.15 on Ticket # 18978
6/15/2023	V1.4.0	Revised audit procedure for recommendation 1.2.22 on Ticket # 18893
6/15/2023	V1.4.0	Edited recommendation 1.2.24 to reflect the new default value configuration on Ticket # 18967
6/15/2023	V1.4.0	Updated reference links to latest documentation on Ticket # 18989
6/09/2023	V1.4.0	Edited audit procedure for recommendation 1.1.17 on Ticket # 18979
6/09/2023	V1.4.0	Removed unsupported overrides portion of the audit procedure for recommendation 1.2.10 on Ticket #18988

Date	Version	Changes for this version
6/09/2023	V1.4.0	Edited recommendation 1.1.23 with revised verbiage to reflect latest version options. Ticket # 18972
6/09/2023	V1.4.0	Revised path for Kubelet CA on recommendation 1.2.6 on Ticket #18929
6/09/2023	V1.4.0	Revised path for certificates for recommendation 1.2.5 on Ticket # 18928
6/05/2023	V1.4.0	Updated default section of recommendation 1.2.9 on Ticket # 18880
6/05/2023	V1.4.0	Modified verification settings to reflect OCP 4.6+ for recommendation 1.2.4 on Ticket #18873
6/05/2023	V1.4.0	Removed references to end-of-life versions of OCP.
6/01/2023	V1.4.0	Revised the insecure port option to specifically reference OCP v4.6+ on Ticket #18892
6/01/2023	V1.4.0	Removed the recommendation that references Kubelet Server Certificates from the Master Node section on Ticket # 18891
6/01/2023	V1.4.0	Adjustment to the audit process made for recommendation 1.2.16 on Ticket # 18888

Date	Version	Changes for this version
6/01/2023	V1.4.0	Added aescgm encryption reference for recommendation 1.2.31 on Ticket #18869
6/01/2023	V1.4.0	Updated remediation default values and audit procedure for recommendation 4.2.12 on Ticket # 18868
6/01/2023	V1.4.0	Simplified the audit procedure to use kubeletconfig API for recommendation 4.2.11 on Ticket # 18867
6/01/2023	V1.4.0	Edited default value since RotateKubeletClientCertificate is enabled by default.
6/01/2023	V1.4.0	Removed the SecurityContextDeny reference in section 5.2 since it is not used.
6/01/2023	V1.4.0	Simplified audit procedure for recommendation 1.2.8 on Ticket # 18879
6/01/2023	V1.4.0	Simplified the RBAC portion of the audit for recommendation 1.2.7 on Ticket # 18877
5/22/2023	V1.4.0	Removed Hostname-override recommendation as it is now always set. Ticket # 18874
5/22/2023	V1.4.0	Improved audit procedure to use runtime config with configz for recommendation 4.2.8 on Ticket # 18864

Date	Version	Changes for this version
5/22/2023	V1.4.0	Simplified audit procedure to use secure proxy configz for recommendation 4.2.7 on Ticket # 18863
5/22/2023	V1.4.0	Added a more detailed audit procedure for recommendation 4.2.2 on Ticket # 18862
5/22/2023	V1.4.0	Updated the audit procedure to pull the configuration from each node instead of machineconfigpools for recommendation 4.2.1 on Ticket # 18861
5/19/2023	V1.4.0	Improved recommendation 4.1.6 on Ticket # 18860
5/19/2023	V1.4.0	Added section to address reviewing access to privileged SCCs on Ticket # 18795
5/19/2023	V1.4.0	Updated defaults for recommendation 5.7.2 on Ticket # 18726
5/09/2023	V1.4.0	Removed recommendation 2.8 as it is same as 1.2.33 on Ticket # 18855
5/05/2023	V1.4.0	Improved remediation of active garage collection process for recommendation 4.2.1 on Ticket # 18854
5/05/2023	V1.4.0	Updated the audit procedure to get the streaming connection timeout runtime config for recommendation 4.2.6 on Ticket #

Date	Version	Changes for this version
5/05/2023	V1.4.0	Simplified audit procedure to only query config maps on Ticket # 18806
5/05/2023	V1.4.0	Updated the audit procedure for client CA file based on the runtime config for recommendation 4.2.4 on Ticket # 18805
5/03/2023	V1.4.0	Worker Node configuration file recommendations edited to require permissions are set to 644 or more restrictive.
5/03/2023	V1.4.0	Edited recommendation 4.1.7 to use runtime CA file on ticket # 18804
5/03/2023	V1.4.0	Fixed literal formatting in the audit procedure 4.1.2 on Ticket # 18797
5/03/2023	V1.4.0	Edited recommendation 5.4.1 to revise the default object returned
5/03/2023	V1.4.0	Updated the network policy reference link for recommendation 5.3.2 on Ticket # 18708
5/01/2023	V1.4.0	Updated audit procedure for recommendation 5.3.1 for Ticket # 18707
4/16/2023	V1.4.0	Modified the audit procedure to lookup all resources in the default namespace for recommendation 5.7.4 on Ticket # 18730

Date	Version	Changes for this version
4/16/2023	V1.4.0	Added audit procedure for inspecting Pods with privileged SCCs. Recommendation 5.7.3 on Ticket # 18729
4/16/2023	V1.4.0	Added checking image configuration to recommendation 5.5.1 on Ticket # 18722
4/16/2023	V1.4.0	Provided an Audit Procedure for finding pods that automatically mount service tokens for recommendation 5.1.6 on Ticket # 18706
4/16/2023	V1.4.0	Added edit procedure to help users find roles that have permissions to act on pods. Recommendation 5.1.4 on Ticket # 18705
4/16/2023	V1.4.0	Improved audit command for recommendation 5.1.3 on Ticket # 18760
4/16/2023	V1.4.0	Updated audit procedure for recommendation 5.1.2 on Ticket # 18761
4/16/2023	V1.4.0	Simplified clusterrole audit process for recommendation 5.1.3 on Ticket # 18704
4/01/2023	V1.4.0	Corrected mismatched command for inspecting roles for recommendation 5.1.3 on Ticket # 18703

Date	Version	Changes for this version
4/01/2023	V1.4.0	Created audit process for finding cluster roles that have access to secrets. Recommendation 5.1.2 on Ticket # 18702
4/01/2023	V1.4.0	Added references to OVN to recommendation 4.1.3 on Ticket # 18799
4/01/2023	V1.4.0	Edited recommendation 4.2.3 to make it easier to audit and remediate. Ticket # 18803
4/01/2023	V1.4.0	Updated the default section of recommendation 5.2.1 on Ticket # 18762
4/01/2023	V1.4.0	Supplied an audit procedure to assist users in finding non-default namespaces. Ticket # 18724

